

Discussion I

Aniket Shivam

Quick reminder

- Lab Assignment I is posted!
- Office Hours:
 - Both Lab Sections are treated as Office Hours.
 - Via Zoom
 - Plus post questions on Piazza


Parallelism

- Data Parallelism
 - Same task(code) on different data
 - Ex: Find a word in a document.
- Task Parallelism
 - Different task(code) on same/different data
 - Ex: Calculate Mean, Median, Mode of some data

Data Parallelism

- Addition of two vectors:

```
for i = 1 to N
    S: A[i] = B[i] + C[i]
end for
```



- Workload balancing and Data Shared
- Machine: Multi-core architecture with 4 cores
- Approach: Create 4 threads, divide dataset into 4 (No shared memory).
- SIMD(Single Instruction Multiple Data)

Task Parallelism

- Four tasks to carry out on same dataset


- Data may be shared/copied between threads

- MIMD(Multiple Instruction Multiple Data)
- MISD(Multiple Instruction Single Data)

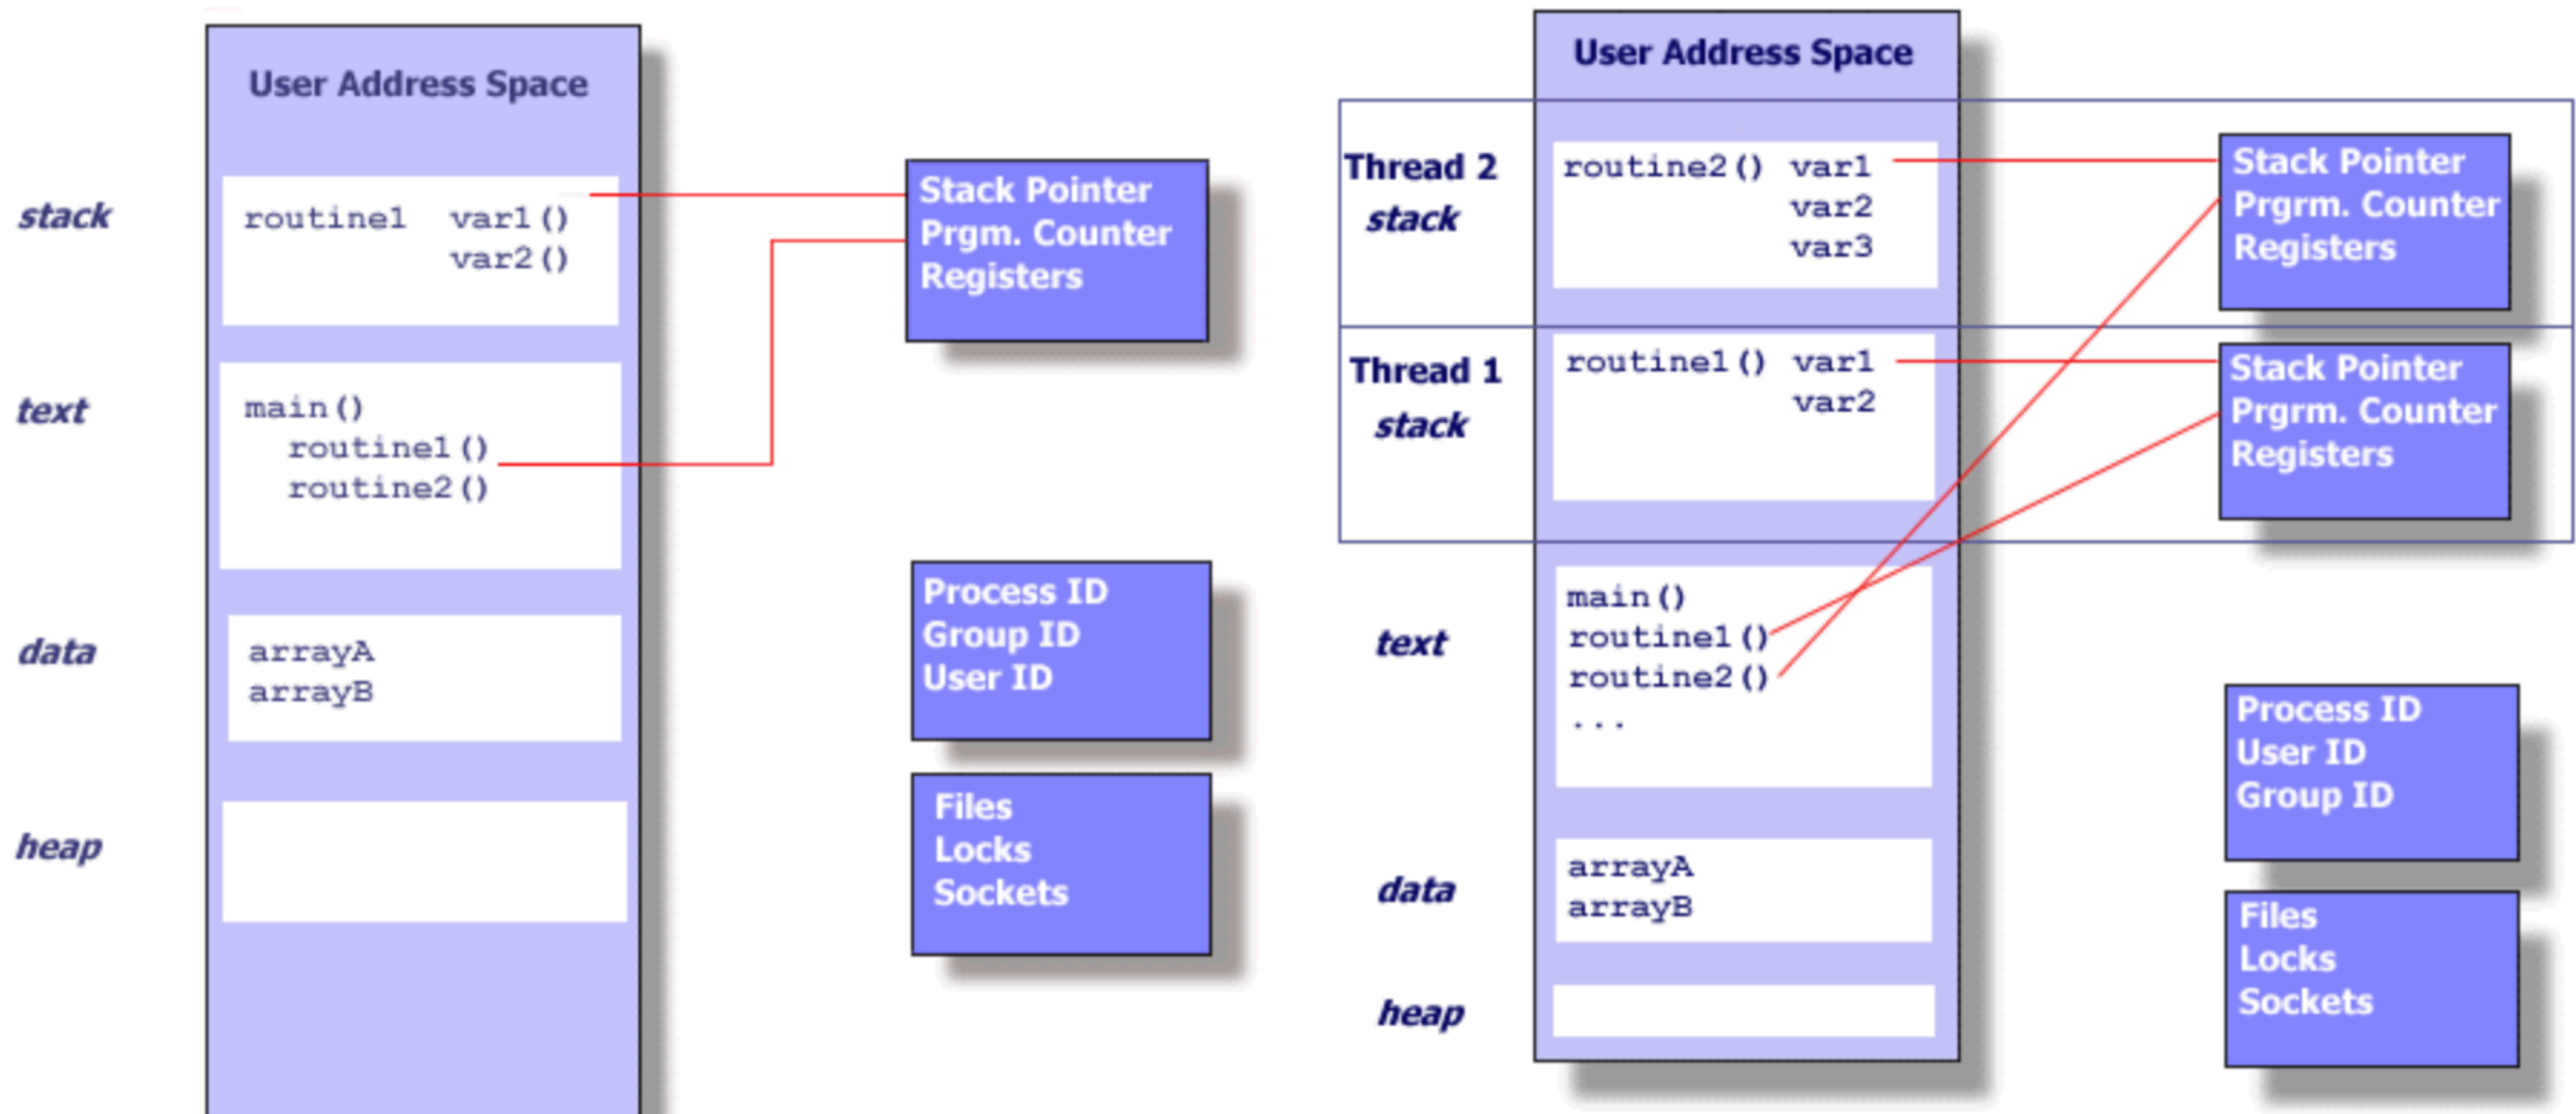
```
Create 4 threads
If thread # == 1:
    Execute Task 1
If thread # == 2:
    Execute Task 2
If thread # == 3:
    Execute Task 3
If thread # == 4:
    Execute Task 4
end if
```

Scope Rules

- Each thread has both shared and private variables.
- Shared global variables
- Private copies of local variables
- Need for mutex(if writing global variables).
- In C/C++:
 - Global: outside all functions.
 - Local: Inside a function or block(`{...}`)



Memory Management



Labs related Info

- Before submitting code, please test on Openlab.
- Compile using flags suggested in the assignment.
- Submit on Canvas
- Only submit essential files (one file per problem)

Login

- ssh to <username>@openlab.ics.uci.edu
- From inside the campus (or housing)
- From outside use the VPN
 - <http://www.lib.uci.edu/connect-campus>
- You will be assigned to any machine that is available

Modules

- Openlab uses the SGE module system
- When you login you only have the basic things
- If you need something you must load the module and it will be available
- More info
<https://www.ics.uci.edu/computing/linux/modules.php>

Using modules

- See what modules are available

```
$ module avail
```

- Load a module

```
$ module load <name>  
Use: gcc/8.2.0
```

C++ threads

- Exist since the C++11 standard came out
- Class that represents a thread of execution
- Is part of the Thread support library
- Underneath the hood uses Pthreads

Creation

- `std::thread` defined in the header `<thread>`
- Receives a function that will be executed in a different thread
- Function must return `void`

Example

```
#include <thread>
void
someFunction()
{
    // Function code ...
}
int
main()
{
    // Some code ...
    std::thread t(someFunction);
    // Main code ....

    return 0;
}
```

Two threads

- You now have two threads
 - Thread `main` executes the rest of `main`
 - Thread `t` executes some function

Just creating threads is not enough!

- The created thread may finish after the main thread
- The operating system will throw an error
- How do we wait for the thread to finish?
- You must wait for the other threads before finishing the main thread

Use `join` to wait for a thread

```
#include <thread>
void
someFunction()
{
    // Function code ...
}
int
main()
{
    // Some code ...
    std::thread t(someFunction);
    // Main code ....

    t.join();

    return 0;
}
```

How to pass arguments

- Arguments are passed as part of the creation process

```
void
someFunction(int a)
{
    // Does something with a
}

int
main()
{
    int mainA;
    std::thread t(someFunction, mainA);
    return 0;
}
```

Passing references doesn't work!

```
void  
someFunction(int& a)  
{  
    // Does something with a  
}  
  
int  
main()  
{  
    int mainA;  
    std::thread t(someFunction, mainA);  
    return 0;  
}
```

You have to use `std::ref`

```
void  
someFunction(int& a)  
{  
    // Does something with a  
}  
  
int  
main()  
{  
    int mainA;  
    std::thread t(someFunction, std::ref(mainA));  
    return 0;  
}
```

Vectors of threads

- To create a vector of threads is simple
- There are two equivalent options
 - `push_back(std::thread(funcNam, arg...))`
 - `emplace_back(func_name, args...`
- Remember the threads will start as soon as you create them!

Example of vector push_back

```
void
someFunction(int a)
{
    // Does something with a
}

int
main()
{
    std::vector<std::thread> threads;
    for(int i = 0; i < 10; ++i)
        threads.push_back(std::thread(someFunction, i));
    for(int i = 0; i < 10; ++i)
        threads[i].join();
    return 0;
}
```

Example of vector emplace back

```
void  
someFunction(int a)  
{  
    // Does something with a  
}  
  
int  
main()  
{  
    std::vector<std::thread> threads;  
    for(int i = 0; i < 10; ++i)  
        threads.emplace_back(someFunction, i);  
    for(int i = 0; i < 10; ++i)  
        threads[i].join();  
    return 0;  
}
```

C++ Atomic

- Each instantiation and full specialization of the `std::atomic` template defines an atomic type.
- If one thread writes to an atomic object while another thread reads from it, the behavior is well-defined.
- Use hardware support to perform operations.
- Check out links in Lab Document for more details

Atomic Compare and Exchange/Swap

- Compare-And-Swap (CAS) used to achieve synchronization.
- C++11 supports this atomic operation on language level to help us write portable multithreaded code that run on all platforms that are standard-compliant.
- It loads the value from memory, compares it with an expected value, and if equal, store a predefined desired value to that memory location.
- All these are performed in an atomic way, i.e., if the value has been changed in the meanwhile by another thread, CAS will fail.
- On x86, performed using a single *cmpxchg* instruction.

Atomic Compare and Exchange/Swap

```
bool compare_exchange_weak( T& expected, T desired,  
                           std::memory_order order =  
                           std::memory_order_seq_cst ) noexcept;  
  
bool compare_exchange_strong( T& expected, T desired,  
                             std::memory_order order =  
                             std::memory_order_seq_cst ) noexcept;
```

Parameters

expected - reference to the value expected to be found in the atomic object.

Gets stored with the actual value of *this if the comparison fails. (NOTE!)

desired - the value to store in the atomic object if it is as expected

order - the memory synchronization ordering for both operations

Return value

true if the underlying atomic value was successfully changed, **false** otherwise.

Atomic Compare and Exchange/Swap

- For the dining philosopher problem, we use to get hold of the forks.
- Use ***compare_exchange_weak*** to occupy the forks one at a time. (Strong version is fine too, for x86 its the same.)

```
bool first_fork = false;  
if( (phil.fork).compare_exchange_weak(first_fork, true) ) {.....}
```
- Use it inside a ***while*** loop, since it will fail a lot.
- Two forks = Two CAS, but just one while loop.
 - Remember to drop the first fork, if thread can't get the second fork.

Mutual exclusion

- To enable serial access to data/streams we have to use mutual exclusion
- We will see three ways:
 - mutex
 - lock_guard
 - unique_lock (we won't use it)

Mutex

- A primitive object to ensure mutual exclusion
- Basic usage

```
std::mutex m;  
int sharedData;  
  
while(!m.try_lock())  
{  
    // manipulate shared data  
    sharedData++;  
    m.unlock();  
}
```

Problems

- What happens if you forget to unlock?
- What happens if there is an exception in the critical section?
- Not recommended to use it in this direct manner

lock_guard

- A better solution than mutex
- An object that holds the mutex and locks it
- Will hold the lock while it is in the scope
- When it is destroyed it releases the lock
- No way to hold the lock forever

lock_guard usage

```
std::mutex m;  
int sharedData;  
  
{  
    std::lock_guard<std::mutex> lock(m)  
    // manipulate shared data  
    sharedData++;  
} // mutex is released here
```


Compiling C++11 Threads

- Load module
- Use the ‘-std=c++11’ and ‘-pthread’ flags
- Add ‘-Wall’ to see warnings




```
$ g++ -std=c++11 -pthread example.cpp -o Example
```

Comments

- We need the shared variables (nextAvailableChunk, outputImage)
- It is better to separate into a function where you are going to update it
 - Assure access is mutually exclusive
- Maybe reduce conflict while writing the image
 - Create temporary partial output image.
 - Reduce chances of Race condition, if error in implementation

Should look like this!

- Solution files are provided for reference.
- Use Linux ***diff*** to check if your solution is correct.
- Usual mistakes:
 - Missing rows around chunks



Implementation

- If you want to change something in skeleton?
 - You can. But do not change CL args structure.
- Seeing some extra variables or structs or globals?
 - Ignore them if you don't need them.

Grading

- 50 pts for each implementation
- 0 pts, if program doesn't even compile.
 - Before submission, check if the code compiles and works on **openlab** server.
- 0 pts, if it just compiles but does nothing useful at all.
- Partial grading, only if there is a minor bug/design error.

Questions?

You can post on Piazza
Or
Attend Office hours

Discussion Slides will be posted on Class Website