


Discussion 5

Basic OpenMP



Pthreads vs OpenMP

- | | |
|---|---|
| <ul style="list-style-type: none">• Low-level model• Not really portable<ul style="list-style-type: none">• Different versions/library for each language• More programming effort• Fine grained control (ex: for Data) | <ul style="list-style-type: none">• High-level API• Good for loop-level parallelism, and even more now• Multi-platform/Portable<ul style="list-style-type: none">• Supports C/C++, Fortran• Less programming effort for similar speedup• Easier work sharing (ex: <i>omp</i> for inside <i>parallel</i> region)• Can be used in conjunction with MPI |
|---|---|

Essential OpenMP

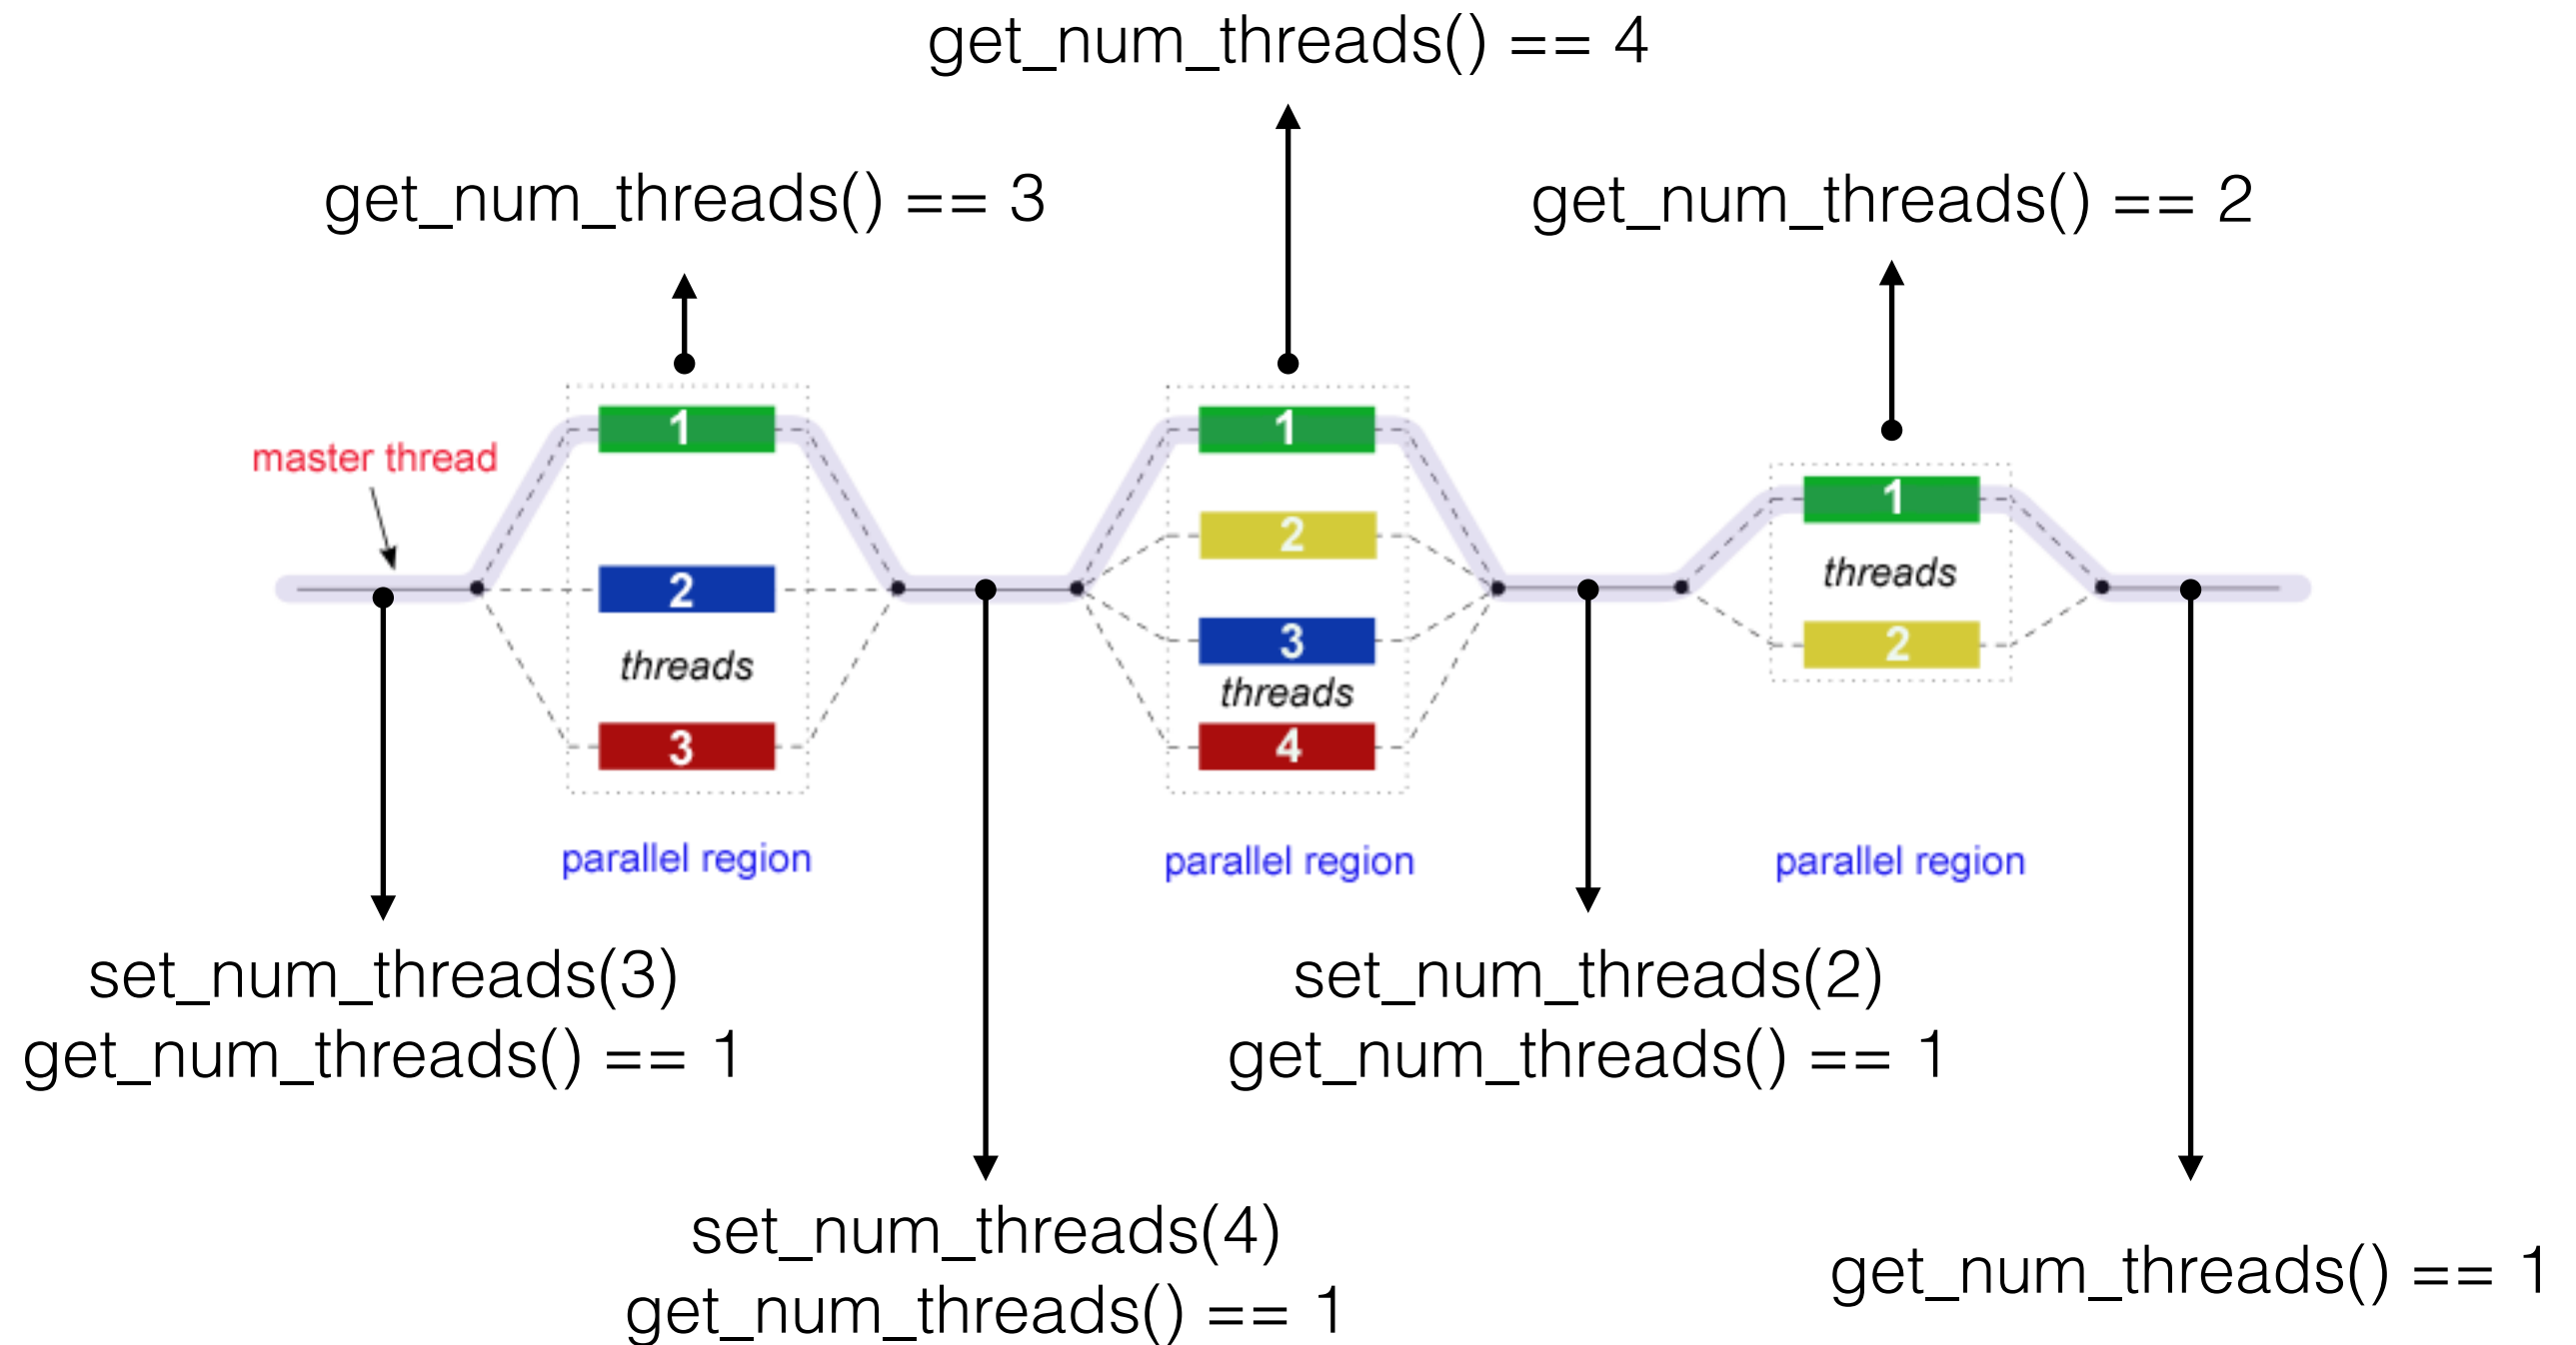
Directives and Functions

- `omp_get_num_threads()`
- `omp_get_thread_num()`
- `omp_set_num_threads(numThreads)` - *Don't use this for the assignment*
- `#pragma omp parallel`
- `#pragma omp parallel for`
 - Static
 - Dynamic
- `#pragma omp critical`
- `#pragma omp task`
- `#pragma omp taskwait`

Set/Get Num Threads

- Set environment variable using CL (Recommended):
 - export/set **OMP_NUM_THREADS**=<number of threads to use>
- **omp_get_num_threads()**:
 - Returns the number of threads executing in the current region (parallel or not)
 - If not in a parallel region it will return 1 (i.e. only the master thread is running)
- **omp_set_num_threads**(numThreads)
 - Overrides the above env variable.
 - Declared outside OMP clauses.
 - Sets the number of threads
 - numThreads must be an int larger than 0


Threads



omp_get_thread_num()

- Returns an Id for the thread that is executing
- Id is unique and always the same in parallel region
- It is between 0 and (omp_get_num_threads() - 1)
- Master Thread always has Id 0

Thread Id's



#pragma omp parallel

- Defines a parallel region whose code will be executed by multiple threads in parallel

```
#pragma omp parallel [clauses]
{
    code_block
}
```

- Implicit barrier/join at the end of parallel region
- Clauses are all optional, all have default values

#pragma omp parallel Clauses

- `num_threads` —> how many threads execute parallel region
 - If not defined it is the default value.
 - Default value is:
 - the one set by `set_num_threads` or
 - the environment variable `$OMP_NUM_THREAD`
- Order of clauses doesn't matter
- Several other clauses.
- See <https://msdn.microsoft.com/en-us/library/68ah4xc7.aspx>

#pragma omp parallel Clauses

- **private** —> list of variables that are private to the threads inside of the parallel region
 - Variables are copied into a local version for each thread
 - Changes to these variable are not shared between threads
 - Changes to these variable will not be reflected after the parallel region
 - If a variable is not private it is shared between all of the threads in the parallel region and changes to it will be reflected outside of the region
 - There are other similar options: **firstPrivate** (initialize with a value), also lastPrivate
 - Good example to understand clauses:
 - <https://msdn.microsoft.com/en-us/library/c3dabskb.aspx>
 - Default clause (Shared or Private): https://www.ibm.com/support/knowledgecenter/en/SSLTBW_2.2.0/com.ibm.zos.v2r2.cbcp01/cupppvars.htm

Example

```
#pragma omp parallel num_threads(4)
{
    int i = omp_get_thread_num();
    printf_s("Hello from thread %d\n", i);
}
```

=

```
set_num_threads(4);
#pragma omp parallel
{
    int i = omp_get_thread_num();
    printf_s("Hello from thread %d\n", i);
}
```

```
Hello from thread 0
Hello from thread 1
Hello from thread 2
Hello from thread 3
```

- Notice that thread id's go from 0 to 3
- Order of print statements may change

#pragma omp parallel for

- Defines a parallel for whose iterations will be divided among several threads

```
#pragma omp [parallel] for [clauses]
{
    code_block
}
```

- If in a parallel region then don't put the parallel...

```
#pragma omp parallel for [clauses]
{
    code_block
}
```

=

```
#pragma omp parallel [clauses]
#pragma omp for [clauses]
{
    code_block
}
```

OMP Parallel For

General Syntax:



```
#pragma omp parallel
#pragma omp for
for (i=0; i<N; i++) {
    // do something with i
}
```

Or

```
#pragma omp parallel for
for (i=0; .....)
```

Reason you may want a separate parallel code segment:

```
#pragma omp parallel
{
    code1();
    #pragma omp for
    for (i=1; i<=4*N; i++) {
        code2();
    }
    code3();
}
```



OpenMP *for* scheduling

```
#pragma omp parallel for schedule(kind [,chunkSize])
```

- A parallel for divides the iterations between the executing threads
- How is the dividing done?
 - It depends on the schedule kind clause(static, dynamic, guided)
- There are several kinds of scheduling, each has its advantages and disadvantages
 - We will use: Static and Dynamic
- See <https://msdn.microsoft.com/en-us/library/x5aw0hdf.aspx>

Static and Dynamic Scheduling

- **Static** scheduling divides iterations evenly between threads
 - Default chunk size: $num_Iterations/num_Threads$
- **Dynamic** scheduling uses an internal work queue to dispatch chunkSized blocks to each thread. When a thread finishes a block, it gets the next block
 - Default chunk size is 1
 - Think Part B Lab 1

Static vs Dynamic Scheduling

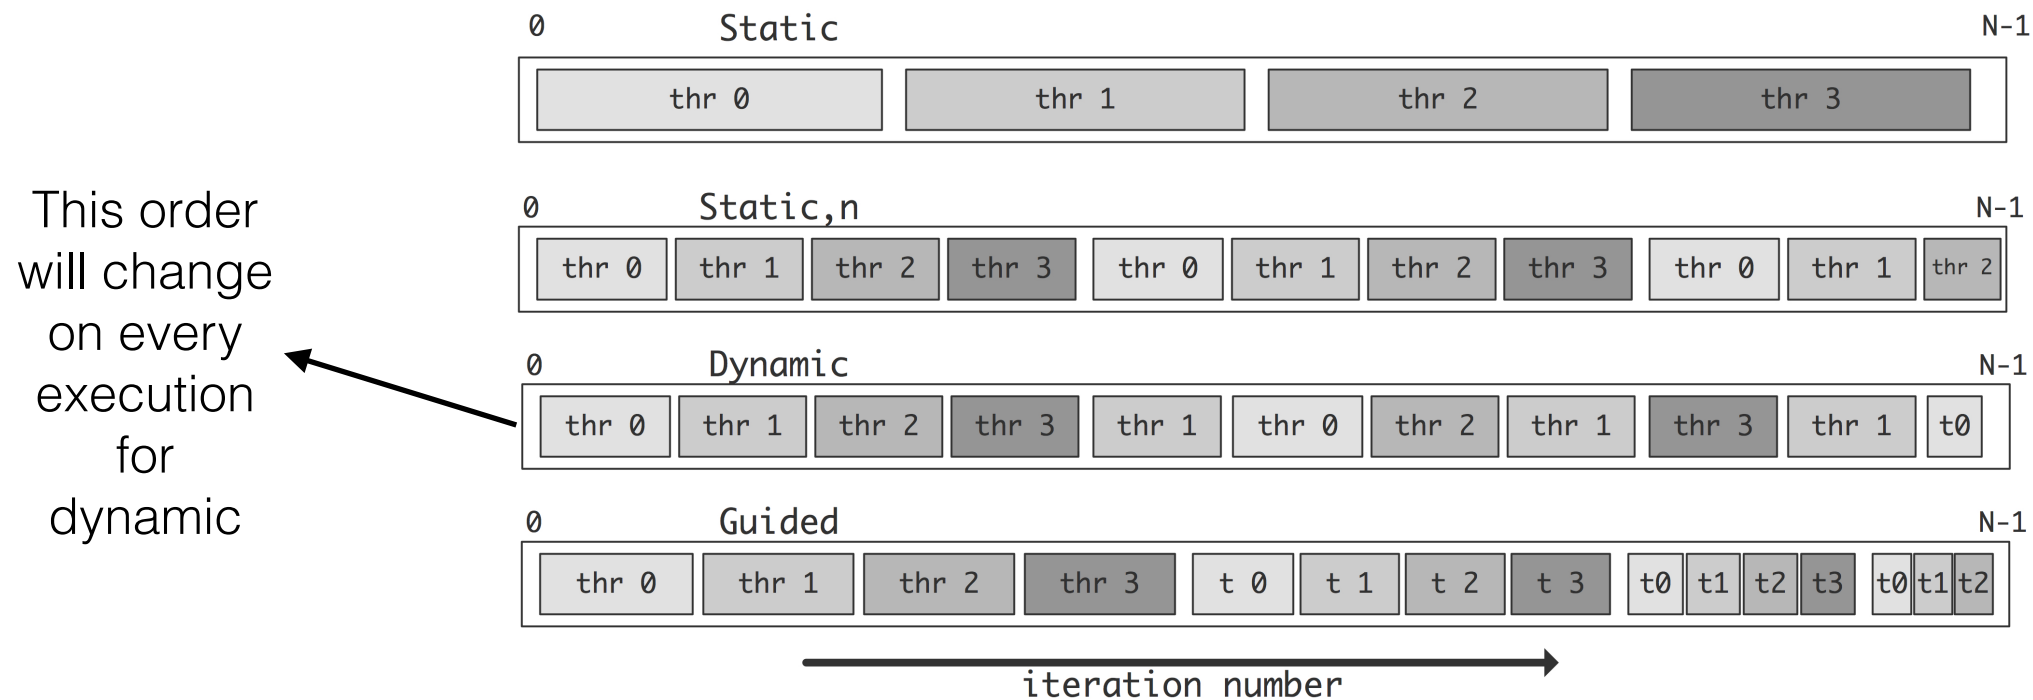



Illustration static round-robin scheduling versus dynamic



#pragma omp task

- When a thread encounters a task construct, it may choose to execute the task immediately or defer its execution until a later time.
- If deferred, the task is placed in a conceptual pool of tasks associated with the current parallel region.
- All team threads will take tasks out of the pool and execute them until the pool is empty.
- A thread that executes a task might be different from the thread that originally encountered it.

#pragma omp task

Clauses:

- depend (list)
- if (expression)
- final (expression)
- untied
- mergeable
- default (shared | firstprivate | none)
- private (list)
- firstprivate (list)
- shared (list)
- priority (value)

```
int main(int argc, char *argv[]) {  
    #pragma omp parallel  
    {  
        #pragma omp single  
        {  
            printf("A ");  
            #pragma omp task {printf("race ");}  
            #pragma omp task {printf("car ");}  
        }  
    } //End omp parallel - taskwait not needed  
    printf("\n");  
    return(0);  
}
```

OUTPUT: A race car / A car race

#pragma omp task

```
float sum(const float *a, size_t n)
{
    // base cases
    if (n == 0) {
        return 0;
    }
    else if (n == 1) {
        return *a;
    }

    // recursive case
    size_t half = n / 2;
    return sum(a, half) + sum(a + half, n - half);
}
```

```
float sum(const float *a, size_t n)
{
    // base cases
    if (n == 0) {
        return 0;
    }
    else if (n == 1) {
        return *a;
    }

    // recursive case
    size_t half = n / 2;
    float x, y;

    #pragma omp parallel
    #pragma omp single nowait
    {
        #pragma omp task shared(x)
        x = sum(a, half);
        #pragma omp task shared(y)
        y = sum(a + half, n - half);
        #pragma omp taskwait
        x += y;
    }
    return x;
}
```

#pragma omp task

```
float sum(const float *a, size_t n)
{
    // base cases
    if (n == 0) {
        return 0;
    }
    else if (n == 1) {
        return 1;
    }

    // recursive case
    size_t half = n / 2;
    float x, y;

    #pragma omp parallel
    #pragma omp single nowait
    {
        #pragma omp task shared(x)
        x = sum(a, half);
        #pragma omp task shared(y)
        y = sum(a + half, n - half);
        #pragma omp taskwait
        x += y;
    }
    return x;
}
```

```
#define CUTOFF 100 // arbitrary

static float parallel_sum(const float *, size_t);
static float serial_sum(const float *, size_t);

float sum(const float *a, size_t n)
{
    float r;

    #pragma omp parallel
    #pragma omp single nowait
    r = parallel_sum(a, n);
    return r;
}

static float parallel_sum(const float *a, size_t n)
{
    // base case
    if (n <= CUTOFF) {
        return serial_sum(a, n);
    }

    // recursive case
    float x, y;
    size_t half = n / 2;

    #pragma omp task shared(x)
    x = parallel_sum(a, half);
    #pragma omp task shared(y)
    y = parallel_sum(a + half, n - half);
    #pragma omp taskwait
    x += y;

    return x;
}

static float serial_sum(const float *a, size_t n)
{
    // base cases
    if (n == 0) {
        return 0.;
    }
    else if (n == 1) {
        return a[0];
    }

    // recursive case
    size_t half = n / 2;
    return serial_sum(a, half) + serial_sum(a + half, n - half);
}
```

#pragma omp critical

- Defines a critical region where only one thread can execute at a time. Good for output and debugging

```
#pragma omp critical [(name)]  
{  
    // code_block  
}
```

- (name) is used to identify the region, two critical regions with the same name are considered the same region
- (name) must be between parentheses
- See <https://msdn.microsoft.com/en-us/library/b38674ky.aspx>
- Similar option for single assignment statement: #pragma omp atomic

#pragma omp atomic

- Similar to *critical*, this directive allows access of a specific memory location atomically.
- Advantage: Fewer locks/mutexes, limited but more efficient (Similar to Compare-and-Exchange/Swap)
- Atomic Clauses:
 - Update (default) -> x++, x--, ++x, --x, binary ops (+, *, -, /, &, ^, |, <<, >>)
 - Read -> y = x
 - Write -> x = y
 - Capture -> y = x++, x--, ++x, --x, binary ops

```
#pragma omp atomic [atomic-clause]
{
    // atomic operation
}
```

```
extern int x[10];
extern int f(int);
int temp[10], i;

for(i = 0; i < 10; i++)
{
    #pragma omp atomic read
    temp[i] = x[f(i)];

    #pragma omp atomic write
    x[i] = temp[i]*2;

    #pragma omp atomic update
    x[i]++;
}
```

- See https://www.ibm.com/docs/en/zos/2.2.0?topic=SSLTBW_2.2.0/com.ibm.zos.v2r2.cbclx01/prag_omp_atomic.htm
- Also: <https://www.openmp.org/spec-html/5.0/openmps95.html>

Important things!

- Be clear on what variables needs to be PRIVATE, FirstPrivate and SHARED
- Make sure you print Threads and the starting of their chunks (after output image is created).
- Reference for OMP Tasks:
 - <http://pages.tacc.utexas.edu/~eijkhout/pcse/html/omp-task.html>