

**Welcome to**

**CS131**

# **Parallel and Distributed Systems**

**Prof. A. Veidenbaum**

# Today's topics

- Distributed systems design goals
- Types of distributed systems
- Reading assignments:
  - Today: Ch. 1, Sec. 1.2
  - Next time: Ch. 1, Sec. 1.3 - 1.3
  - Plz complete reading assignments before next class

# Summary of last lecture

- **A distributed system is a collection of independent computers that appears to its users as a single coherent system.**
- **Middleware sftw is what makes it a DS**
  - **Somewhat like an OS for a single computer**
  - **Middleware provides**
    - » **Resource management**
    - » **Facilities for inter-application communication.**
    - » **Access management**
    - » **Security services.**
    - » **Masking of and recovery from failures.**

# Major DS Design Goals

- **Allow users to easily connect to resources**
- **Transparency**
  - hide the fact that resources are distributed
- **Openness**
  - allow extendibility, interoperability
- **Scalability**
  - An ability to increase system size without problems
- **Security, Fault-tolerance, etc**

# Why Connect Users/Resources

- **Distributed systems allow resource sharing**
  - printers, cycle servers, files, documents, storage devices, Web pages, networks, etc
- **Makes sense on many levels**
  - economic reasons
  - ease of information exchange and sharing
  - collaboration
- **Creates new problems**
  - Access management, consistency
  - Security
  - Fault tolerance
  - Fair and controlled resource sharing

# Transparency

- **Makes a DS look like a single computer**
- **Users do not need to know where the resources are**
  - Or that they are shared
  - Or that they fail
  - Or that they migrate or are replaced
- **How is it achieved?**
  - Through multiple layers of software!

# Transparency (Cont'd)

- **Can be applied at multiple levels in a DS**
  - Access
  - Location
  - Relocation
  - Migration
  - Replication
  - Concurrency (of use)
- **Is complete transparency always desired?**
  - Hint: think of performance

# Openness

- Allows components from different vendors
- Allows new components to be added
- Allows old components to be replaced
- Implies interoperability and portability
- **Uses an Interface Definition Language (IDL)**
  - specifies syntax for a process to talk to another process which provides a given interface
    - » E.g. function names, parameters, etc
  - NOT the implementation (policy vs mechanism)



# Openness (Cont'd)

- **Use of an IDL allows**
  - Different implementations of a given interface
  - Many to one mapping for a service
- **A DS needs an ability to change**
  - For instance, change/modify/improve a component
- **Separating policy from mechanism allows this**
  - Book example – a Web cache with a number of policies
    - » Where is the data cached?
    - » What is the replacement policy?
    - » Is it shared or private?
    - » Consistency with the source

# Scalability

- An ability to grow the system
  - One of the key advantages of distributed systems
    - » Adapt to new requirements, such as more users
- Centralization of resources limits scalability
- Replication and Caching improves scalability
  - BUT they lead to **consistency** problems
    - » Multiple copies exist *and* may be different
      - Have you encountered this problem?

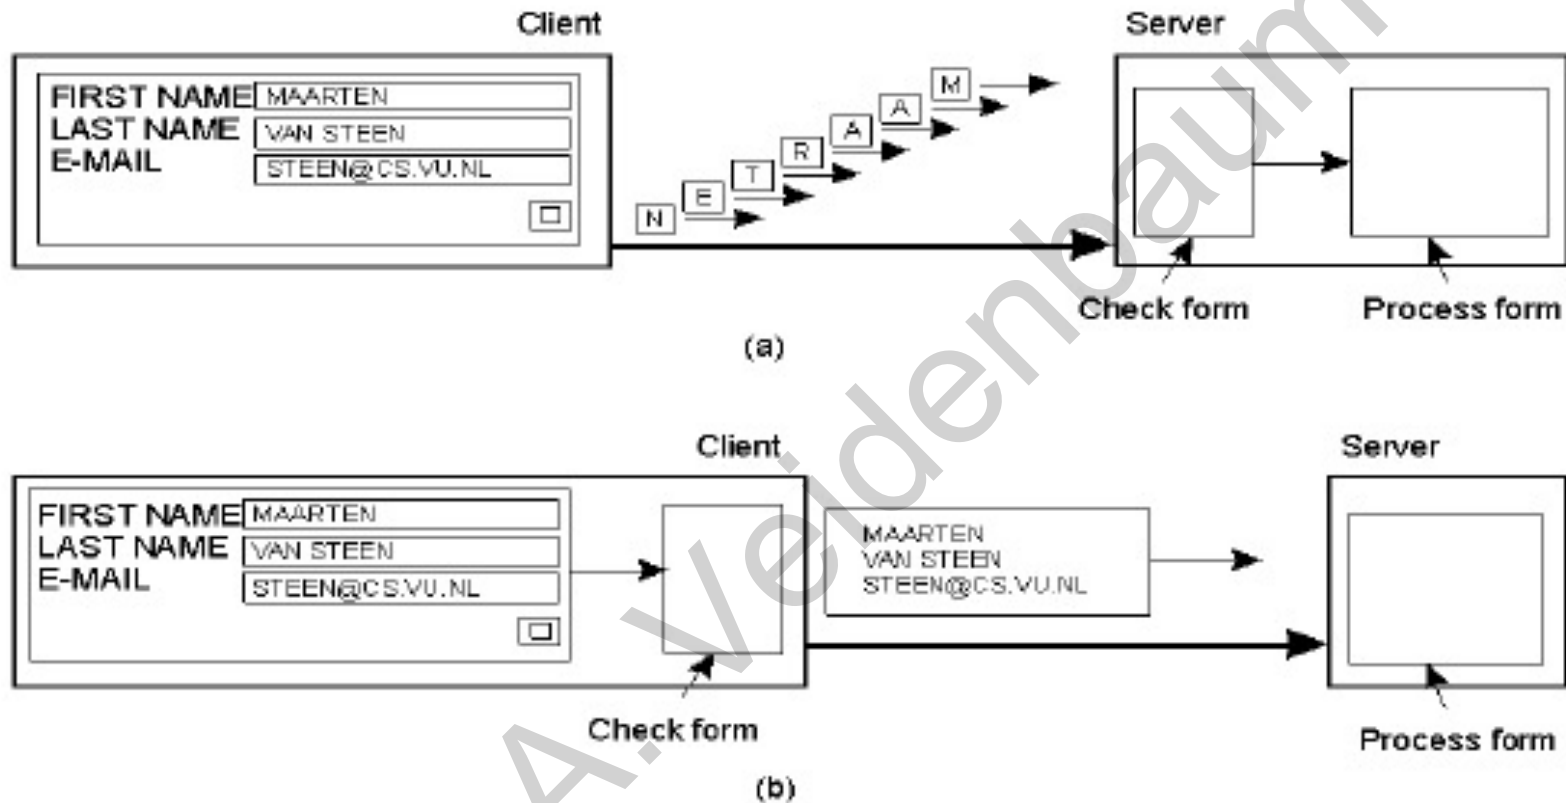
# Scalability

- Can be measured in 3 different dimensions
  - **Size scalability**
    - » can more users/resources be added
  - **geographical scalability**
    - » with respect to the distance
  - **administrative scalability**
    - » across multiple administrative domains

# Examples of Scalability Problems

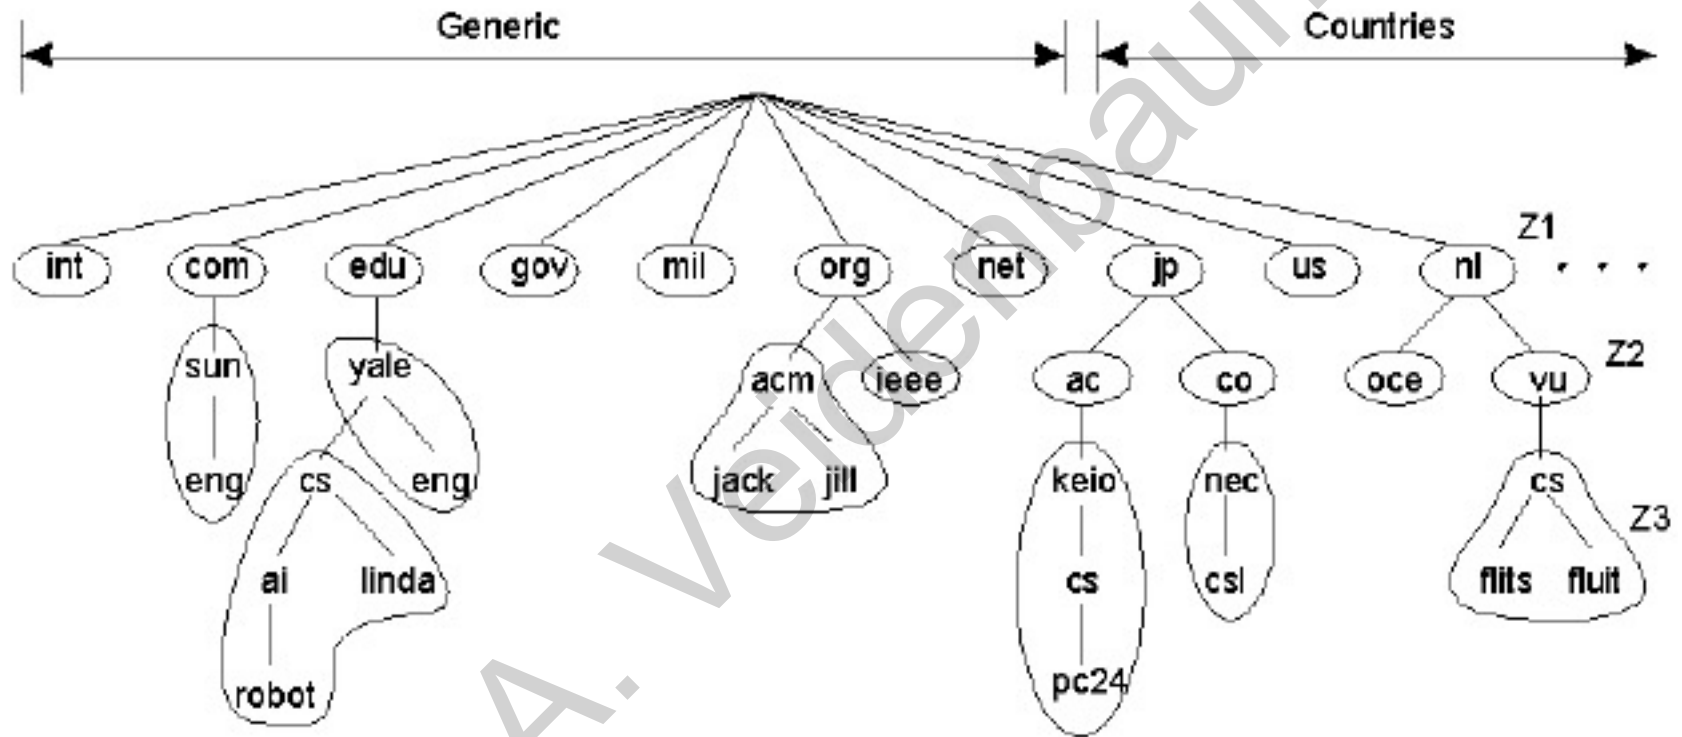
| • Concept                       | • Example  |
|---------------------------------|--|
| • <b>Centralized services</b>   | • <b>A single server for all users</b>               |
| • <b>Centralized data</b>       | • <b>A single on-line telephone book</b>             |
| • <b>Centralized algorithms</b> | • <b>Doing routing based on complete information</b> |

# Scaling Techniques (1) - hiding latency



- Who does the work of checking forms as they are being filled?
- a) is not efficiently using the network and the server

# Scaling Techniques (2) – partitioning, distribution



Distribution of service: dividing DNS name space into zones.

# Scaling Techniques (3) – replication, caching

- Replication helps limit performance problems
  - Local servers
  - Multiple resources - Google data centers
- Caching is a special form of replication
  - Making a copy of data
- Replication/caching lead to **consistency** problems
  - Multiple copies may differ
    - » Think of browser (local) caching

# Possible implementations

Consider a large database with many users. How to best implement it?

1. One central database
  - What will this lead to?
2. Multiple database partitions distributed somehow
  - Better?
3. Each computer has a full database
  - Problems?
4. Replicate partially and Cache locally



# Pitfalls in designing DSs

- **Basically, dubious simplifying assumptions made during the DS design**
  - The network is reliable
  - The network is secure
  - The network is homogeneous
  - The topology does not change
  - Latency is zero
  - Bandwidth is infinite
- **Such assumptions need to be carefully considered!**