

CompSci 131

Parallel and Distributed Systems

Prof. A. Veidenbaum

Today's topics

- Last lecture review
- Types of Distributed Systems
- Reading assignment:
 - Today's: Sec. 1.3.1-1.3.2
 - Next: 2.1-2.2
 - » Complete the assignment before next class

Summary of last lecture

- **A distributed system is a collection of independent computers that appears to its users as a single coherent system.**
- **Design goals**
 - **Allow users to easily connect to resources**
 - **Hide the fact that resources are distributed (Transparency)**
 - **Openness to allow extendibility, interoperability**
 - **Scalability to allow growth without performance loss**

‘First Time’ DS Design Pitfalls

- **Are standard sftw engineering rules enough?**
- **No, pitfalls include assuming that**
 - Network is reliable
 - Network is secure
 - Network is homogeneous
 - Topology does not change
 - Latency is zero
 - Bandwidth is infinite
 - Transport cost is zero
 - There is only one administrator
- **Most of the DS design principles relate to the above**
 - Solve problems caused by one or more of these being false


Types of Distributed Systems

- **Distributed Computing Systems**
 - Targets increase in application performance
- **Distributed Information Systems**
 - Supports enterprise computing
- **Distributed Pervasive Systems**
 - Collections of small devices, not even always on
 - Adaptable to local environment

Distributed Computing

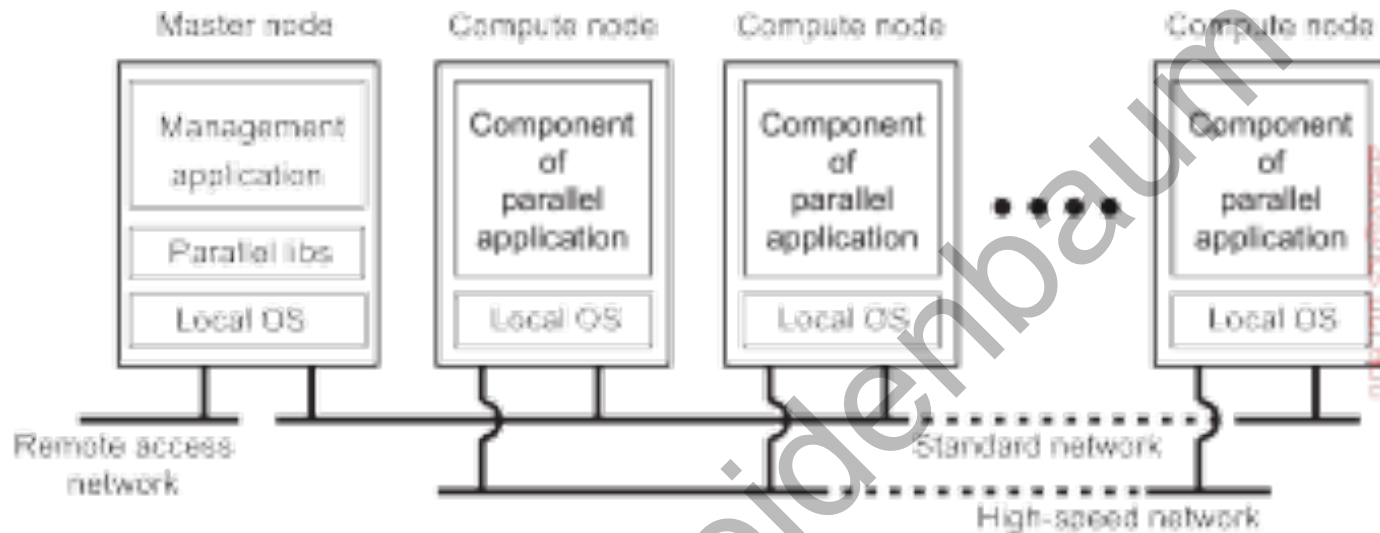
- **Clusters - integrate a small number of local systems**
 - Master node: parallel libraries, application management
 - Multiple computing nodes
 - Middleware layer to tie together
- **Small can mean 128 or even a few K nodes**
 - ICS openlab is an example
- **You will write a program using a middleware layer called MPI**
 - Submit via a master queue using SLURP sftw
- **Can be considered a small version a supercomputer**

Multiprocessor vs multi-computer



- Today high-performance computing nodes are multiprocessors
- Clusters are multi-computers
- Supercomputers are very large clusters with custom interconnect

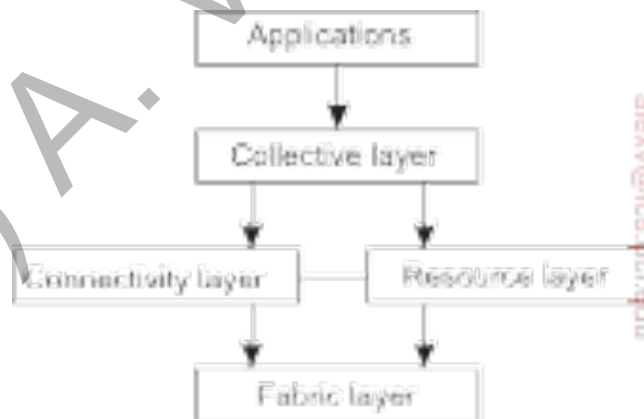
Cluster Computing



- Only master is remotely accessible
 - Job submission, file exchange, etc
- Compute nodes run independent portions of a parallel application

Grid computing

- Integrate resources on the internet
 - Another layered software organization
 - » **Fabric layer:** interface to local resources at a given site
 - » **Communication layer**
 - Authentication, communication
 - » **Resource layer:** allows access to a given resource
 - Access control, configuration data access, specific operations
 - » **Collective layer:** access to multiple resources
 - Resource discovery, allocation, scheduling
 - » **Application layer**



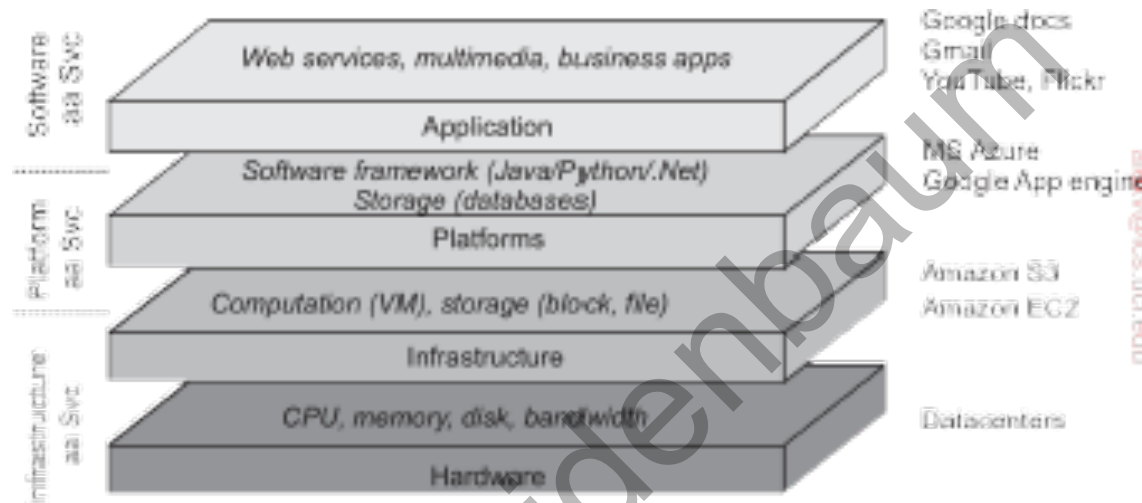
Grid computing

- Integrates resources from different organization
 - Another name – federation
 - Creates a virtual organization for collaboration
- *Pre-cloud*, allowed many more nodes than a cluster
- What is the main difference with cluster?
 - other than size

Cloud computing

- **Grid -> utility computing -> clouds**
 - Another name – federation
 - Creates a virtual organization for collaboration
- ***A cloud is a pool of virtualized resources***
- **Typically four layers**
 - Hardware
 - Infrastructure (employs virtualization)
 - Platform (somewhat similar to OS: job submission, files)
 - Application (Google docs, Acrobat DC, etc)

Cloud computing



- These layers are offered to customers as
 - Infrastructure as a SVC (IaaS)
 - Platform as a SVC (PaaS)
 - Software as a SVC (SaaS)

Distributed Information Systems

- Typically enterprise level systems
 - Often separating processing from data
- Transaction processing systems deal with data

Transaction Processing Systems

- Transactions protect operations on *data*
 - Access/modify many data items in one atomic operation
 - **An All or Nothing operation**
 - » Either all data is accessed/modified or everything is aborted
 - Aborted = no state change occurs
- The origin is in business transactions
 - Negotiation followed by contract
- Distributed system' transactions are similar
 - Negotiate, perform, commit
 - back out on crashes, other problems

Programming Transactions

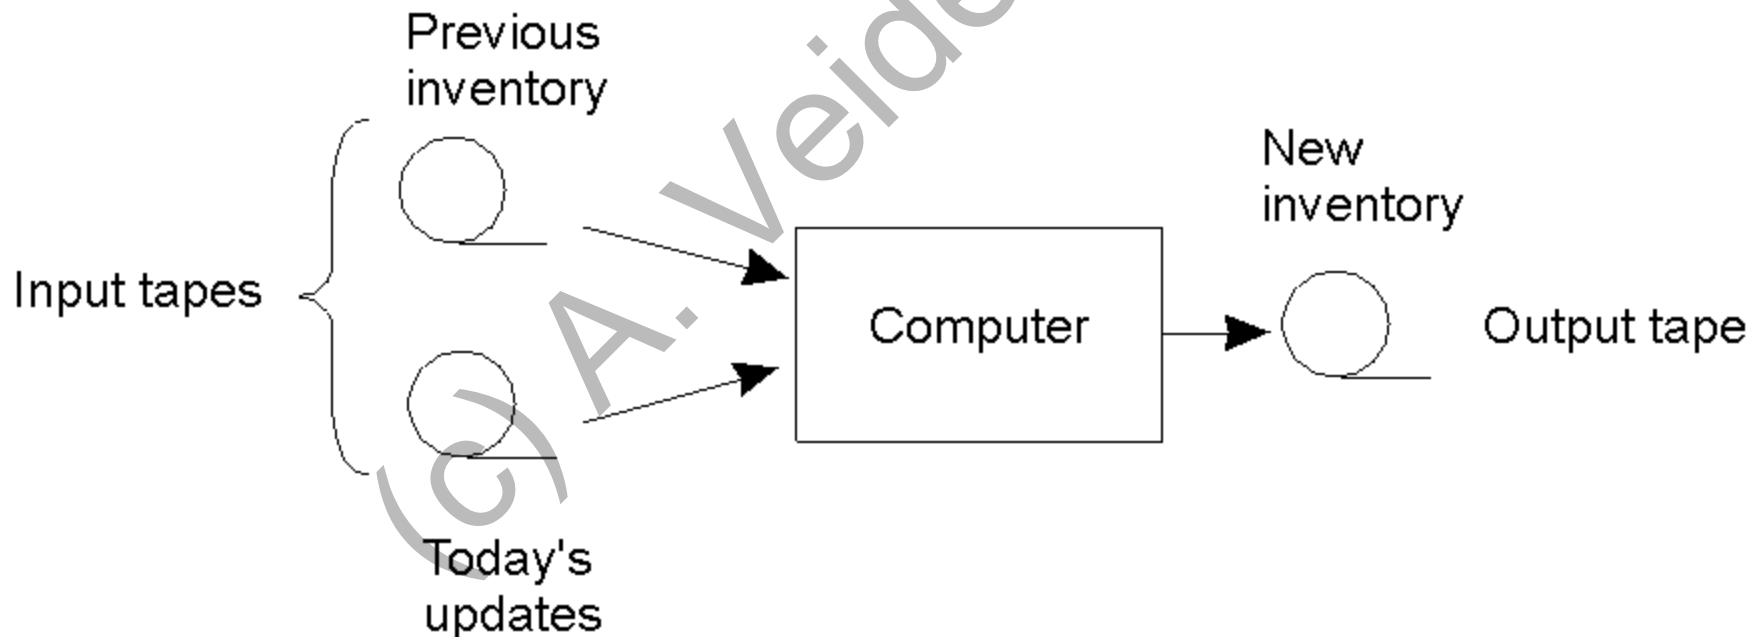
- Basic operations to implement transactions

Primitive	Description
BEGIN_TRANSACTION	Make the start of a transaction
END_TRANSACTION	Terminate the transaction and try to commit
ABORT_TRANSACTION	Kill the transaction and restore the old values
READ	Read data from a file, a table, or otherwise
WRITE	Write data to a file, a table, or otherwise

- Transactions may be nested!

Example from the Old Days

- **Store Inventory computation**
 - Input: tapes of last inventory, today's sales+deliveries
- **Computation can be aborted at any time**
 - Rewind tapes, start over



Today's example

- **Online banking**
 - Withdraw \$1000 from account A
 - Deposit \$1000 into account B
- **What is the problem here?**
- **Can this be done using mutual exclusion?**
- **Transaction model**
 - Perform both or abort

Properties of Transactions

1. Atomic

- All or nothing, looks like everything happens at once

2. Consistent

- Maintains system invariants
 - During a transaction may not be consistent
 - No one can observe that!

3. Isolated or serializable

- Concurrent transactions do not affect each other
- Result is as if they ran sequentially in some order

4. Durable



- After commit the results ARE made permanent

Transaction Models

- **Flat** - a series of operations run to completion
 - Or aborted
 - Has limitations
- **Nested**
 - Top-level transaction can fork off child processes
 - Children can execute their own transactions
 - » Commit of a child transaction can be undone by parent
- **Distributed**
 - A flat transaction on distributed data
 - Requires distributed locking of data, commit

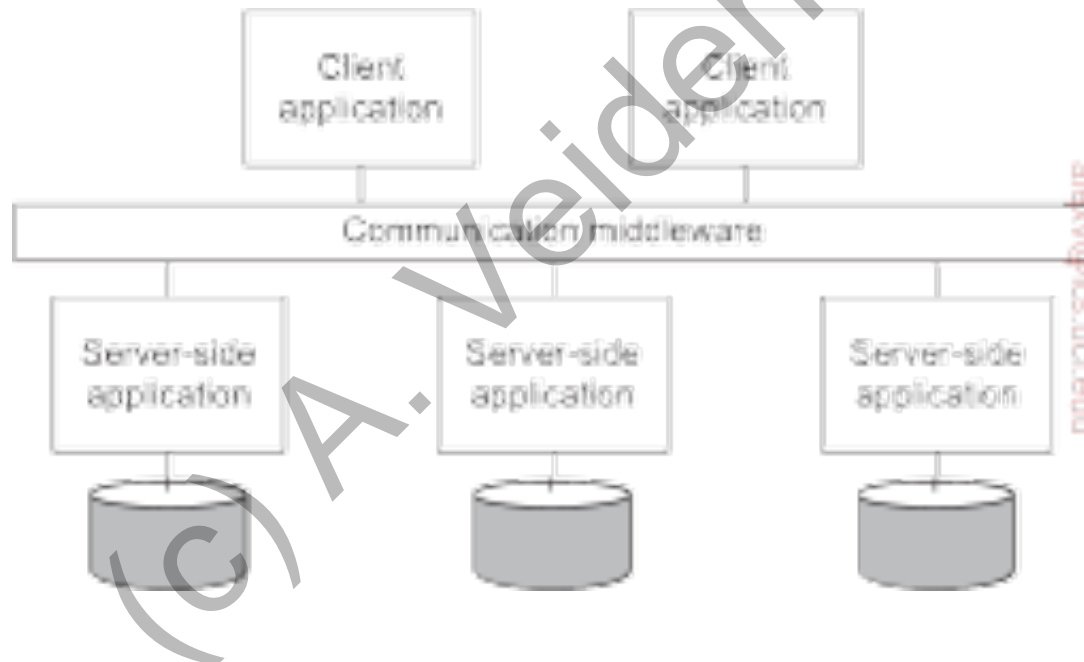
Examples

- a) **Nested transactions - logically separate**
- b) **Distributed transactions - on separate data**




Enterprise application integration

- Application components can communicate directly:
 - Need communication middleware
 - » Remote procedure call (RPC)
 - » Remote method invocation (RMI) for objects



Pervasive Systems

- A pervasive system blends into our environment
 - » Mobile
 - » Sensor networks
- Nodes may move or not be up all the time
- A sensor network example



- All data processed by operator
 - » Sensors may form an ad-hoc network