

CompSci 131

Parallel and Distributed Systems

Prof. A. Veidenbaum

Today's topics

- Last Lecture Summary
- Middleware organization
- System Architecture
- Reading assignment:
 - Today's: Sec. 2.2-2.3
 - Next time: Sec. 2.4 (on your own), 3.1
 - » Also start reading this pthreads tutorial
 - » <https://computing.llnl.gov/tutorials/pthreads/>
 - Complete the assignment before next class

Review of last lecture

- **Architectural styles in DS**
 - Layered
 - Object-based
 - RESTful
 - Event-based
 - Publish/subscribe

Middleware organization


- It is distinct and separate from *software* architecture
- There are two important types of *design patterns* used in middleware organization:
 - Wrappers or adapters
 - Interceptors
- Both target the same goal – DS openness
- Arguably, the ultimate openness is achieved when the middleware can be composed at runtime
 - Modifiable middleware

Wrappers/Adapters

- Solve the problem of expanding a component interface
 - The old component and its interface may remain unchanged for legacy or other reasons
 - A new interface provided by a wrapper may be translated to the old component interface
- Have their origin in OO programming
 - More complex in DS
- Extensibility problem – unique application to application interfaces were used to achieve it
 - N applications would require $O(N^2)$ unique wrappers
- Special middleware may be used for such a reduction
 - For instance use a *broker*, such a message broker

Wrappers/Adapters and Brokers

- A broker knows how to handle any message and sends it to an appropriate component
 - Uses 1-to-N and N-to-1 interfaces, e.g. $O(N)$

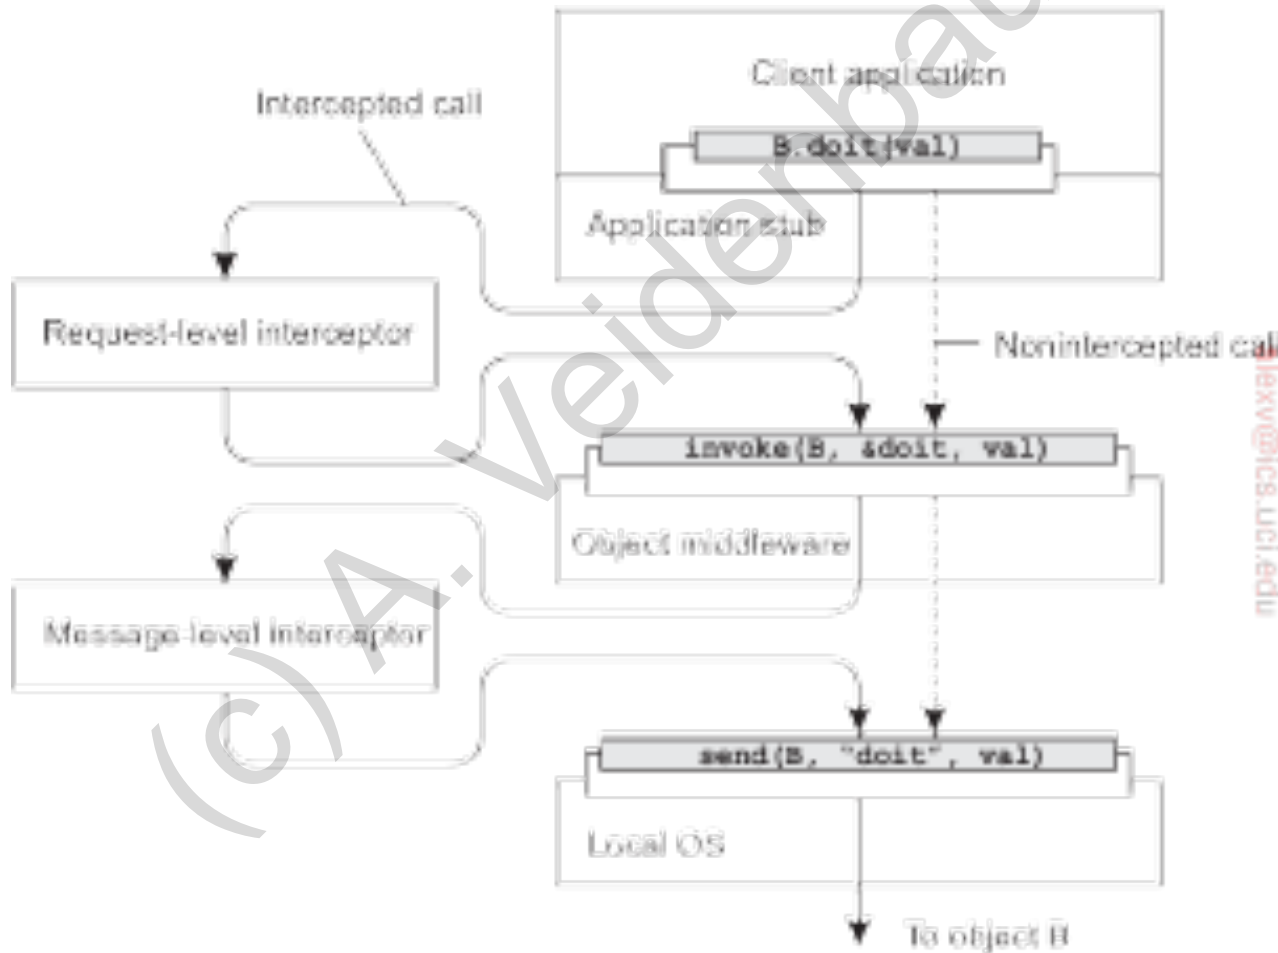


Interceptors

- **A *software* construct that can break normal control flow and allow an app-specific code to be executed**
 - E.g. it adapts middleware to the needs of an application
 - This makes middleware more open
- **Let us consider an object-based DS in which an object A can call a method in object B**
 - The objects reside in different nodes
- **This is done as follows**
 1. **Obj. A is offered a local interface equivalent to B's**
 2. **The call by A is transformed into a generic obj. invocation**
 - » **The general obj. invocation interface is offered locally by the middleware**
 3. **The generic obj. invocation is transformed into a message and sent to B**

Interceptors

- The example shows two interceptors
 1. A message interceptor we discussed on previous slide
 2. A request interceptor that can be used for replicated obj.



Modifiable middleware

- Wrappers and interceptors are a means to adapt to continuous changes in DSs.
 - Failures, delays, replication, system scaling, etc.
- Middleware is made to adapt to such changes
 - Applications are blissfully unaware of any
- Scalability, in particular, makes it very desirable to modify middleware without bringing the DS down
 - Adaptive or modifiable middleware does this
 - » Replacing a component at runtime is a good example
- One approach is to *dynamically construct* middleware from components
 - As opposed to statically configuring it at design time

Modifiable middleware

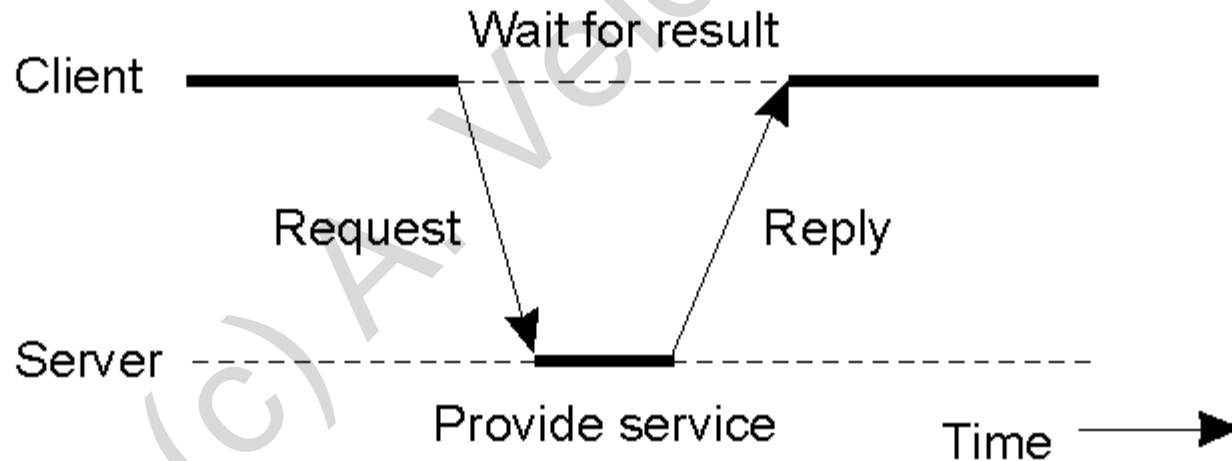
- **Dynamic configuration is complex**
 - Need support for late binding
 - Need to know the effect of component replacement on all other affected components
- **A minimum requirement is to provide dynamic loading/unloading of components**
- **This is complex and is the subject of current research**

System Architectures

- Recall that *A system* architecture defines software instantiated on a real system
 - After final configuration choices were made
 - Component placement is decided
- There are two main types
 - Centralized
 - » Client-server, with clients sending requests
 - Distributed
 - » Peer-to-Peer

Client/Server Architectures

- A good model for interaction in a distributed systems
 - Servers provide a service abstraction
 - Clients request services
- Here is a typical interaction between a client and a server
 - Or a request-reply behavior



- **Can use a simple connectionless communication**
 - **If the network is reliable or if message loss acceptable**
- **Or use TCP/IP otherwise**
 - **But it is less efficient**
- **Server will interpret and do the requested work**

A Client/Server Example

- The *header.h* file used by the client and server.

```
/* Definitions needed by clients and servers. */
#define TRUE 1
#define MAX_PATH 255 /* maximum length of file name */
#define BUF_SIZE 1024 /* how much data to transfer at once */
#define FILE_SERVER 243 /* file server's network address */

/* Definitions of the allowed operations */
#define CREATE 1 /* create a new file */
#define READ 2 /* read data from a file and return it */
#define WRITE 3 /* write data to a file */
#define DELETE 4 /* delete an existing file */

/* Error codes. */
#define OK 0 /* operation performed correctly */
#define E_BAD_OPCODE -1 /* unknown operation requested */
#define E_BAD_PARAM -2 /* error in a parameter */
#define E_IO -3 /* disk error or other I/O error */

/* Definition of the message format. */
struct message {
    long source; /* sender's identity */
    long dest; /* receiver's identity */
    long opcode; /* requested operation */
    long count; /* number of bytes to transfer */
    long offset; /* position in file to start I/O */
    long result; /* result of the operation */
    char name[MAX_PATH]; /* name of file being operated on */
    char data[BUF_SIZE]; /* data to be read or written */
};
```

An Example of a Server

```
#include <header.h>
void main(void) {
    struct message m1, m2;           /* incoming and outgoing messages */
    int r;                           /* result code */

    while(TRUE) {                   /* server runs forever */
        receive(FILE_SERVER, &m1);  /* block waiting for a message */
        switch(m1.opcode) {         /* dispatch on type of request */
            case CREATE:    r = do_create(&m1, &m2); break;
            case READ:      r = do_read(&m1, &m2); break;
            case WRITE:     r = do_write(&m1, &m2); break;
            case DELETE:    r = do_delete(&m1, &m2); break;
            default:        r = E_BAD_OPCODE;
        }
        m2.result = r;              /* return result to client */
        send(m1.source, &m2);      /* send reply */
    }
}
```

Client File Copy Example

```
#include <header.h>
int copy(char *src, char *dst){
    struct message ml;
    long position;
    long client = 110;

    initialize( );
    position = 0;
    do {
        ml.opcode = READ;
        ml.offset = position;
        ml.count = BUF_SIZE;
        strcpy(&ml.name, src);
        send(FILESERVER, &ml);
        receive(client, &ml);

        /* Write the data just received to the destination file.
        ml.opcode = WRITE;
        ml.offset = position;
        ml.count = ml.result;
        strcpy(&ml.name, dst);
        send(FILE_SERVER, &ml);
        receive(client, &ml);
        position += ml.result;
    } while( ml.result > 0 );
    return(ml.result >= 0 ? OK : ml result);
}
```

/* procedure to copy file using the server */
/* message buffer */
/* current file position */
/* client's address */

/* prepare for execution */

/* operation is a read */
/* current position in the file */
/* how many bytes to read*/
/* copy name of file to be read to message */
/* send the message to the file server */
/* block waiting for the reply */

/* operation is a write */
/* current position in the file */
/* how many bytes to write */
/* copy name of file to be written to buf */
/* send the message to the file server */
/* block waiting for the reply */
/* ml.result is number of bytes written */
/* iterate until done */
/* return OK or error code */

Multitiered architectures

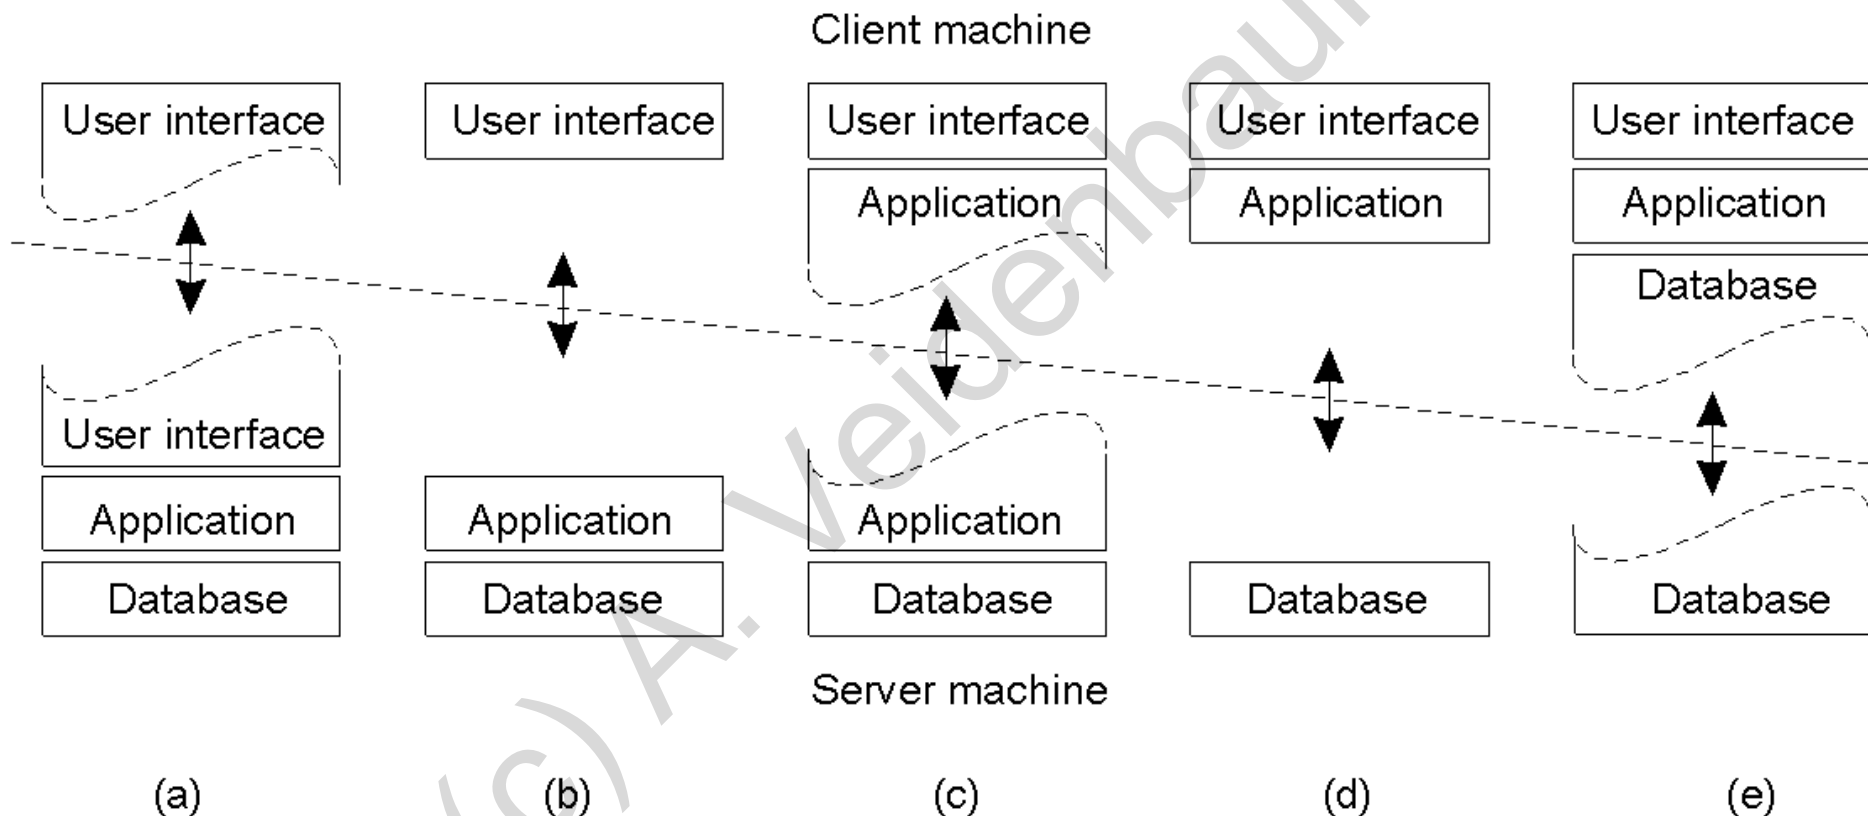
- **Recall the three levels we discussed**
 - **User interface level**
 - » Deals with input, output functions like display, getting input
 - **Processing level**
 - » This is the application itself
 - **Data level**
 - » Stores the state or data and defines primitive operations on it
- **Many different implementations are possible**
 - Along the three levels above
- **Let us start with a (physically) 2-tiered architecture**
 - **Client system and server system**

Client-Server Architectures

- **Simplest:**
 - Client implements the user interface level
 - Server implements the rest
- **What is the problem with this system?**
- **Many other ways possible**
 - Can actually distribute programs in a level/layer
 - Two-tiered architecture

Multitiered Architectures

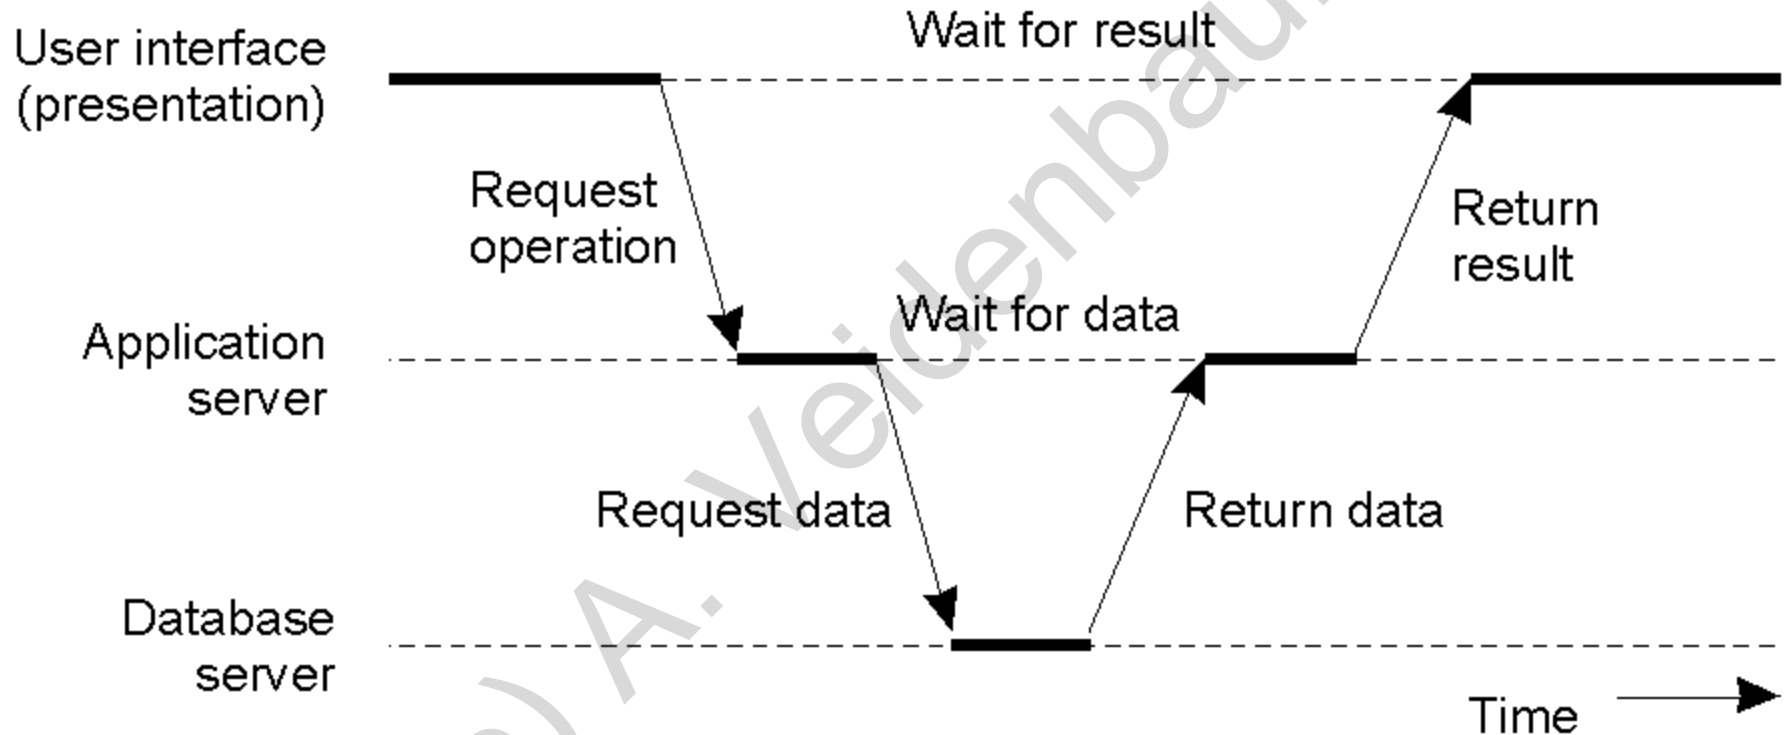
- Alternative client-server organizations (a) – (e).



- **Organizations d, e are very popular**
 - A data base is a backend, just serves requests
- **Overall, a somewhat limited organization**
 - A server cannot act as a client!
- **Solution: three-tiered architecture**
 - Processing level distributed among servers
 - » Common in transaction processing
 - Web sites

Three-tiered Architecture

- An example of a server acting as a client.



Decentralized Organizations

- So far we have seen a **vertical distribution** in a DS
 - Tiers or levels
- Distributed Systems also need replication
 - This leads to a **horizontal distribution**
 - » Client or server can be physically split into logic parts
 - Each part equivalent but operates on its own share of data
 - » Peer to peer systems are an example
 - The load is distributed, statically or dynamically
- Let us look at peer to peer systems (P2)
 - Processes are all equals in P2P systems
 - But the system function is carried out by all processes
 - » Each can serve as a server and a client (servant)


- The P2P system are symmetric
- All processes have to be able to communicate
 - All to all type of communication
- So it's a question of a communication network
 - TCP/IP works just fine
 - » Allows arbitrary communication patterns
 - » But perhaps slow as there is only one link per system
- Sometimes may want to restrict the *all to all*
 - Overlay networks, where nodes have a limited number of connections to the network

Structured P2P systems

- Use a specific, deterministic topology
 - Allows efficient data lookup
- Generally based on a *semantic-free index*
 - Each item has a unique key used as an index
- One common approach is to use a hash
 - $\text{key}(\text{data item}) = \text{hash}(\text{data item's value})$
 - The whole P2P system now stores (key, value) pairs
- Each node is assigned an identifier from the hash space and stores a subset of key with that hash
- Access via a lookup function
 - $\text{existing node} = \text{lookup}(\text{key})$


P2P example

- Uses a 4-D hypercube w/ 16 nodes
- A value is hashed to one of the 16 keys
 - A lookup is local
- A node that does not have an item sends the request
 - Toward the node with the correct key



Hybrid Architectures (HA)

- HAs combine features of client-server and decentralized architectures
- An important example are Edge-server systems
 - This is how end users connect to the Internet



Users connect to the internet via edge servers