


**CompSci 131**

# **Parallel and Distributed Systems**

**Prof. A. Veidenbaum**

# Parallel Programming on SMPs

- **SMP = Symmetric Multi-Processor**
  - **Shared Memory, “tightly coupled”**
    - » **Processors are connected to memory via a bus or switch**
      - A bus-based multiprocessor



- **Caches are a (partially) replicated memory**
  - **Improve performance by lowering average memory latency**
  - **But they create a memory coherence problem**
    - » **We'll look at how this can happen**

# Multiprocessors (1)

- **Multi-core vs multiprocessor:**
  - A core is a processor in a single-chip microprocessor
    - » Intel Xeon Gold is a multi-core processor w/ up to 28 cores
  - A multiprocessor today is built from multiple multi-cores
    - » Intel terminology: a processor = socket = 1 Xeon Gold
- **Multi-cores require additional hardware support for**
  - » Inter-core communication and interrupts
  - » Mutual exclusion
  - » Indivisibility of memory RMW
  - » Cache coherence
- **May use processor or core interchangeably**
  - Should be clear from the context

# Multiprocessors (1)

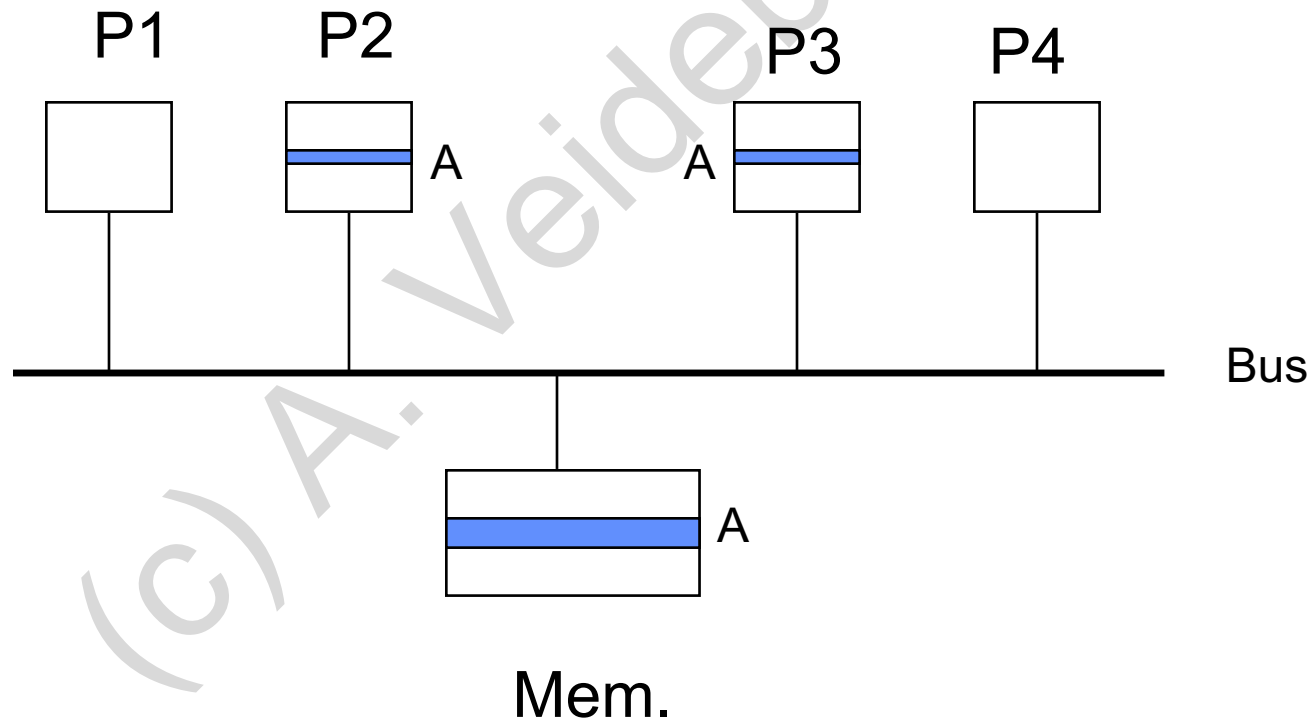
- **Multi-cores require additional OS support for**
  - **Process isolation and management**
  - **Mutual exclusion**
  - **VM management**
- **Mutual exclusion is tied to process management**
  - **A process/thread blocked on a semaphore is suspended**
- **Let's look at hardware support for mutual exclusion**

# Multiprocessors (2)

- Intel instructions for locking, indivisibility
  - Aka atomic processor instructions
    - » atomically update a memory location
  - An atomic instruction uses a lock prefix on the instruction and has its destination operand in memory.
  - These instructions can use a lock prefix
    - » ADD, ADC, AND, BTC, BTR, BTS, CMPXCHG, CMPXCH8B, DEC, INC, NEG, NOT, OR, SBB, SUB, XOR, XADD, and XCHG
- Inter-processor interrupt (IPI)
  - Each core has an APIC
  - A core puts the IV, destination ID into its local APIC
    - » Its local APIC delivers it to the remote APIC

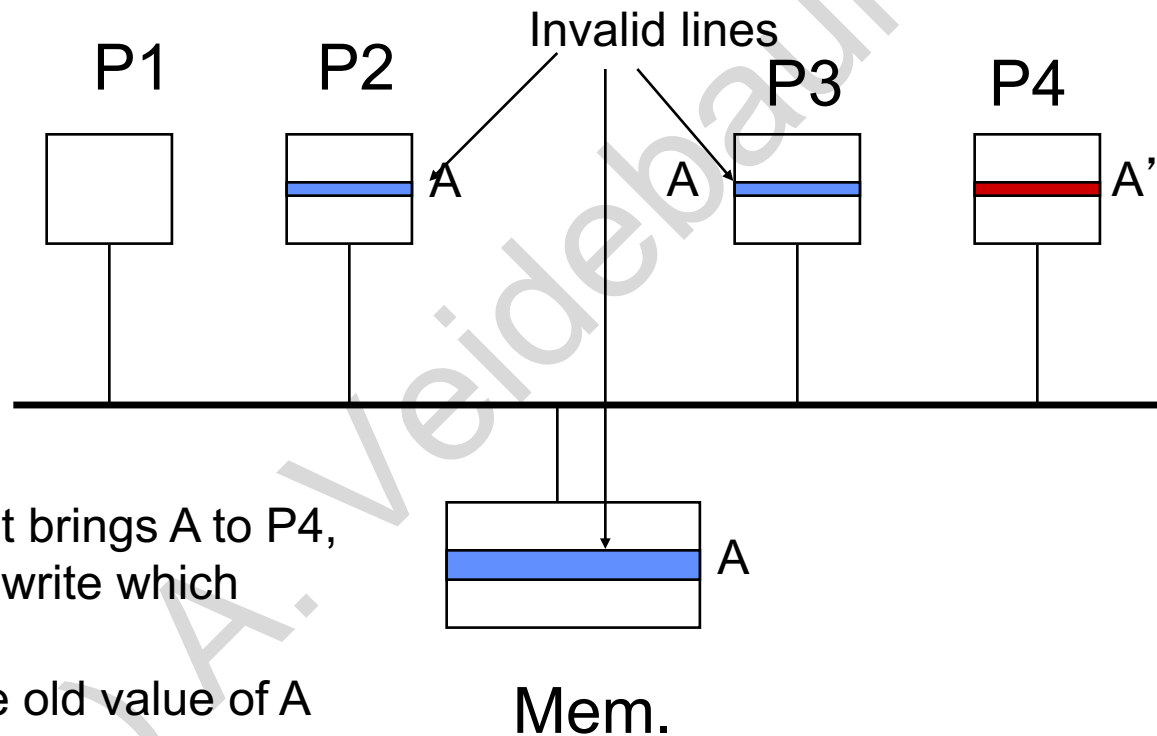
# Cache Coherence: The problem

- P2 reads A, it is automatically stored in its cache
- P3 reads A, same deal



# Now a Write occurs

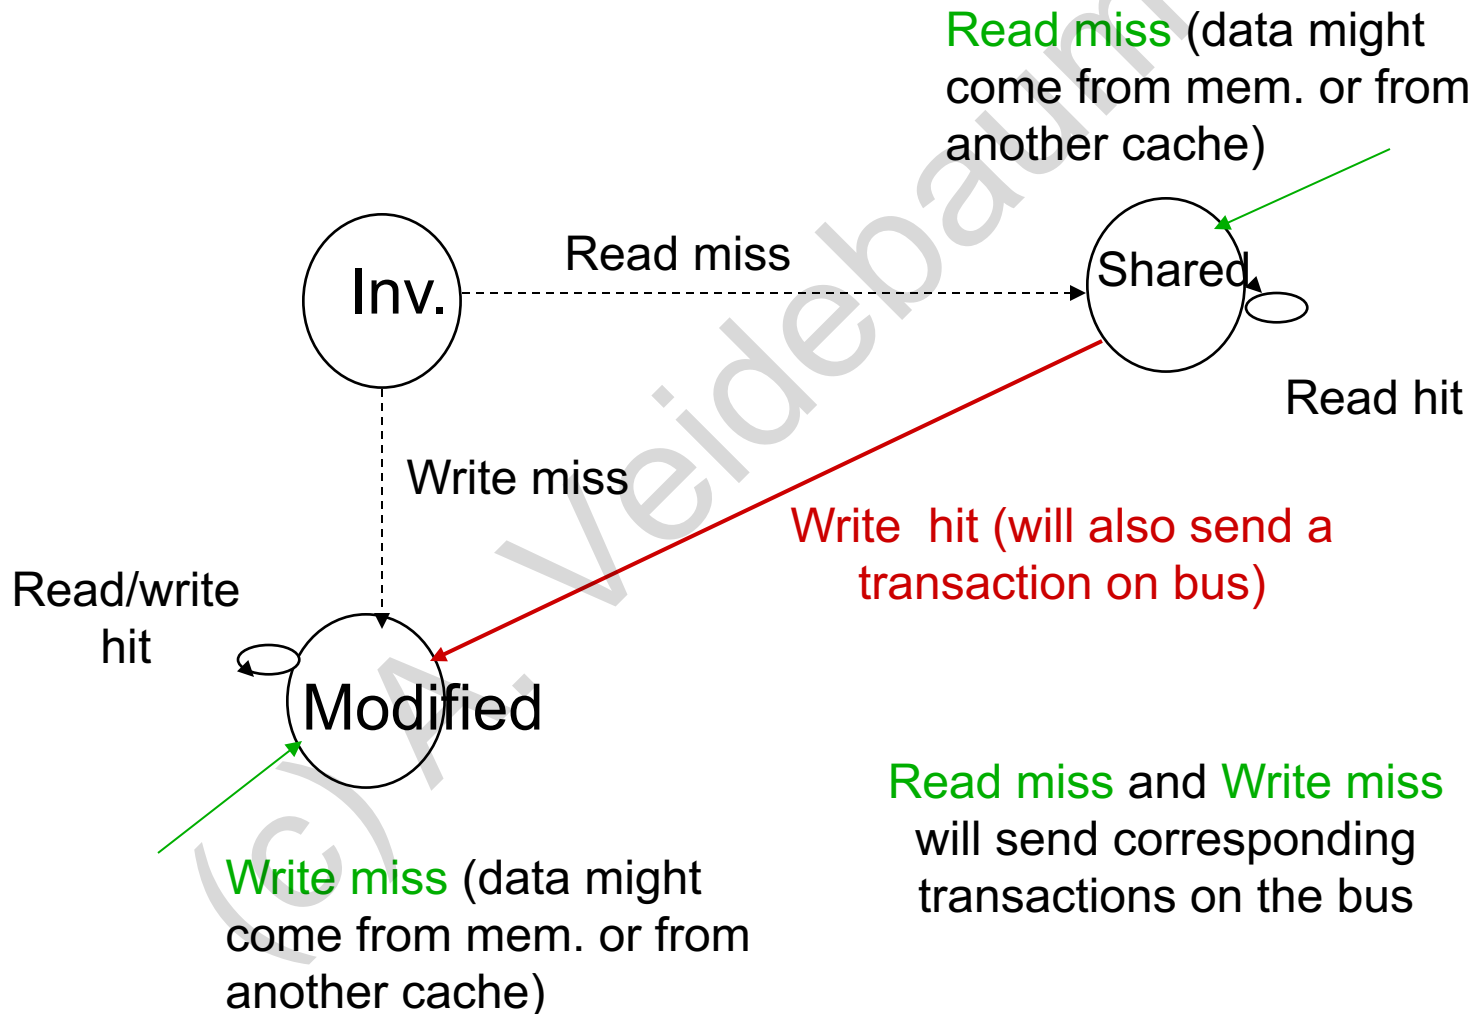
- P4 has a write miss on line A



- A write miss first brings A to P4,
- followed by the write which creates A'
- P2, P3 have the old value of A
- So does memory!

- Multiple copies are no longer identical, this is the coherence problem
- Solved by a hardware cache coherence protocol

# A 3-State Protocol: Processor Actions





# SMP Programming Models

- The abstract model is parallel RAM (PRAM)
  - Concurrent reads, writes *to the same address* are OK
    - » What does the latter mean?
  - Good for theoretical modeling, not for programming
- All SMP models are based on multiple threads
  - Each thread is largely a von Neumann model
    - » But how do they interact, access memory?
- All models define and enforce a memory model
  - Access ordering to *same and different* RAM locations
    - » From same and different threads

# Time ordering of memory accesses

- Time ordering of accesses to a *DRAM* address
  - Accesses from a thread/core are (program) ordered
  - Accesses from different threads/cores are not ordered
  - The only guarantee of ordering is via synchronization
- Time ordering of accesses to different RAM addresses (same or different threads)
  - *Cannot be observed in general*
    - » *Why for the same thread?*
  - Can only be guaranteed via synchronization

- **The “sequentially consistent” model**
  - **Memory accesses within a thread are strictly sequential**
  - **Accesses to memory are an interleaving of sequential accesses from each thread**
    - » **Access order for a thread remains the same in all interleavings**
    - » **All threads see the same ordering**
- **This is said to be preferred by programmers**

# Parallel programming on SMPs

- Goal – find parallelism in a program to speed it up
- Approach - define concurrent tasks to run on each core
  - Divide data or program into independent “chunks”
    - » Data parallelism or thread parallelism
    - » A chunk = unit of work to be done sequentially
  - Assign each chunk to a different core to run in parallel
- Data parallelism is easiest to grasp
  - Independent operations on multiple data elements
    - » find all occurrences of x in an array, add two arrays, etc.
  - Sometime operations are dependent, e.g. array sum
    - »  $x = x + A[i]$ 
      - Can something be done in parallel here?
    - » Update all elements of a linked list
      - Is there parallelism here?

# Data Parallel Programming

- The program remains quite similar to sequential
  - Each processor runs the same code on a subset of data
    - Hence data parallelism
  - A sequential example

```
for (i=1; i< N; i++)  
    A[i] = B[i] * C[i]
```
  - A data parallel version

```
for (i=1; i< N; i+Chunksize)  
    spawn sum( i, ChunkSize, A[i], B[i], C[i]);
```

    - Assume “spawn” means “create a new thread”
- What is a good chunk size?
  - Application and system dependent
  - Large chunks – less parallel work, less contention
  - Small chunks – more parallelism, larger overhead, contention

# Thread-Parallel Programming

- This is more general than Data-parallel
  - Threads may execute different code on different data
- A sequential version of main

```
...  
call func1(...)  
...  
call func2(...)
```

- A thread parallel version

```
...  
rc1 = pthread_create( tid1, NULL, func1, (void *) arg1)  
...  
rc2 = pthread_create( tid2, NULL, func2, (void *) arg2)  
...  
pthread_join(tid1,NULL)  
pthread_join(tid2,NULL)
```

- **A pthread exits when it completes func execution**
  - Or executes `pthread_exit`
- **A program is finished when all threads finished**

# Another way to look at parallel programming

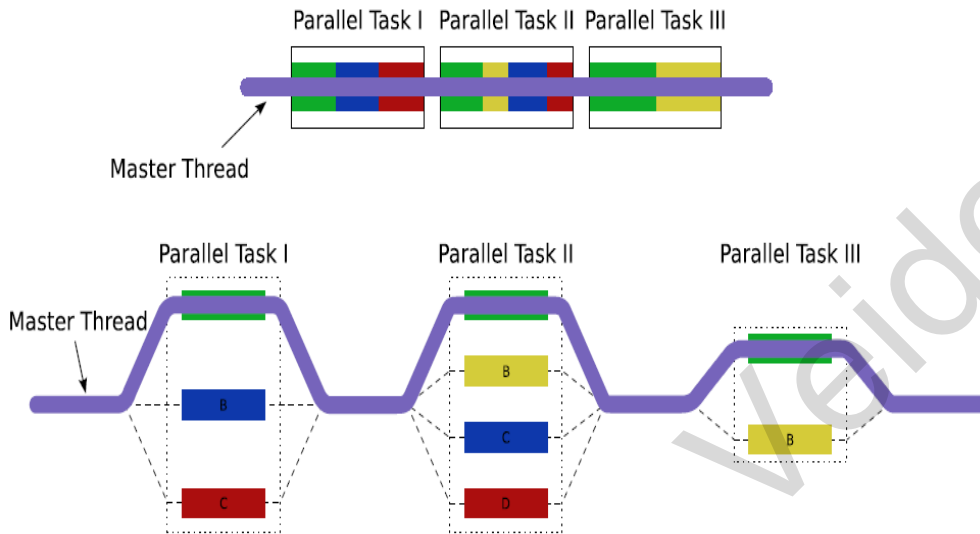



figure from Wikipedia: "OpenMP"

## fork-join



dag



# OpenMP

- Thread programming is difficult and low-level
- Another approach is to use OpenMP
  - Another standard, simplifies loop-based parallelization
    - » Can also do tasks
  - Based on Prof. Kuck's work on parallel loops at UIUC
- A user inserts directives (pragmas) in her code

```
#pragma omp parallel for schedule(static, ChunkSize)
for (i=1; i< N; i++)
    A[i] = B[i] * C[i]
```

  - Compiler does the rest
- A “parallel for” must have independent iterations
- Has an implied barrier at the end (unless "nowait")

# Data Dependence

- **Independent** = no data dependencies across iterations
  - A programmer needs to understand and satisfy/resolve data dependencies between program or data chunks
- **Data dependence: output, data, and anti- dependencies**
  - **Data (flow) dependence** (between iterations  $i$  and  $i+3$ )  

```
for (i=1; i< N; i++)  
  A[i] = ...  
  ... = A[i-3]
```
  - **Output dependence** (between all iterations)  

```
for (i=1; i< N; i++)  
  x = x + A[i]
```
  - **Anti-dependence** (use  $A[i+3]$  before iteration  $i+3$  overwrites it)  

```
for (i=1; i< N; i++)  
  ... = A[i+3]  
  A[i] = ...
```

↑ a data race

- **OpenMP *parallel for* cannot have any dependencies**
  - **A compiler does not check!**
- **Programming approach:**
  - **Select compute intensive loops in the code**
  - **Make sure there are no dependencies between loop iterations**
  - **Apply omp constructs**
- **An “omp parallel for” loop has an implicit barrier at the end**
  - **Waits for all tasks to finish**

# OpenMP

- **C/C++/Fortran compilers support OpenMP**
  - e.g. Intel ICC, GNU GCC, MS Visual Studio, LLVM
    - » A compiler cannot check if the resulting parallel code is correct
- **A programmer defines number of tasks**
  - Env. variable, dynamically in code
- **OpenMP run-time schedules tasks**
  - relies on pthreads and kernel scheduler
  - thread to core mapping can be forced to stay fixed (*affinity*)
  - Implements multiple scheduling algorithms
    - » Static, dynamic, guided

# OpenMP

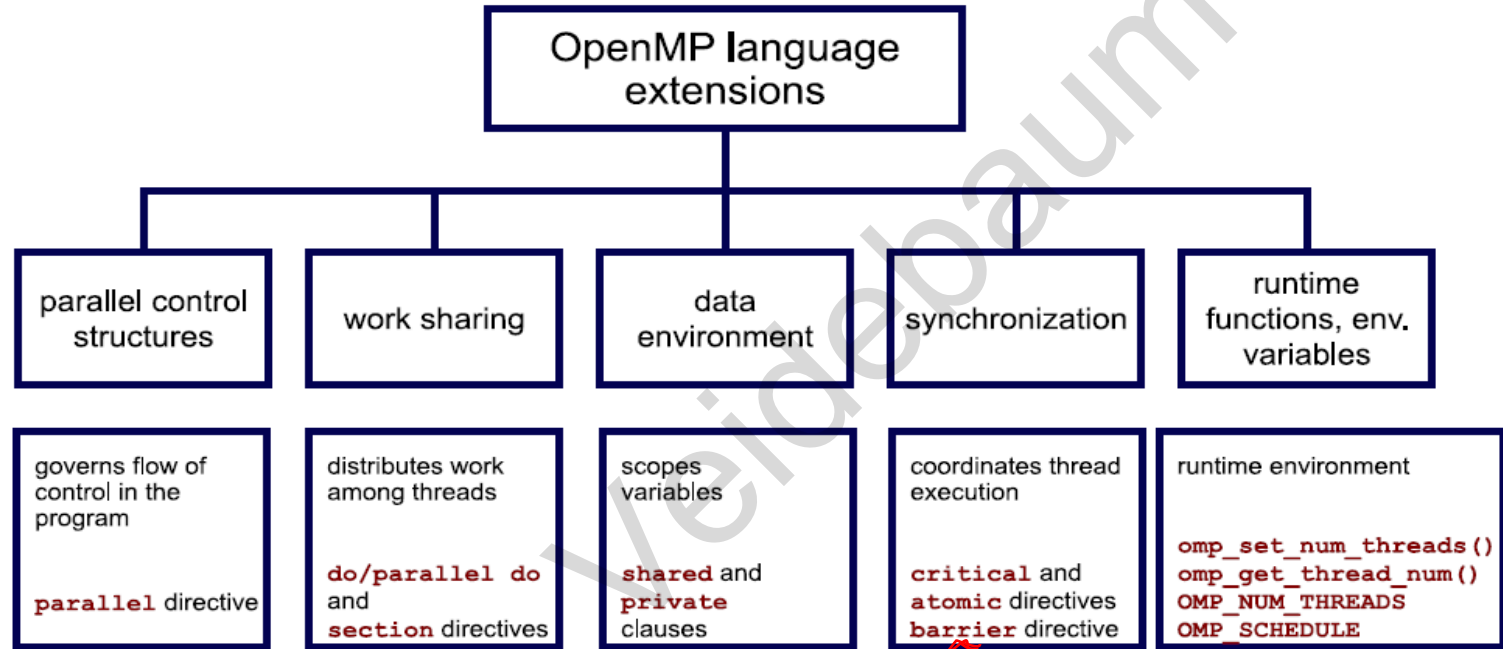


figure from Wikipedia: "OpenMP"

threads launched

pragma

lib func

OpenMP 2.5+

Current OpenMP version is 5

# OMP matrix transpose and scale

```
void transpose(int *M, int X, int n) {  
    #pragma omp parallel  
    {  
        int i, j;  
        #pragma omp for schedule(dynamic,20)  
        for (i=0; i<n; i++) {  
            for (j=i+1; j<n; j++) {  
                M[i,j] = M[j,i] * X  
            }  
        }  
    }  
}
```

- Tasks are created
- Shared variables
  - M[], X, n
- Private variables
  - i, j
- Scheduling options
  - static, dynamic, etc
  - chunk size
- Implicit barrier where tasks join

# OpenMP compilation, execution

- **Compilation**

```
$ gcc -fOpenMP my_OpenMP.c -o my_OpenMP
```

**NOTE:** the compiler ignores the OpenMP directives if `-fOpenMP` is omitted

- **Execution**

```
$ export OMP_NUM_THREADS=3
```

```
$ ./my_OpenMP
```

# Cilk

- Developed by Leiserson at MIT, starting in 1994
- An extension to C
  - for constructing multithreaded parallel programs
  - Uses a set of keywords: cilk, spawn, sync, etc.
- Uses a source-to-source compiler and a standard C compiler
- Intel TBB is based on CILK



# Cilk


- **Software developer exposes parallelism**
  - identifies functions free of side effects, which can be treated as independent tasks and executed in parallel.
    - » annotate such functions with the cilk keyword
    - » invoke them with the *spawn* keyword
  - Places the *sync* keyword which indicates that execution of a function/main cannot proceed until all previously spawned functions have completed and returned their results to the parent

# Cilk

- Runtime takes care of task distribution/scheduling
  - Tasks are stored in dequeues, one per processor
  - Employs *work stealing*
    - » each processor fetches tasks from the top of its own queue
    - » when a processor queue is empty it picks another processor at random and "steals" a task from the bottom of its dequeue,
      - LIFO order
- While OMP parallel for does what is called *work sharing*
- Work stealing is better for dynamic workloads
  - DAGs
  - Widely varying execution times per task

# cilk example: Fibonacci series

```
int fib (int n) {  
    int x, y;  
    if (n<2) return n;  
    x = fib (n-1);  
    y = fib (n-2);  
    return x+y;  
}
```



```
cilk Int fib (int n) {  
    int x, y;  
    if (n<2) return n;  
    → x = spawn fib (n-1);  
    → y = spawn fib (n-2);  
    sync;  
    return x+y;  
}
```

- Sample output

```
fib(1)=1  
fib(2)=1  
fib(3)=2  
fib(4)=3  
fib(5)=5  
fib(6)=8  
fib(7)=13  
fib(8)=21  
fib(9)=34  
fib(10)=55
```

- Compile


```
$ cilk -O3 fib.cilk -o fib
```

- Execution

```
$ ./fib -nproc 4 30
```

# cilk example: Fibonacci series

```
int fib (int n) {  
    int x, y;  
    if (n<2) return n;  
    x = fib (n-1);  
    y = fib (n-2);  
    return x+y;  
}
```



```
cilk Int fib (int n) {  
    int x, y;  
    if (n<2) return n;  
    x = spawn fib (n-1);  
    y = spawn fib (n-2);  
    sync;  
    return x+y;  
}
```

- Sample output

```
fib(1)=1  
fib(2)=1  
fib(3)=2  
fib(4)=3  
fib(5)=5  
fib(6)=8  
fib(7)=13  
fib(8)=21  
fib(9)=34  
fib(10)=55
```

AV/UCI

- Compile

**`$cilkc -O3 fib.cilk -o fib`**

- Execution

**`./fib -nproc 4 30`**

In general the software developer must restructure/express the code in a “*divide & conquer*” fashion.

The topology of the dag is implicitly conveyed by the placement of the *spawn* keywords.