

CompSci131

Parallel and Distributed Systems

Prof. A. Veidenbaum

Today's topics

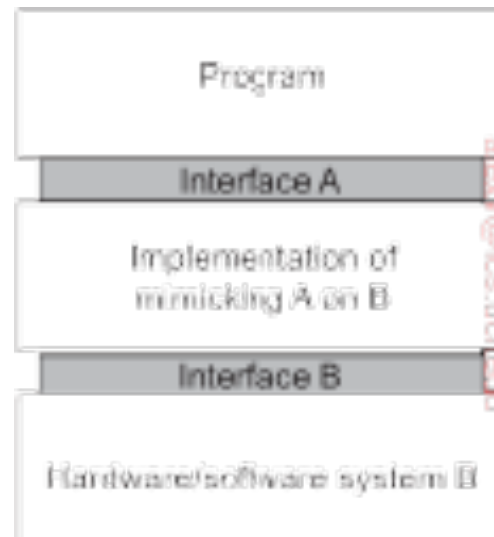
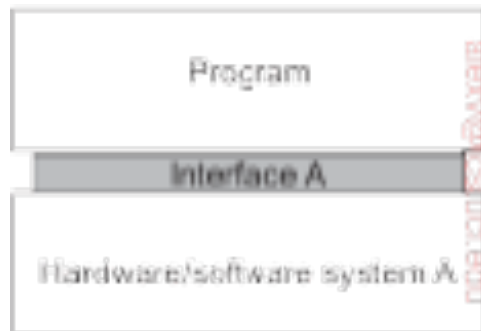
- **Virtualization**
- **Code migration**
- **Reading assignment:**
 - **Today: 3.2, 3.5**
 - **Next time: 4.2-4.3**

Last Lecture Covered

- **SMP parallel programming**

Virtualization

- **DS virtualization helps port sftw to new systems**
 - E.g. provide a new interface to old sftw
 - More general, higher level than wrappers
 - Also helps share resources
 - So IBM created a “hypervisor” and could run either OS on it




- **Networking is pervasive, so can be on remote system**

(Resource) Virtualization

- **Already saw an OS ‘virtualizing’ a single processor**
 - For a process
 - But different OS’s do it differently
- **IBM had to maintain backward compatibility and run the old OS on new hardware in the 80’s**
 - They also had a new OS
 - So they created a “hypervisor” and could run either OS on it
- **The hypervisor is ‘a private OS’ on top of raw hrdw**
 - Intel’s term=Virt. Machine Monitor (VMM)
 - » Intel provides VM hardware extensions (VMX)
- **Will see what processor support is required**


Types of Virtualization

- Computers generally provide 4 types of interfaces
 - E.g. provide a virtual machine
- These interfaces are at 3 different levels
 - Hrdw/Sftw level – instruction set architecture (ISA)
 - » User and priveleged instructions
 - OS level – system calls
 - API level – usually library calls
 - » May hide some system calls



Types of Virtualization


- Virtualization mimics these interfaces
- Two ways of doing it
 1. A runtime system providing an abstract ISA to apps
 - » Java virtual machine - interpreted
 - » Windows apps running on Linux – emulated
 - Quite hard for Windows sys calls
 - » This is virtualization for a single process



Types of Virtualization

2. A system layer shielding the original hrdw


- » Offers the hrdw interface to layers above
 - Same as underlying hrdw OR different hrdw
- » Known as “native VMM” ‘cause it runs on original hrdw
 - Can still clearly see the 4 interfaces



- Note that the VMM interface can be used by multiple apps simultaneously


Types of Virtualization

- The native VMM has to provide access to and manage various resources
 - E.g. storage, networking, etc through drivers
- But the OS already does it, so why re-implement?
- Another virt. type is known as a hosted VMM



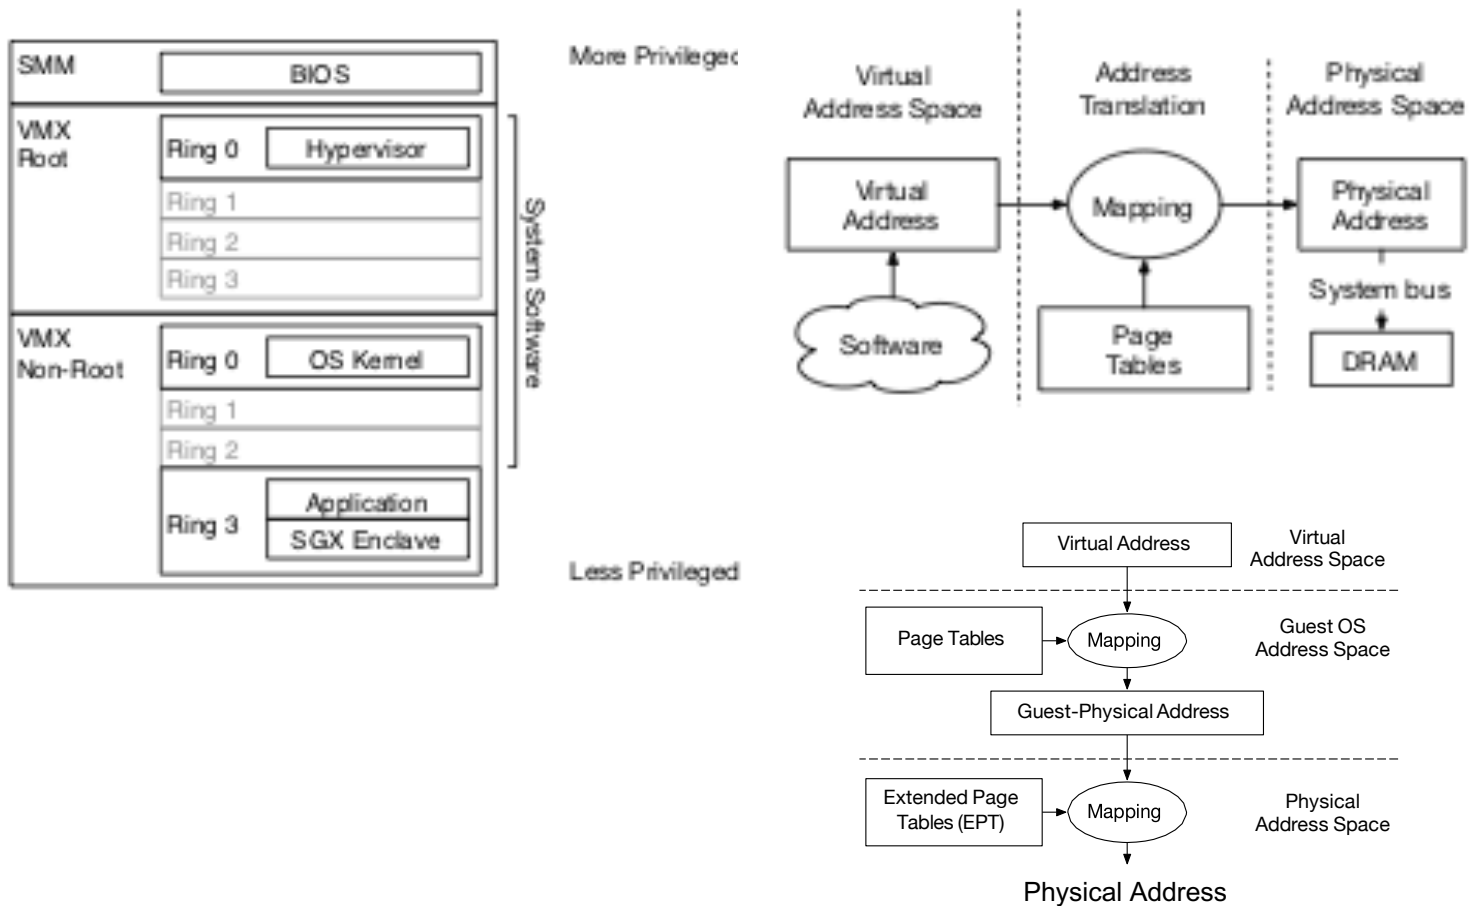
Virtualization Performance

- Modern virtualized systems are highly efficient
 - They run as much code as possible natively
Directly on underlying hrdw for most non-priveldged instructions



Processor Virtualization

- Hypervisor divides up hardware resources to guest OS's: page frames in DRAM, access to I/O devices,...



An example of system virtualization

- Recall the concept of Infrastructure-as-a-Service
 - Aka IaaS, is built on virtualization
- Consider Amazon EC2
 - It allows one to create a number of virtual servers
 - » Using preconfigured machine images (AMIs)
 - An OS kernel with a number of services
 - » One such image is LAMP
 - Linux kernel, Apache Web server, MySQL, PHP libraries
 - » A launched AMI is an EC2 instance
 - Each instances gets 2 IP addresses: public & private
- A customer cannot tell
 - Where his code is running
 - Who, if anyone, is she sharing the EC2 resources
 - What fraction of a CPU, etc does she get

Code migration

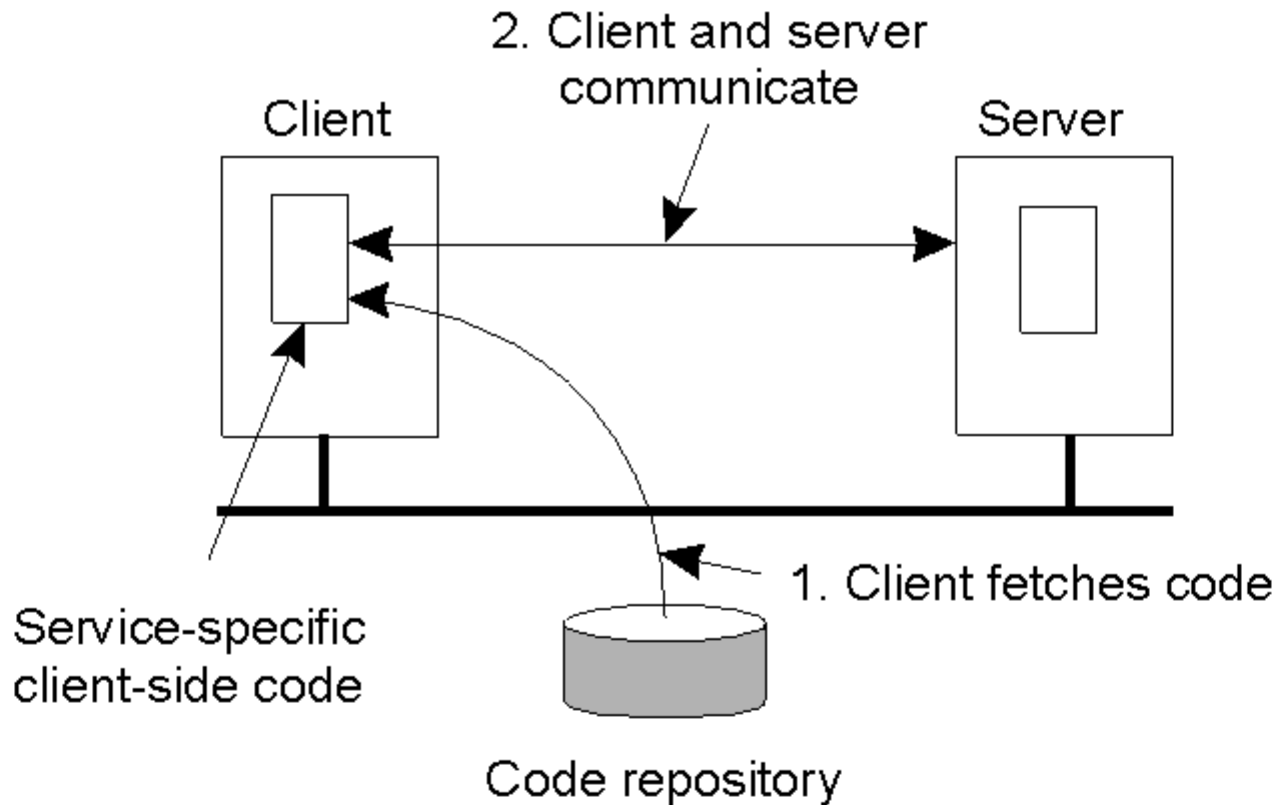
- It is sometimes desirable to move programs or objects
 - Traditionally done as process migration
 - » Move a process from one DS node to another
 - » Process is a running program!
 - A very difficult and costly move, must be sure it is worth it
 - Usually to get better performance due to load balancing
- Today the goal is to get better loading of each machine
 - Allows data centers to serve more customers -> get more \$\$\$
 - Done by migrating complete virtual machines + their apps
- At the same time, network delays may dominate
 - So it may be worth moving programs to data
 - » E.g. when extensively using a large database

Advantages of code migration

- **Allows load consolidation**
- **Can improve load balancing and thus performance**
 - Or exploit parallelism and create many copies to do work
- **Reduce communication cost**
- **Flexibility in design**
 - don't need to decide where application runs
 - Can configure clients when they run, not at design time
 - » **Dynamic client configuration**

Dynamic Client Configuration

- The principle of dynamically configuring a client to communicate to a server. The client first fetches the necessary software, and then invokes the server.
- Allows client server protocols to be changed transparently



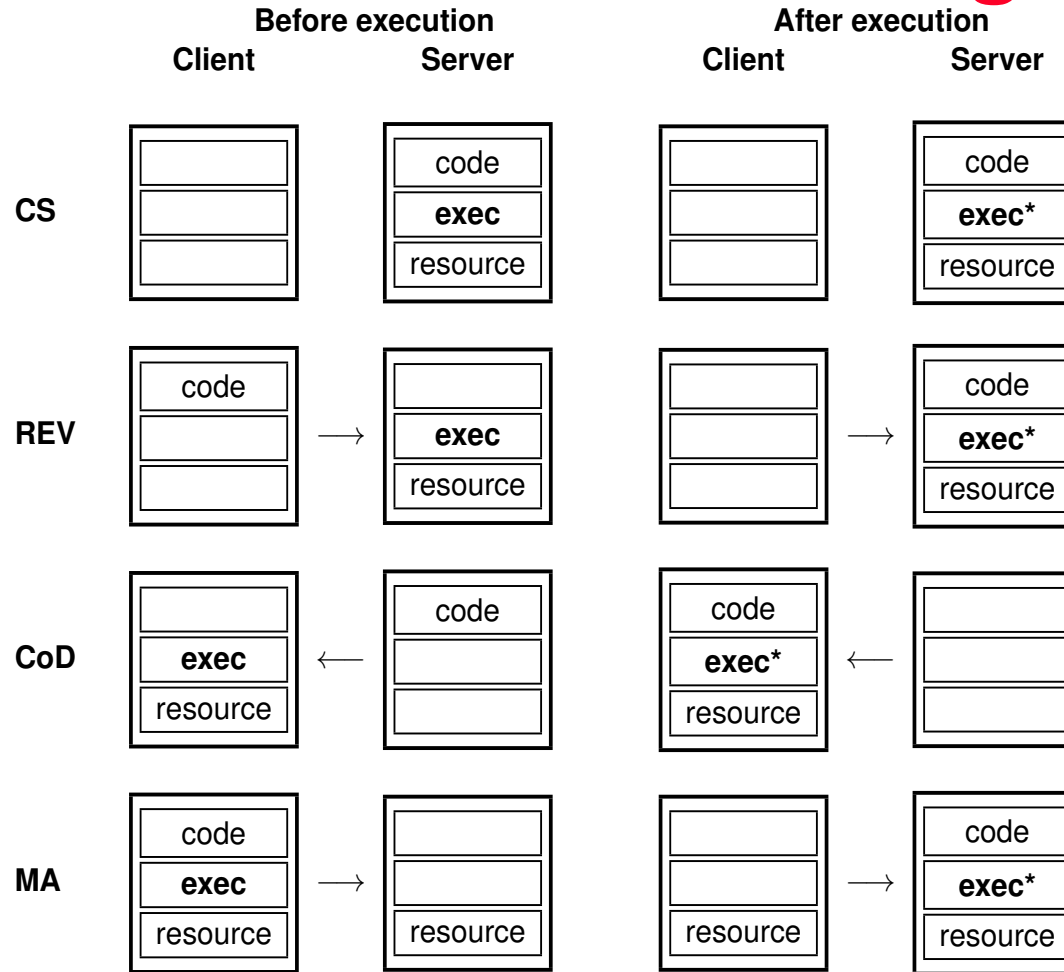
Models for Code Migration

- **Models assume a process consists of**
 - A Code Segment
 - A Resource Segment
 - An Execution Segment (state, including data)
- **Weak mobility migrates only the code segment**
 - (Re)starts the program from initial state
 - » **Advantage: simplicity**
 - Example - Java applets
- **Strong mobility also migrates execution segment**
 - Allows a running process to be stopped, moved, restarted

Mobility characteristics

- **Sender-initiated**
 - **Send code to a server**
 - » **E.g. send a program to a compute server**
- **Client-initiated (or receiver-initiated)**
 - **Browser loads an applet**
 - » **This is harder! Why?**
- **Lastly, can run as same process or a new one**
 - **New process is expensive but provides protection**
- **For strong mobility can also clone a process**
 - **An exact copy**
 - **Runs in parallel with the original process**

Alternatives for code migration

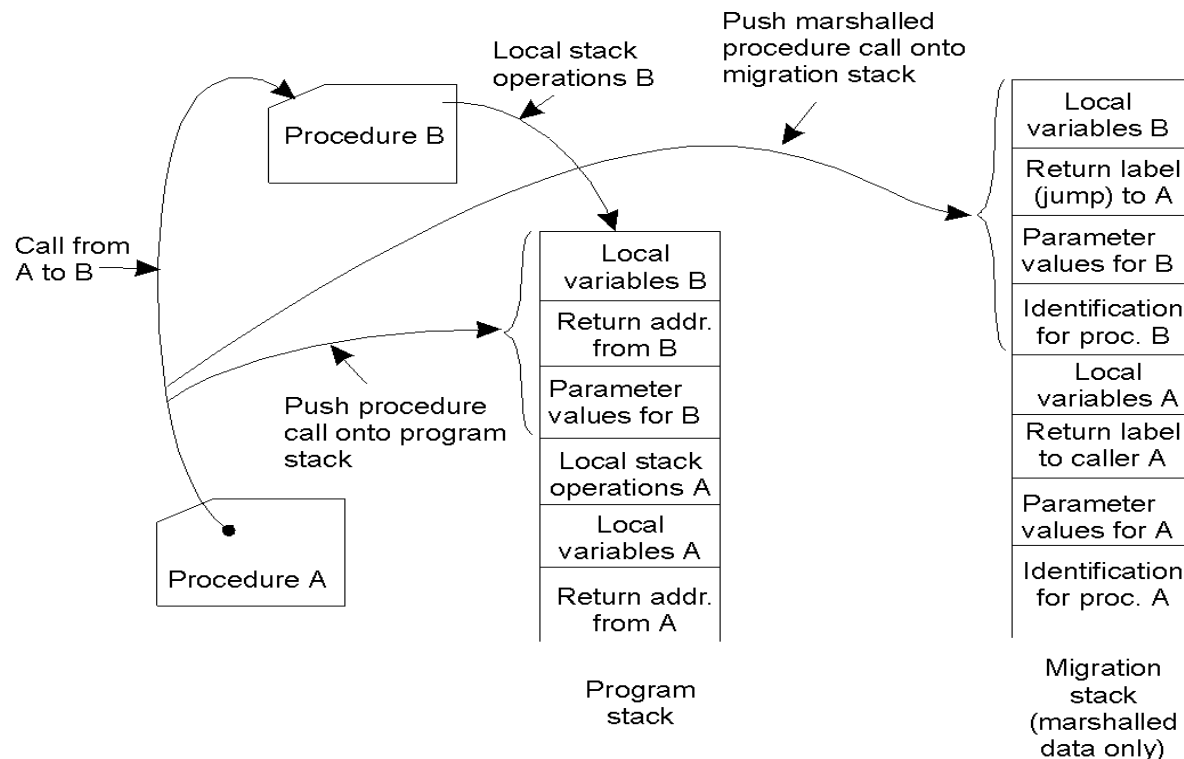


CS: Client-Server
CoD: Code-on-demand

REV: Remote evaluation
MA: Mobile agents

Migration in heterogeneous systems

- Requires executing code on different types of system
- Hard for strong mobility, main problem is migrating the execution segment
- Can use a migration stack to migrate an execution segment, for C and Java
 - Restricted to specific points in the program - function calls
 - » Using machine independent execution stack (migration stack)
 - Maintained in parallel with a regular stack



- **Best supported by a compiler that generates both normal and migration code/stack for C**
- **Another solution is using VMs (e.g. Java),**
 - **An intermediate language, compiler-generated**
 - **A restricted set of libraries**
 - **Protection of local resources**
 - **Interpretation or JIT**