

**CompSci 131**

# **Parallel and Distributed Systems**

**Prof. A. Veidenbaum**

# Today' s topics

- Naming in DS
- Reading assignment:
  - Today: Naming, 5.1-5.3
  - Next time: Coordination, 6.1-6.2
    - » Complete the assignment before next class

# Last Lecture Covered

- More MPI functions, implementation

# Naming *and* locating

- A name is a string of bits identifying an entity
  - Entity can be anything: device, file, connection, address...
  - May or may not be human-friendly
- Entities are operated on/accessed *via an access point*
  - Access points are themselves entities!
    - » Their name is called an *address (aka location)*
- A name needs to be resolved to a location
- It is good to separate a name and an address!
  - Allows migration, re-organization, multiple access points
- Need a DS *naming service, which may be distributed*
  - Scalable, but harder to implement

# Types of Names

- An entity name that is independent from its address is called location independent
  - An FTP server vs an address of a host where FTP server resides
    - » or mobile phone numbers
- A true (or unique) identifier
  - An identifier the refers to at most one entity
  - Each entity is referred to by at most one identifier
  - An identifier always refers to the same entity
    - » is never reused
- An identifier is thus a unique reference to an entity
  - An address may or may not be an identifier

# Locating entities

- How does one map a name to a location or an address?
  - Aka resolve a name
- Simplest way – a table of <Ent, Addr> pairs
  - But does not work well in many DS settings, for instance when entities migrate
- There are two main ways of locating entities
  - Using flat names
  - Using hierarchical names

# Flat Names

- Let's start with simple solutions for locating an entity with a flat name
  - Using a broadcast
    - » ARP is an example for a local net
  - Using forwarding pointers
    - » Client-server stubs
    - » Can short-circuit after first access

# Locating an entity using broadcast

- **A local or a wireless network can**
  - Use broadcast to send an entity identifier
  - All network nodes check if they have the entity
  - Only nodes that offer an access point reply
- **Used in address resolution protocol (ARP)**
  - IP address to ethernet address
    - » No routing tables required!
- **Can also use restricted multicasting**
  - Limited to a group of hosts
    - » Supported in the Internet
- **Not a scalable solution**

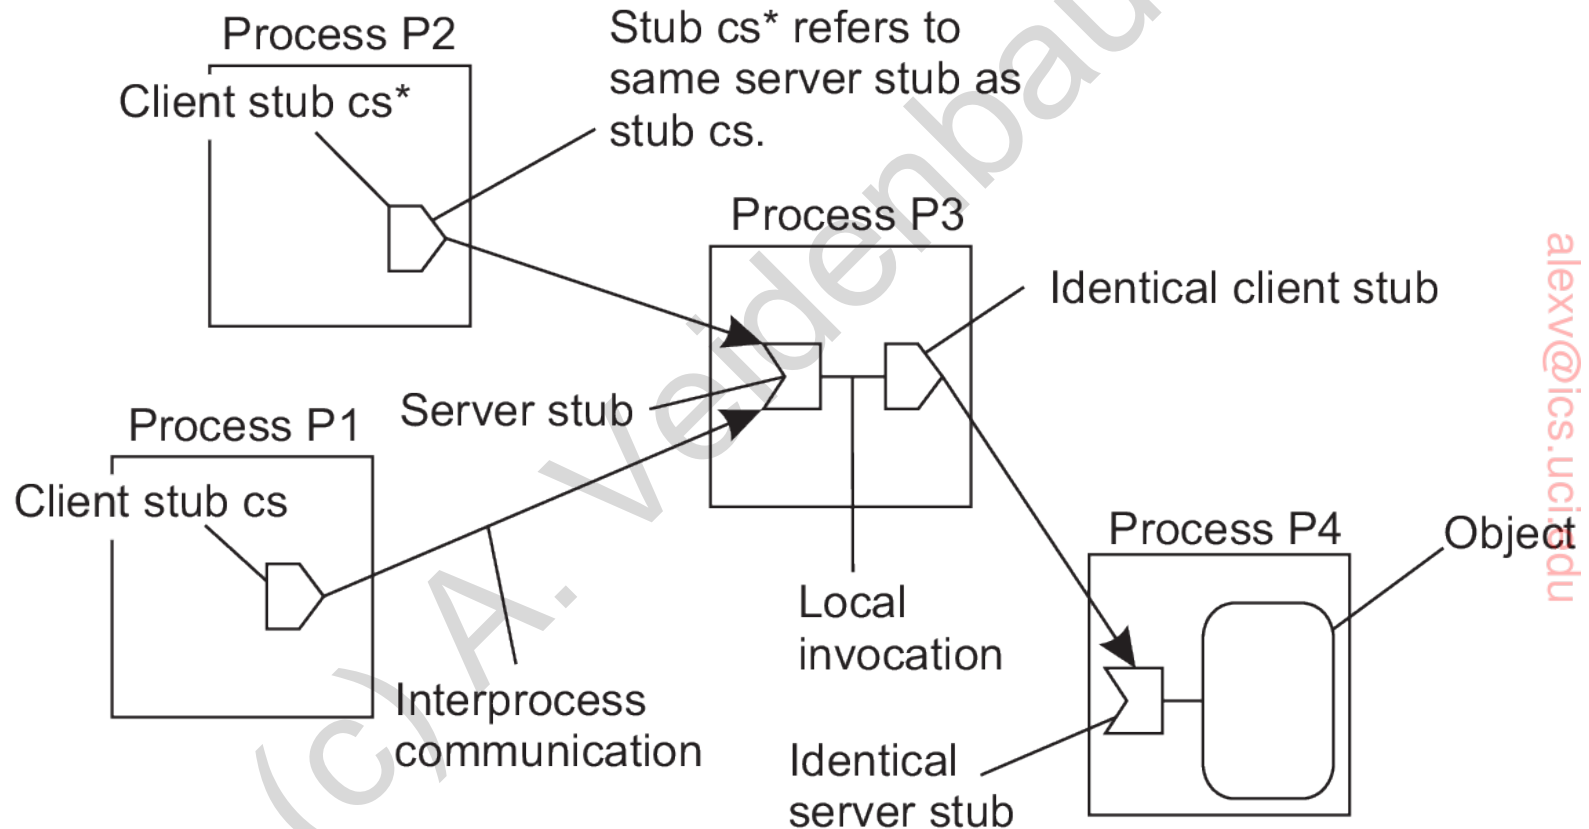


# Using Forwarding Pointers

- **A relocated entity leaves a forwarding pointer at original location**
  - **Very simple**
  - **Scalable? Not for frequently moving servers!**
    - » **Can create a very long forwarding chain**
- **Example: used with RPC**
  - **A moving server leaves behind a forwarding stub**
    - » **Just to send the request to a new location**
  - **The server at the new location handles the request**
- **Want to short-cut this for performance**
  - **Send new location to requestor for update**

# Example of Forwarding Pointers

- A relocated entity leaves a forwarding address

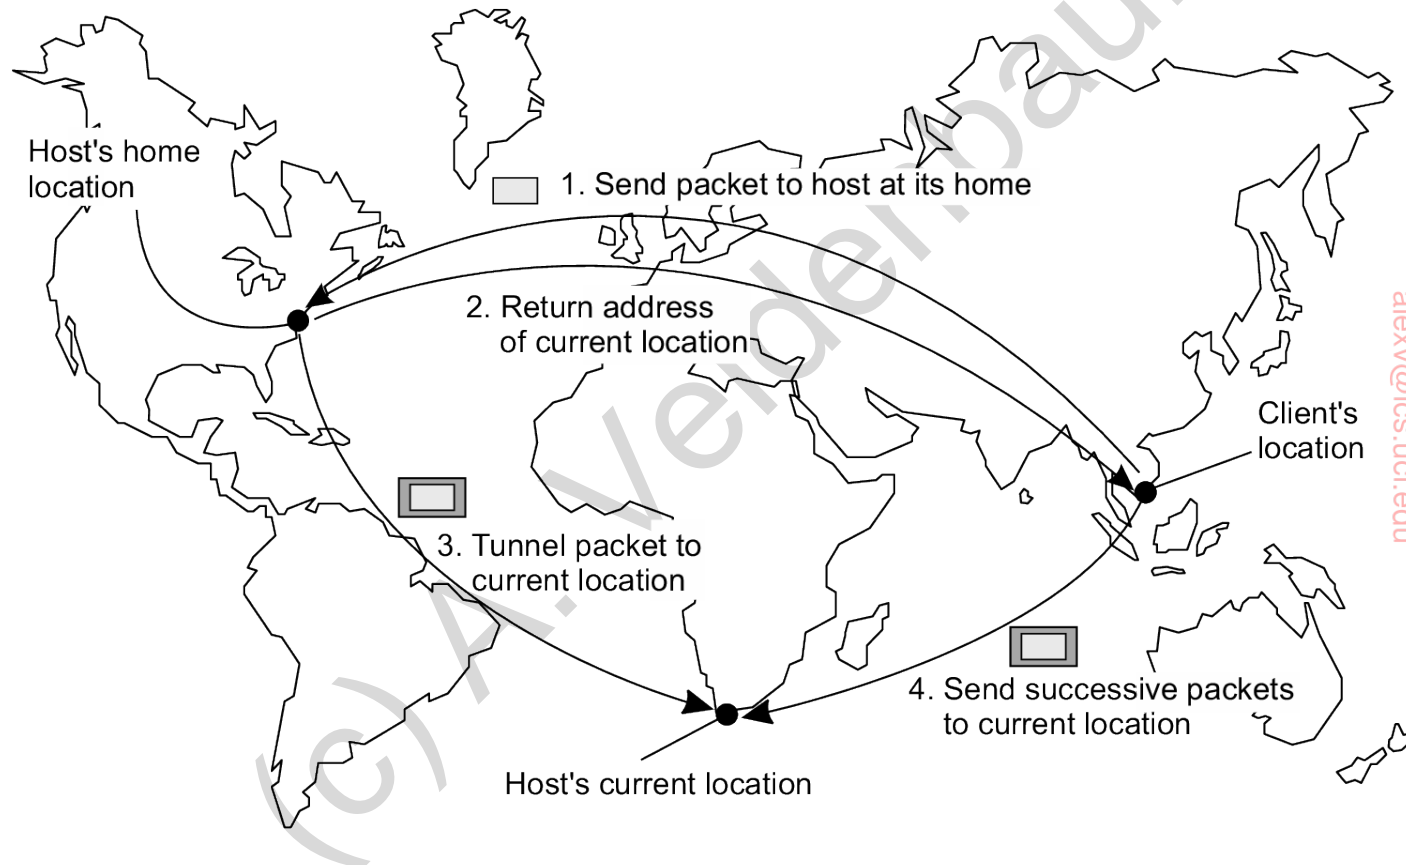


# Home-based Approaches

- **A home location is unique**
  - Can always find an entity starting at its home
- **Used for mobile entities in large networks**
  - Cell phone roaming
  - Mobile IP, using a “home agent”
- **Uses a forwarding pointer as entity moves**
- **Largely invisible to applications**
  - They only refer to original IP address
- **Drawback – long delays possible**

# Home-based Example

- Messages travel around the world



# Hierarchical Approaches


- A Network is divided into sets of domains
- Domains may be 'nested' inside each other
- Two node types in a domain (plus sub-domains):
  - A directory node
    - » Contains an entry for each entity in the domain
  - A leaf node
    - » Typically a local network or a mobile cell
      - Contains the entity address
- One root node (no predecessors)
  - E.g. a directory tree

# Hierarchical Approaches

- **A leaf domain contains:**
  - An *address* for each entity in the domain
- **A directory node contains ‘location’ records**
  - For each entity in the domain it knows the sub-domain the entity is in
    - » But *not* the address
- **There can be replicas of an entity!**
- **Lookup of entity N in this organization**
  - originates at a leaf and goes up
  - Finds the first directory that contains the entry for N
  - This exploits locality!

# Hierarchical Domain Example

- A tree of domains



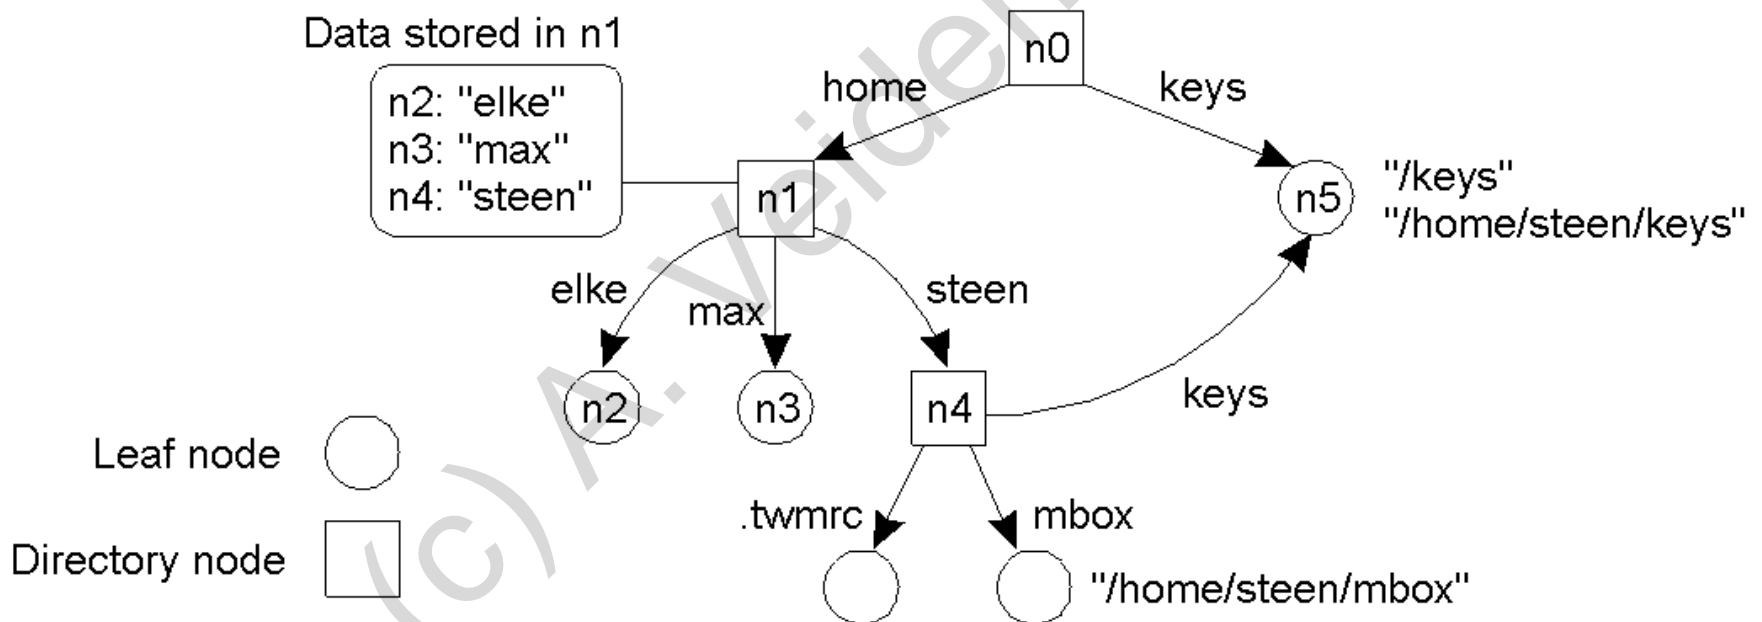
# Structured naming

- Flat names are not human-friendly
- Structured names are human readable names
  - File names, domain names
- Names are organized into *Name Spaces (NS)*
- NSs are represented as directed graphs
  - Leaves are named entities
  - Directory nodes have *labeled* outgoing edges
- Any node in this graph is itself an entity
  - Has an identifier
  - Directory nodes have a dir table: (edge,node id)



# An example

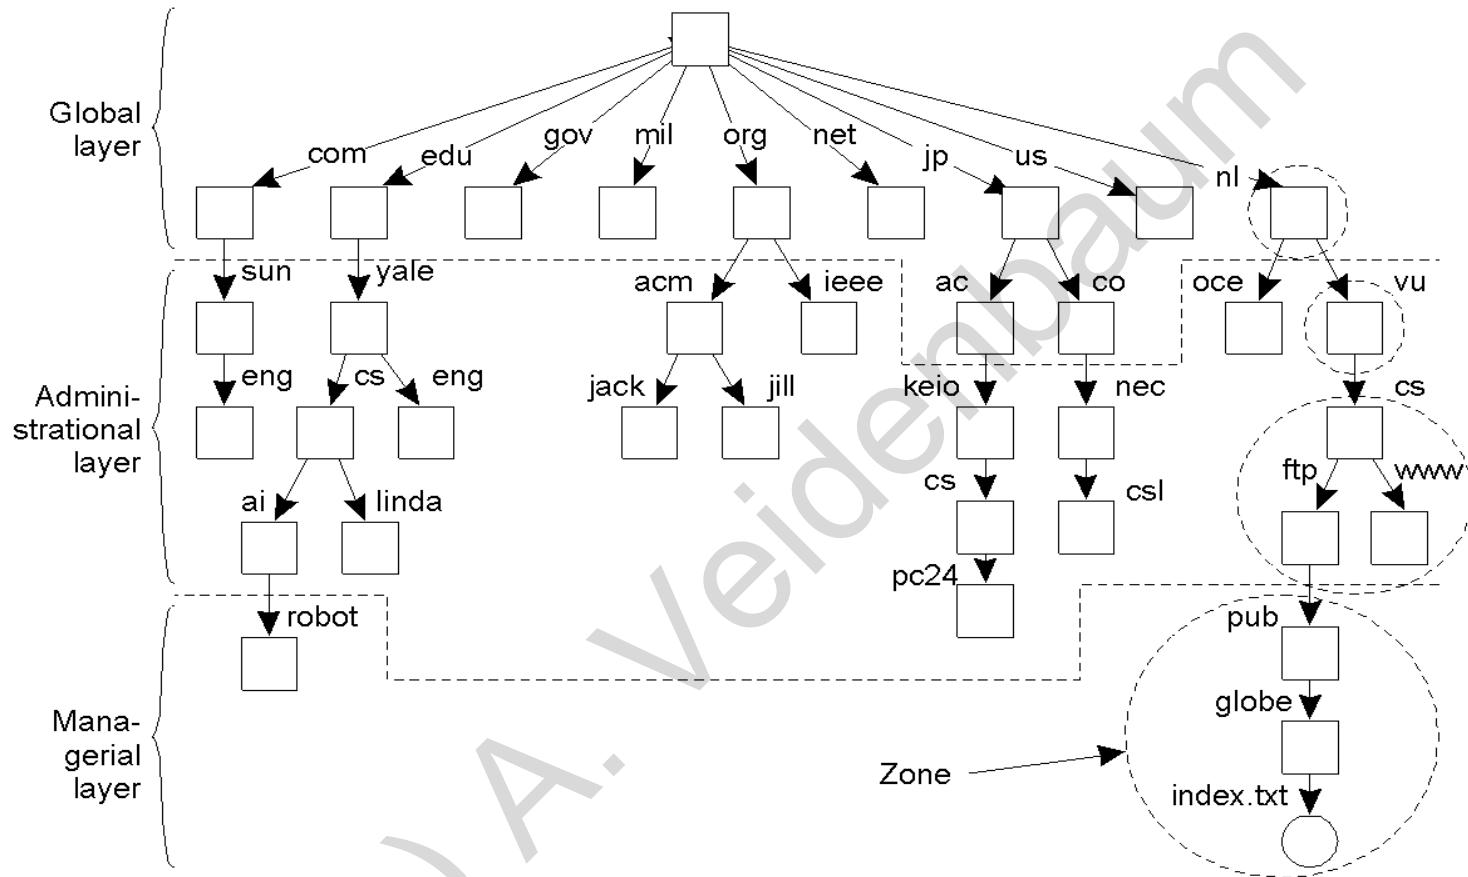
- **General, but close to most file systems**
  - Path name = a sequence of edge labels
  - From root: absolute path
  - Otherwise a relative path



# Name Resolution

- The process of looking up a name
  - Go through edges in naming graph, until node found
- Closure mechanism: where and how to start
  - i.e. where is root?
    - » Hard-coded as inode0 in Unix
  - Local/global names
    - » HOME environment label

# Name Space Distribution (DNS)



- **Global layer:** “stable” nodes (non-overlapping “zones”)
- **Administration layer:** a node per one organization, relatively stable

# Properties

| Item                            | Global    | Administrational | Managerial   |
|---------------------------------|-----------|------------------|--------------|
| Geographical scale of network   | Worldwide | Organization     | Department   |
| Total number of nodes           | Few       | Many             | Vast numbers |
| Responsiveness to lookups       | Seconds   | Milliseconds     | Immediate    |
| Update propagation              | Lazy      | Immediate        | Immediate    |
| Number of replicas              | Many      | None or few      | None         |
| Is client-side caching applied? | Yes       | Yes              | Sometimes    |

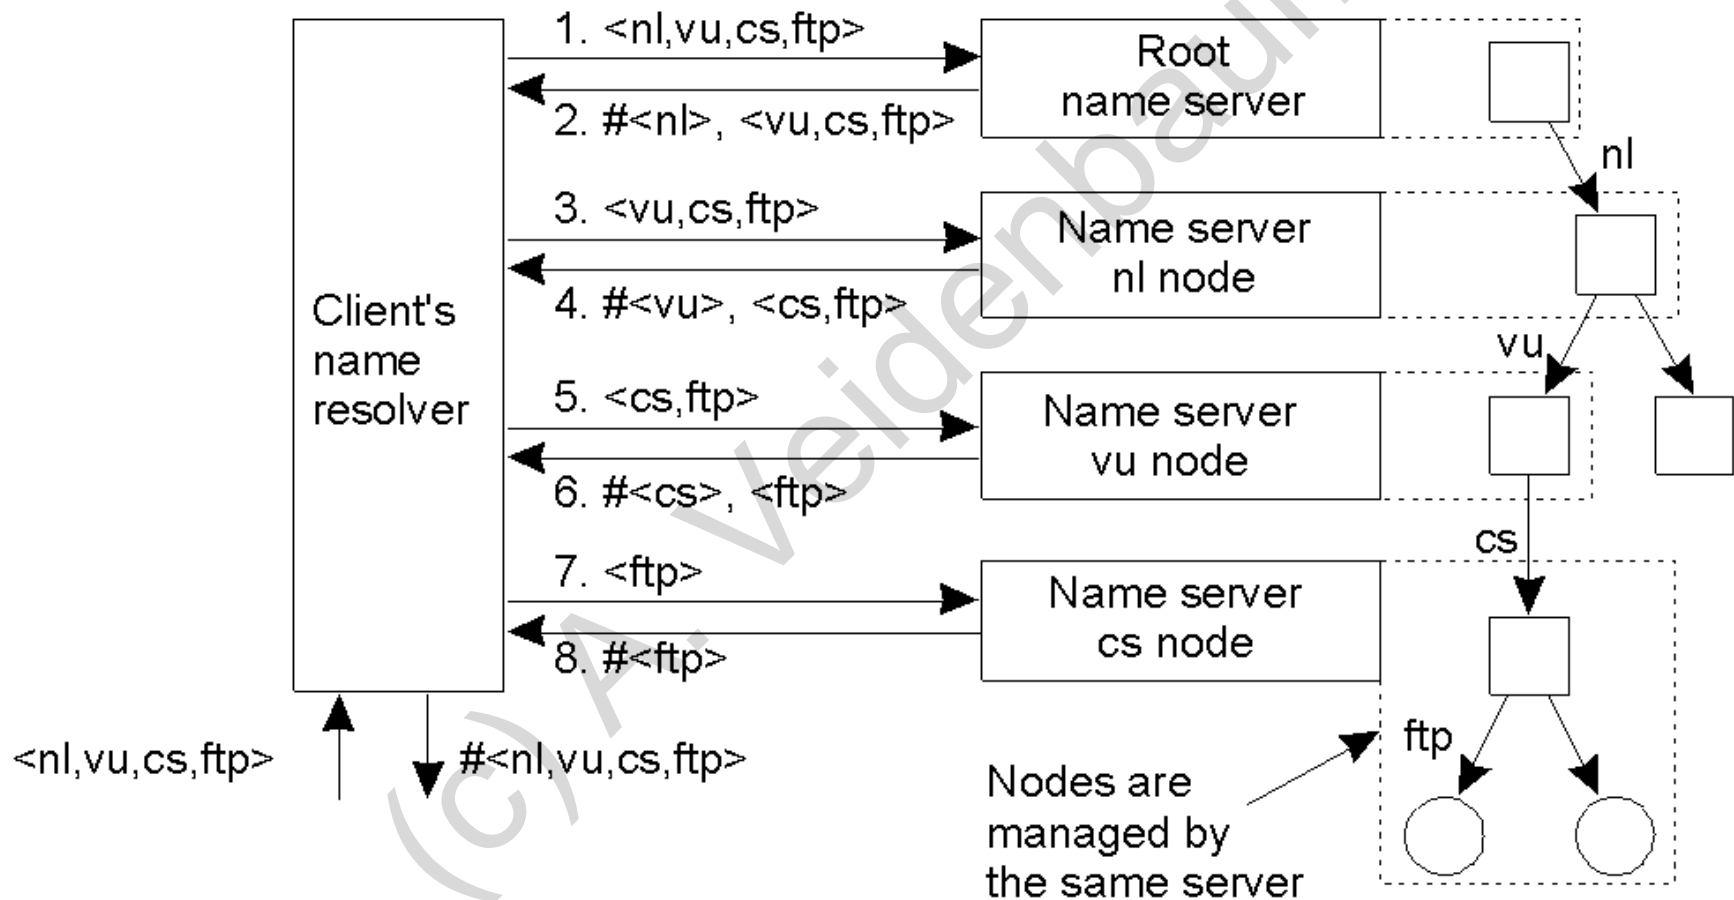
- **Used in DNS, but is generally applicable**
- **Properties affect availability and performance**
  - **Caching, replication**
    - » **Caching allows faster resolution, but creates consistency problems**

# Name Resolution

- **DS have “name resolvers”**
  - A client has access to a local name resolver
- **Two ways of resolving:**
  - Iterative
  - Recursive

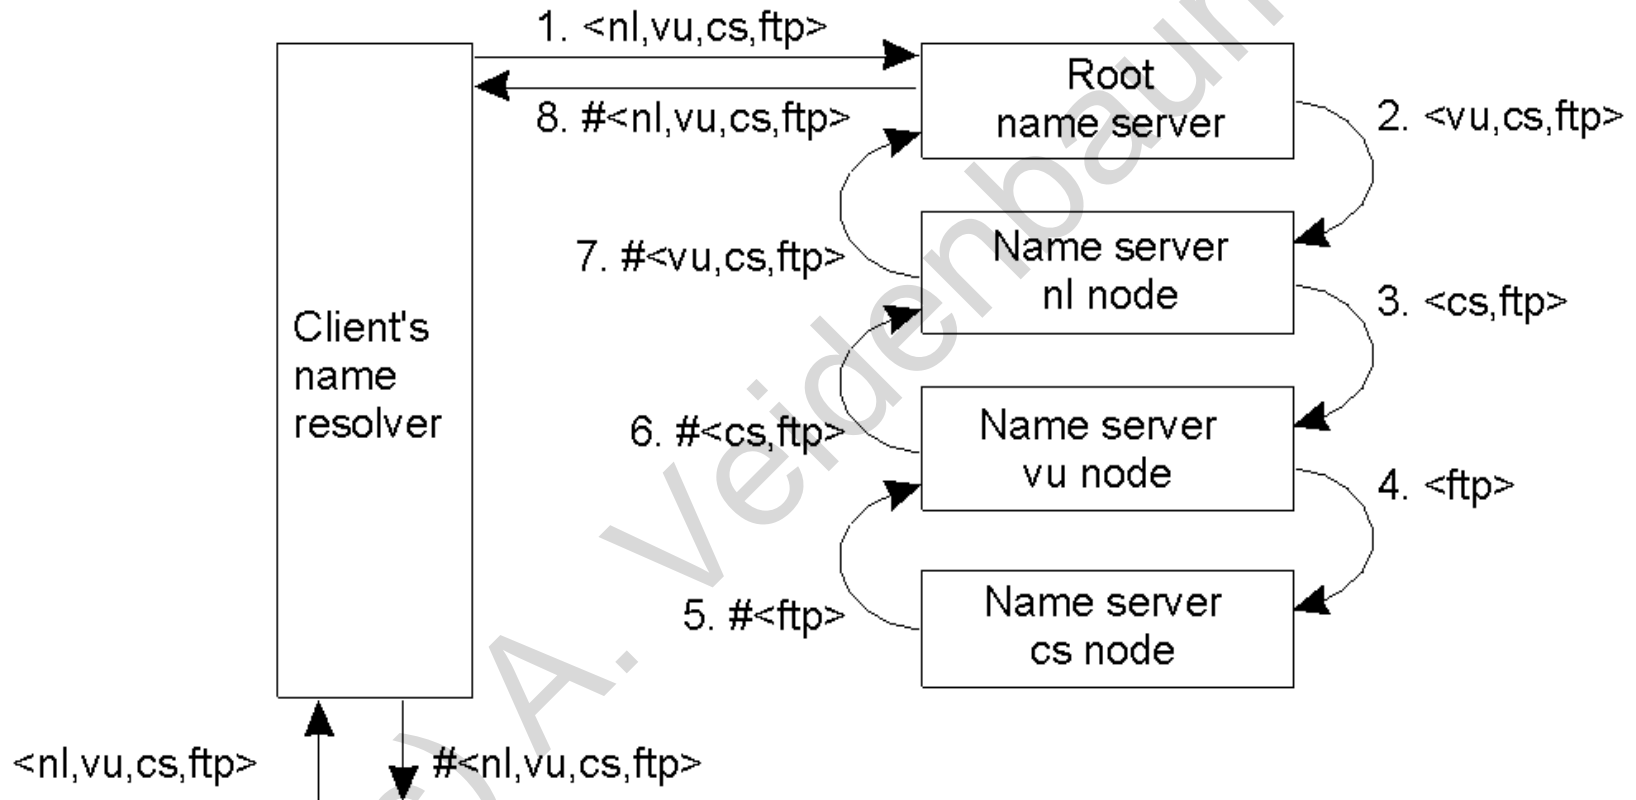
# Iterative Name Resolution

- Each name in the path is resolved by new client-server exchange




# Recursive Name Resolution

- Only result returns to client, name servers communicate



# Implementation of Name Resolution (4)



- **What is an advantage of one vs the other?**



# Domain Name System

- Name to IP address resolution
  - Primarily for hosts and mail servers
- 30 years old, but working just fine
- Main reason: design simplicity

# The DNS Name Space

- The most important types of resource records forming the contents of nodes in the DNS name space.

| Type of record | Associated entity | Description   |
|----------------|-------------------|---|
| <b>SOA</b>     | <b>Zone</b>       | Holds information on the represented zone                     |
| <b>A</b>       | <b>Host</b>       | Contains an IP address of the host this node represents       |
| <b>MX</b>      | <b>Domain</b>     | Refers to a mail server to handle mail addressed to this node |
| <b>SRV</b>     | <b>Domain</b>     | Refers to a server handling a specific service                |
| <b>NS</b>      | <b>Zone</b>       | Refers to a name server that implements the represented zone  |
| <b>CNAME</b>   | <b>Node</b>       | Symbolic link with the primary name of the represented node   |
| <b>PTR</b>     | <b>Host</b>       | Contains the canonical name of a host                         |
| <b>HINFO</b>   | <b>Host</b>       | Holds information on the host this node represents            |
| <b>TXT</b>     | <b>Any kind</b>   | Contains any entity-specific information considered useful    |

# DNS database for the zone *cs.vu.nl*.

- 3 name servers
- 3 mail servers
  - Diff. priorities
- www, ftp servers
- Laser printers

| Name              | Record type | Record value                              |
|-------------------|-------------|---|
| cs.vu.nl          | SOA         | star (1999121502,7200,3600,2419200,86400) |
| cs.vu.nl          | NS          | star.cs.vu.nl                             |
| cs.vu.nl          | NS          | top.cs.vu.nl                              |
| cs.vu.nl          | NS          | solo.cs.vu.nl                             |
| cs.vu.nl          | TXT         | "Vrije Universiteit - Math. & Comp. Sc."  |
| cs.vu.nl          | MX          | 1 zephyr.cs.vu.nl                         |
| cs.vu.nl          | MX          | 2 tornado.cs.vu.nl                        |
| cs.vu.nl          | MX          | 3 star.cs.vu.nl                           |
| star.cs.vu.nl     | HINFO       | Sun Unix                                  |
| star.cs.vu.nl     | MX          | 1 star.cs.vu.nl                           |
| star.cs.vu.nl     | MX          | 10 zephyr.cs.vu.nl                        |
| star.cs.vu.nl     | A           | 130.37.24.6                               |
| star.cs.vu.nl     | A           | 192.31.231.42                             |
| zephyr.cs.vu.nl   | HINFO       | Sun Unix                                  |
| zephyr.cs.vu.nl   | MX          | 1 zephyr.cs.vu.nl                         |
| zephyr.cs.vu.nl   | MX          | 2 tornado.cs.vu.nl                        |
| zephyr.cs.vu.nl   | A           | 192.31.231.66                             |
| www.cs.vu.nl      | CNAME       | soling.cs.vu.nl                           |
| ftp.cs.vu.nl      | CNAME       | soling.cs.vu.nl                           |
| soling.cs.vu.nl   | HINFO       | Sun Unix                                  |
| soling.cs.vu.nl   | MX          | 1 soling.cs.vu.nl                         |
| soling.cs.vu.nl   | MX          | 10 zephyr.cs.vu.nl                        |
| soling.cs.vu.nl   | A           | 130.37.24.11                              |
| laser.cs.vu.nl    | HINFO       | PC MS-DOS                                 |
| laser.cs.vu.nl    | A           | 130.37.30.32                              |
| vucs-das.cs.vu.nl | PTR         | 0.26.37.130.in-addr.arpa                  |
| vucs-das.cs.vu.nl | A           | 130.37.26.0                               |

# DNS Implementation (2)

| Name                 | Record type | Record value         |
|----------------------|-------------|----------------------|
| <b>cs.vu.nl</b>      | <b>NS</b>   | <b>solo.cs.vu.nl</b> |
| <b>solo.cs.vu.nl</b> | <b>A</b>    | <b>130.37.21.1</b>   |

- **Part of the description for the *vu.nl* domain which contains the *cs.vu.nl* domain.**