

**CompSci 131**

# **Parallel and Distributed Systems**

**Prof. A. Veidenbaum**

# Today' s topics

- Coordination in DS
- Clocks
- Clock Synchronization
- Reading assignment:
  - Today: 6.1
  - Next time: 6.2
  - Complete the assignment before next class

# Last Lecture Covered

- Naming in DS

# Coordination

- **Activities of multiple processes in DS need to be coordinated.**
  - **Saw an example of synchronization in SMP systems**
- **DS have a much more complex case:**
  - **A general event ordering**
- **How to deal with this?**

# Synchronization

- In the context of SMPs
  - Mutual exclusion
    - » Enforce one-at-a-time access to a code section
      - Or object
  - Post/Wait
    - » Ordering enforcement – Post before Wait
- In the context of DS
  - » A somewhat simple case: Send/Recv in MPI
    - very similar to Post/Wait
- Barriers in both types of systems
- DS have a much more complex case:
  - general event ordering

# Synchronization in DS

- One way to achieve event ordering is using time
  - Need to know the actual time *very precisely*
- Often just need to know the relative order
  - Does even A occur before event B
- The question is, how to figure out this order?
  - Will look at some general solutions

- One solution: a central coordinator
- But – processes *must* agree on which one it is
  - Statically
    - » What are the issues in using a static coordinator?
  - Dynamically
- ***Can elect*** one process as a coordinator
  - How to run such an election?


# Time-based ordering

- **A computer has a clock**
  - **Driven by complex analog circuitry**
    - » Includes an accurate oscillator
    - » Generates “clock ticks”
  - Ticks are counted for a timer interrupt
  - Today there is also a 64b performance counter
    - » Continuously counts *clock cycles*
- **Processes on the computer see the same clock**
  - Can read clock register or value
- ***Thus no need for anything else to determine order***
  - What about multi-cores?



# Clock Synchronization in DS

- Each machine has its own clock, time may differ
- Programs like make have a problem
  - What is local time?




- **Clock synchronization** makes clocks “the same” on all nodes
  - How close is “the same”?

# Measuring time

- **Modern clocks are only about 600 years old**
  - Before it was just seasons or solar/lunar days and months
  - Romans had a solar “time of day clock”
- **Atomic clocks were introduced in 1948**
  - Based on Cesium 133 electron state transitions
    - » Under microwave radiation
  - accuracy =  $2 \text{ to } 3 \times 10^{-14}$ th, i.e. 0.000000000000002 Hz
- **US Naval Observatory operates ~70 clocks**
  - Coordinated Universal Time is the average

# Atomic Clocks



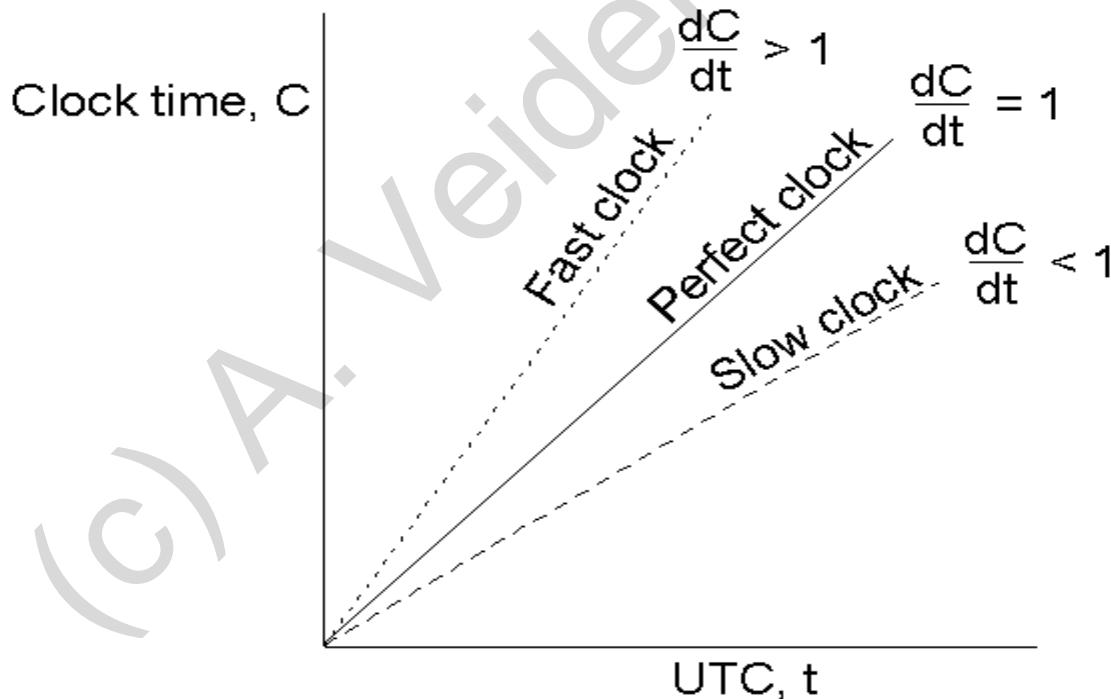
- **TAI (French!) seconds are of constant length, unlike solar seconds**
  - Leap seconds are introduced to keep in phase with the sun
- **The result is *Universal Coordinated Time* (UTC)**
  - Replaced Greenwich mean time
  - Is the civilian time keeping
- **Broadcast by NIST on WWW (radio station ID), 1msec accuracy**
  - But 10msec due to prop delay variations...

# GPS

- **Uses satellites with several atomic clocks each**
  - These are periodically calibrated from Earth
- **A satellite broadcasts its position plus a time stamp**
  - They are geostationary
  - Receiver computes delay from satellite  $i$ , plus time correction
    - »  $\Delta_i = (T_{\text{now}} - T_i) + \Delta_r \Rightarrow \text{distance } D_i = c \cdot (T_{\text{now}} - T_i)$
  - Now triangulate to find a receiver position
- **4 satellites give you position and time**

# Clock Synchronization

- Clocks drift with respect to UTC time  $t$ , i.e.  $C(t) \neq t$
- Can define drift rate (max is  $\pm x$ )
- Need to check, synchronize clocks frequently to avoid drift
- May be used with WWW (or not)




# Back to Computer Clocks

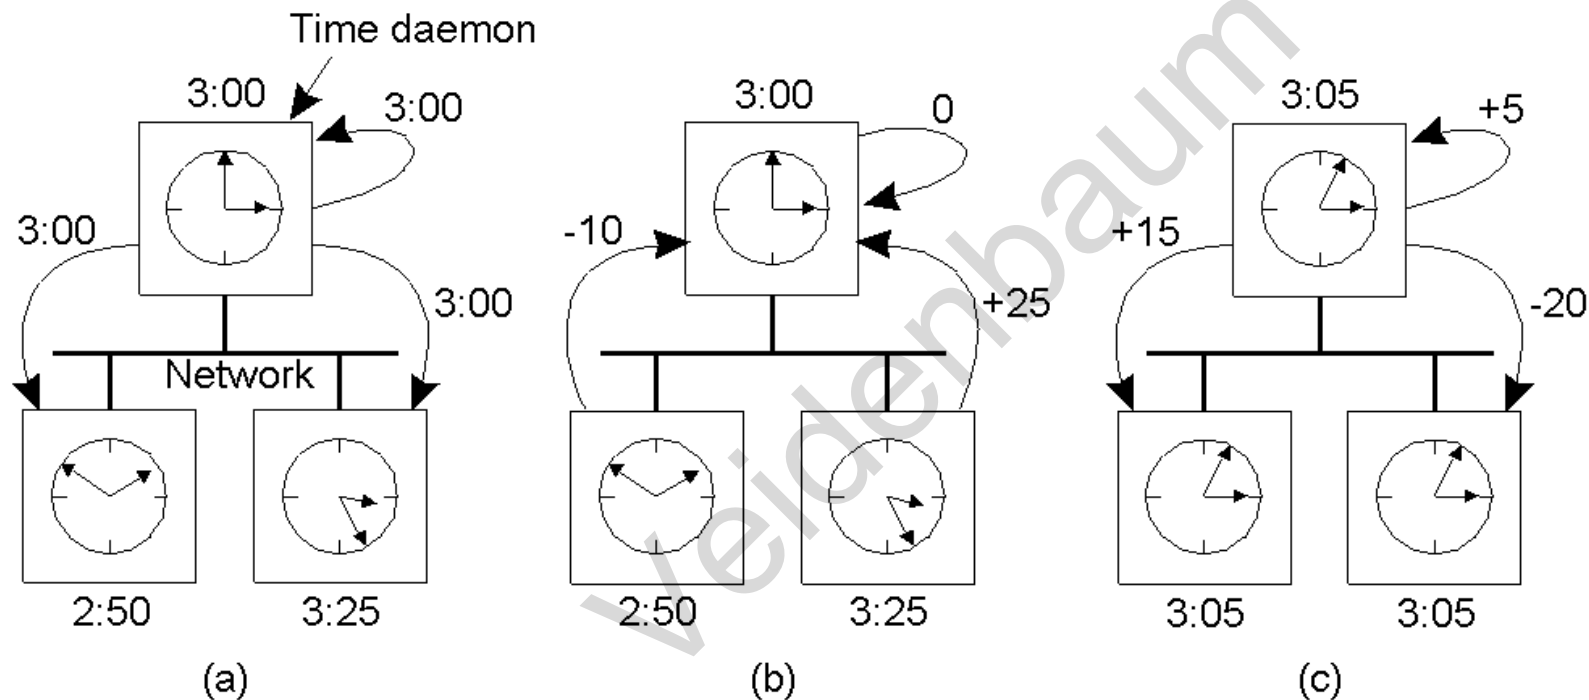
- Variation gets them out of sync
- A crystal oscillator is very precise, *but not perfect*
  - Oscillators with the same nominal frequency differ
    - » Error accumulates over time
- The counter generates an interrupt after N sec
  - $N = 1$  or 10 msec today
- Sftw handler adds 1 to counter C. C is the local clock
  - Interrupt service time may also vary!
- Thus need clock synchronization!

# Network Time Protocol (NTP)

- Get the current time from a time server
- Send time stamp  $T_0$ , server records arrival time  $T_2$ , sends at  $T_3$ 
  - Server returns  $T_2$  and  $T_3$
  - $\Delta$  to Tserver:  $T_2 = (T_0 + \delta) + T_{prop}$ ,  $(T_1 + \delta) = T_3 + T_{prop}$ 
    - » Assumes:  $T_{prop} \approx T_2 - T_0 \approx T_3 - T_1$ ,  $(T_i + \delta)$  is client's adjusted time
- Buffer 8 pairs  $\{\Delta, \delta\}$ , **pick minimal  $\delta$  and corresponding  $\Delta$**
- Adjustment has to be gradual and cannot go backwards
  - To slow a clock down add less time on each timer interrupt



# The Berkeley Algorithm



- a) The time daemon asks all the other machines for their clock values
- b) The machines reply
- c) The time daemon tells everyone how to adjust their clock

- **Centralized!**



# Synchronization in Wireless nets

- **Reference Broadcast Synchronization (RBS)**
  - One sender node  $S$ , multiple receivers
  - Does not aim to sync to UTC
    - » *Just figure out receiver nodes' pairwise offsets*
- **Assumes one-hop routing**
  - Thus prop delay is basically constant
- **But there is a receiver clock offset  $\delta$  (to  $S$ )**
  - all receiver nodes *exchange*  $T_{recv_i}$
  - Can compute the relative (average) offset
    - » *As well as individual offsets*