

CompSci 131

Parallel and Distributed Systems

Prof. A. Veidenbaum

Today' s topics

- Mutual exclusion
- Reading assignment:
 - Today: 6.3
 - Next time: 6.4
 - Complete the assignment before next class

Last Lecture Covered

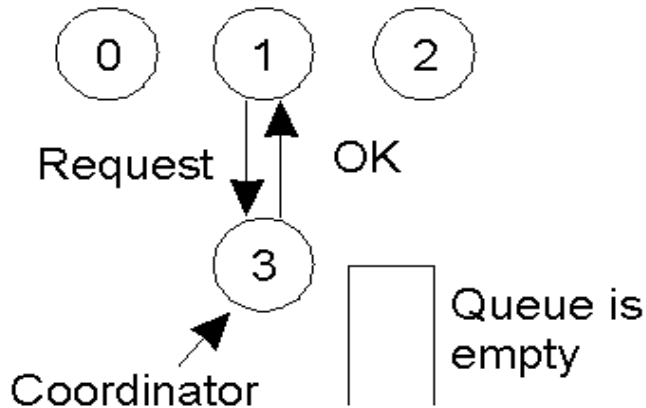
- Logical Clocks
- Vector Clocks

Mutual Exclusion *in DS*


- Same goal: protect access to critical resources
- Used semaphores, etc in a single system
 - Harder in distributed systems
 - » Because it happens on different systems
 - No common memory, clock, etc. etc.
- Algorithms for mutual exclusion in DS have to be
 - Fair, with no starvation
 - Deadlock-free
 - Fault-tolerant

- **Two main approaches in distributed systems**
 - **Token-based – pass a token around in a deterministic way**
 - » **Fair**
 - » **What about fault tolerance?**
 - **Permission-based**
 - » **Get permission from all other processes/nodes**
 - For instance, send a request and wait for all others to reply
- **Let us first look a different approach though:**
 - **A centralized algorithm**
 - » **Have one node that can grant permission to a lock**
 - Called a *coordinator*
 - Now can use an SMP-like algorithm
 - » **Issues:**
 1. ?
 2. ?


A Centralized Algorithm



(a)



(b)



(c)

- a) Process 1 asks the coordinator for permission to enter a critical region
 - Permission is granted
- b) Process 2 then asks permission to enter the same critical region.
 - The coordinator does not reply, but queues up the request
- c) When process 1 exits the critical region, it tells the coordinator
 - Coordinator then replies to first in queue – 2 in this case
- Queue holds state of CS requests

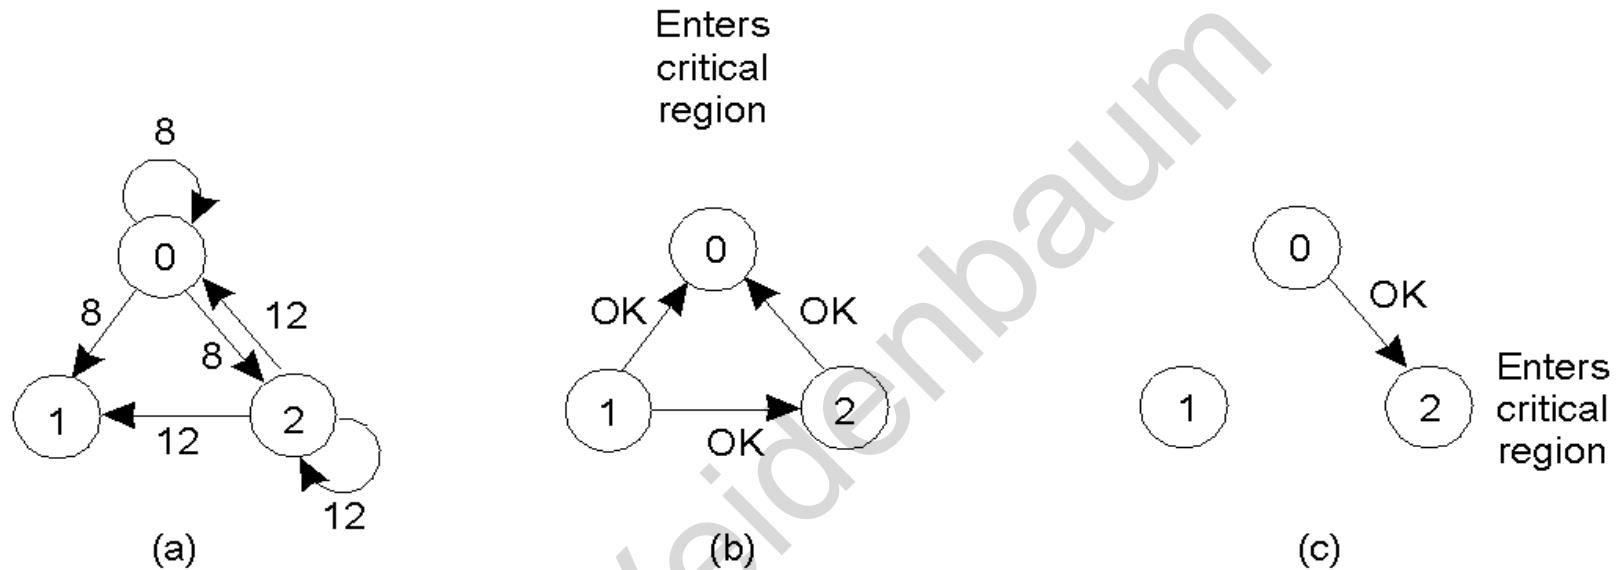
Problems

- **Fault tolerance**
 - What happens when the coordinator goes down?
- **Scalability**

A Permission-based Distr. Algorithm

- A request to lock a region is sent to **all** with a time stamp and PID
 - Assumes reliable message delivery
- Three cases are possible
 1. A receiver is not accessing, nor wants to access the lock
 - » It ACK's
 2. A receiver is currently holding the lock
 - » It does not ACK and queues the request
 3. A receiver wants to access the lock concurrently
 - » It is waiting for ACKs to its request
 - » Uses timestamps
 - » Process with the lowest timestamp wins
 - Lower pid, $\min(i,j)$, wins when $ts_i = ts_j$
 - » The receiver does not reply if it wins
- When a process is done, it sends an ACK to those waiting.


Example




- a) Processes 0 and 2 try
- b) A request to lock a region is sent to **all** with a time stamp and PID
- c) Process 0 wins
 - » does not reply to other requestors
- d) When process 0 is done, it sends an OK, so 2 can now enter the critical region.

- **No deadlock or starvation**
- **Reliability issue is even worse**
 - **Any process failing will not send an ACK**
 - » **One solution – do ACK and NACK, use timeouts**
- **What about scalability?**

A Token Ring Algorithm



(a)



(b)

- a) An unordered group of processes on a network.
- b) A logical ring is constructed in software
- c) A single token is given to process 0 at initialization
 - Circulates around the ring
- d) A process possessing the token can access resource

- **Can make token circulate fast if no action pending**
- **Correct and fair**
- **Fault-tolerance is hard**
 - **What happens if token is lost?**

A Decentralized Algorithm

Essentially extends the centralized algorithm

- a) Replicate a resource n times,
 - each replica has its own local coordinator
- b) A process sends requests to all n coordinators
 - And does not reply to other requestors
- c) The process can enter a critical region if it receives $k > m/2$ replies
 - Otherwise it waits and tries again
- Why is this correct?
 - What does correct mean?
- This approach is only probabilistically fault tolerant
 - To single coordinator failure

Comparison

Algorithm	Messages per entry/exit	Delay before entry (in message times)	Problems
Centralized	3	2	Coordinator crash
Distributed	$2 (n - 1)$	$2 (n - 1)$	Crash of any process
Token ring	0	0 to $O(n)$	Lost token, process crash