

CompSci 131

Parallel and Distributed Systems

Prof. A. Veidenbaum

Today' s topics

- Elections in large-scale systems
- Location Systems
- Reading assignment:
 - Today: 6.4-6.5
 - Next time: 6.7 7.1-7.2
 - Complete the assignment before next class

Last Lecture Covered

- Election algorithms

Election in Large-Scale Systems

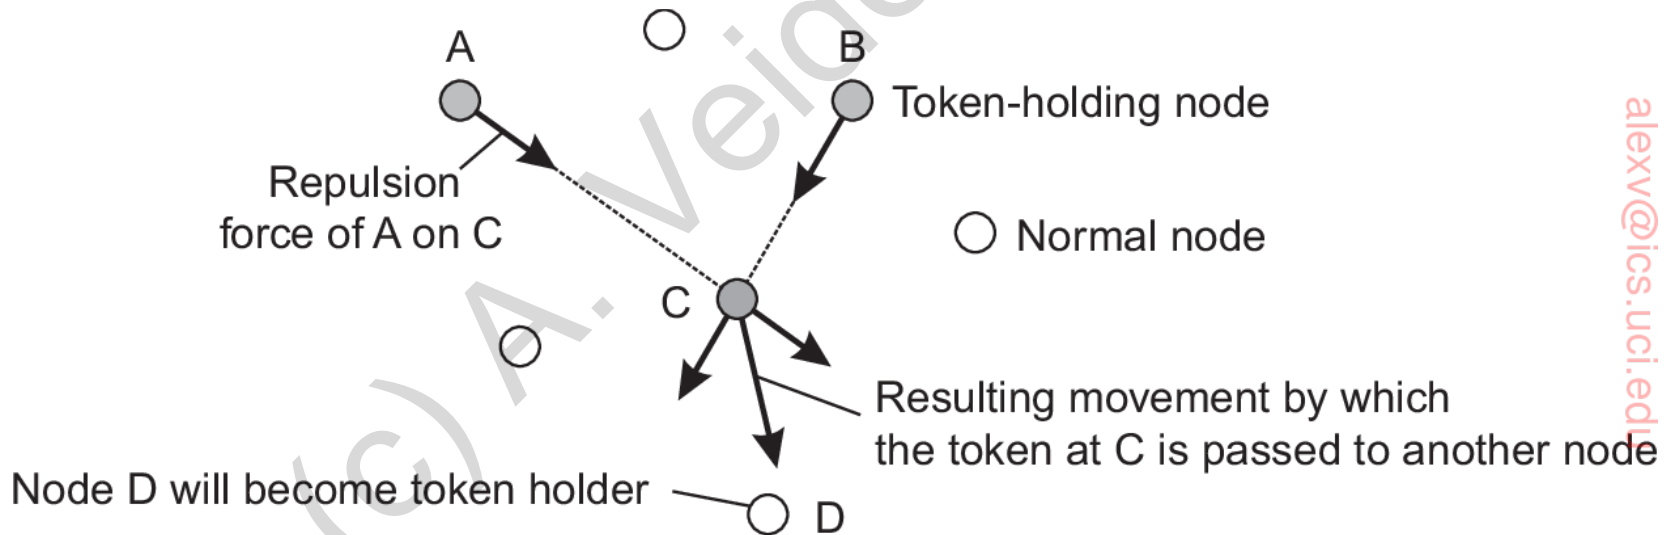
- So far looked at algorithms to elect a single leader
 -
- Sometimes may need to elect *more than one* leader
 - Example – super-peer selection in P2P systems
- Requirements for super-peer (S-P) selection
 1. Normal nodes should have low-latency access to S-Ps
 2. S-Ps should be evenly distributed across the network
 3. There should be a predefined portion of S-Ps relative to the total number of nodes in the overlay network.
 4. Each S-P should not need to serve more than a fixed number of normal nodes.

Election in Large-Scale Systems (2)

- Approach: position N nodes in an m -dimensional space
 - 2^m nodes in the system, N S-Ps
 - The space is geometric -
- Spread N tokens to N randomly selected nodes
 - No more than one token per node
- Each token represents a repelling force
 - Other tokens are *likely* to move away from it
 - » S-Ps should be evenly distributed across the network
- Nodes then need to know where the other tokens are
 - How to do this?

Election in Large-Scale Systems (3)

- Use a gossip protocol to initially disseminate a token's force throughout the network
- Then look at forces acting on a node
 - Shaded nodes have a token



Epidemic protocols for info dissemination

- Disseminate information as viral infections do
- An epidemic protocol classifies nodes as
 - *Infected* if it has data and is willing to disseminate
 - *Susceptible* if it has not yet seen the data
 - *Removed* if it is no longer willing/able to spread data
- An epidemic protocol works as follows
 - A node picks another node and pulls update from it
 - » Or they exchange updates
 - A round = period during which every node initiated an update once
- In $O(\log(N))$ rounds can propagate an update to all

Gossip protocols for info dissemination

- A variant of an epidemic protocol
 - Disseminate information as rumors are spread
- An gossip protocol works as follows
 - A node picks another node and sends it an update
 - A node that received an update will start spreading it too
 - A node stops when it reaches a node that already got the update
- This is a “push” algorithm and it is not a great
 - Pull is much better
- The gossip algorithm is not guaranteed to update all nodes in a large system

Location Systems

- So far did not worry about location of nodes
 - Even super-peers were kinda randomly pushed around
-
- Sometimes need to know proximity
 - E.g. which update or ftp server to contact
- Use location based techniques to place processes
-
- One approach – use GPS
 - But as we have discussed, it may not work for every node
- What else can we do?

Location Systems (2)

- Use logical , proximity-based locations
- Most common approach is to use internode latency
 - Can be measured
 - » Recall how clock synchronization algorithms compute latency
- Consider a replicated server Q with N replicas
 - Q can redirect the request to the closest replica
 - » Or can have delay to all replicas available and go directly
- Another example is the problem of replica placement
 - Compute positions of N servers that minimize average client-replica time

Location Systems (3)

- Centralized positioning works by contacting nodes with known locations
 -
- Again, use internode latency to estimate distance to these nodes with location info
- Problems:
 - Latency measurement accuracy and variability
 - A different position is obtained if known nodes change
 - » One of them may have gone down
 -
- Can use decentralized positioning
 - Assumes nodes are connected by springs and compute what happens

m

Location Systems (4)

- In decentralized positioning a node P_i repeatedly execute the following steps
 - An apostrophe indicates a vector in coordinate space
 - Measure the latency d_{ij} to node P_j , and also receive P_j 's coordinates x'_j .
 - Compute the error $e = d_1(P_i, P_j) - d_2(P_i, P_j)$
 - » A perturbation of P_i coordinate to current coordinates
 - Compute the direction $u' = u(x'_i - x'_j)$
 - Compute the force $F = e * u'$
 - » Force proportional to distance
 - Adjust own position by moving along the force vector
 - » $x'_i \leftarrow x'_i + \delta * u'$
 - Selecting δ is a key issue, sometimes use an adaptive value