

CompSci131

Parallel and Distributed Systems

Prof. A. Veidenbaum

Today's topics

- Gossip-based coordination
- Consistency and replication
- Consistency models
- Reading assignment:
 - Today: 6.7, 7.1-7.2
 - Next time: 7.2

Last Lecture Covered

- Elections at scale, with multiple leaders
- Epidemic and gossip protocols


Gossip based coordination


- **Examples of using epidemic protocols**
 - Aggregation
 - Large-scale peer sampling
 - Overlay construction
- **Aggregation: collect or aggregate info from all nodes**
 - Let's consider it in terms of numerical values
 - A node P_i has a number v_i , contacts P_j which has v_j
 - Both update $v_i, v_j = (v_i + v_j) / 2$
 - Eventually all nodes will have the same value – the average
- **A variation - P1 starts with 1, every other node with 0**
 - The result is $1/N$, where n is the number of nodes
 - Now all nodes know N

Gossip based coordination (2)

- How does a node choose a peer randomly?
 - Need to know all nodes in the network
 - » Works for small networks only
- Solution: a peer sampling service (PSS)
 - Fully decentralized
 - Can be done using epidemic protocols
 - Using *partial views*
- A partial view is a list of c neighbors a node maintains
 - Ideally each node is a randomly chosen, *live* node
 - » Nodes regularly exchange entries from their partial view
 - » Each entry has an age associated with it

PSS (2)

- The following operations are used
 - selectPeer: Randomly select a neighbor from the local partial view
 - selectToSend: Select some other entries from the partial view,
 - » add to the list intended for a selected neighbor
 - Includes its own entry
 - » List size is $c/2 + 1$
 - selectToKeep: Add received entries to partial view
 - » remove repeated items and shrink the view to c items.
 - ReceiveFromAny
 - Two ways to construct new views upon receipt
 - Discard the entries they sent to each other, i.e. swap them
 - Discard as many old entries as possible
 - » Still maintaining a total of c
- 



PSS (3)



1 <u>selectPeer</u> (&Q);	1
2 <u>selectToSend</u> (&bufs);	2
3 <u>sendTo</u> (Q, bufs);	3 <u>receiveFromAny</u> (&P, &bufr);
4	4 <u>selectToSend</u> (&bufs);
5 <u>receiveFrom</u> (Q, &bufr);	5 <u>sendTo</u> (P, bufs);
6 <u>selectToKeep</u> (p_view, bufr);	6 <u>selectToKeep</u> (p_view, bufr);

Diagram illustrating the sequence of operations for two peers (left and right) in the PSS (3) protocol. Red arrows indicate the flow of data and control:

- Peer 1 (left) sends to Peer 2 (right) at line 3.
- Peer 2 (right) receives from Peer 1 (left) at line 3.
- Peer 2 (right) sends to Peer 1 (left) at line 5.
- Peer 1 (left) receives from Peer 2 (right) at line 5.

- As long as peers regularly exchange partial views, **selecting from a partial view is indistinguishable from randomly selecting from the entire network**

Consistency

- Data replication creates multiple copies
 - For performance, fault tolerance
- It is multiple copies that cause the consistency problem
 - A modification of one copy requires update of all copies
 - How can it be done?
 - » Must prevent old (inconsistent) values from being used
- Can try the following: send update to all replicas
 - while doing a local update, and wait for all to ACK
 - Problems?
- Will also consider consistency for shared memory
 - The issue here is what to expect when reading/updating data


- **Data consistency is hard to implement efficiently at large scale**
 - This is because all copies need to be updated synchronously
 - Sometimes, a *simpler* model will suffice
- **The consistency problem has no general solution**
 - It can only be solved by loosening some requirements!
 - Or for special cases, like shared memory systems
- **We will look at**
 1. Keeping replicas consistent
 2. Content distribution to replicas

(Replicated) Data Consistency

- An intuitive definition: data is consistent when ***ALL* copies *ALWAYS* have *the same* value**
 - Hard/expensive to achieve
 - » Recall totally ordered multicast
 - Requires an atomic update: update all copies “at once”
 - » May not be scalable
 - Replication was supposed to solve scalability problems!
- Can we loosen the consistency constraints?
 - For instance, do not require the atomicity of update
 - **BUT:**
 - » require that a DS still has “*correct*” execution/results

Data-Centric Consistency

- A logical data store - shared memory, file, database
 - physically distributed and replicated across multiple nodes
 - » Each node can access a local (or nearby) copy
- A *write* is an operation that changes data
- Otherwise, it's a *Read*



A Consistency Model

- A contract between processes and the data store
 - If processes obey certain rules the store will work “correctly”
- Intuitively, correctly means a read gets the most recently written data value
 - That of the “last” write
 - » Or a well-defined approximation
- What is the last write?
 - Hard to say without a global clock
- A consistency model defines/restricts which (write) values a read can return
- Will look at such models next

Continuous Consistency

- An application defines consistency it can tolerate
- Three types of *inconsistency* can be distinguished
 - A numerical value deviation
 - » Define a tolerance range if data behavior is known
 - “Staleness” between replicas
 - » Replica is several updates behind
 - It may be acceptable to have a slightly older data...
 - Deviation in update ordering
 - » Among different replicas
 - Typically in a local update, while waiting for global agreement
 - Often used with roll-back
- A **consistency unit (conit)** is a data unit on which consistency is defined – an update unit
 - a byte, a word, a cache line, a page, etc

Example

- Tracking a fleet of cars
 - 3-variable conit, initially all values = 0

Replica A

Conit

d = 558 // distance

g = 95 // gas

p = 78 // price

Operation	Result
< 5, B> g ← g + 45	[g = 45]
< 8, A> g ← g + 50	[g = 95]
< 9, A> p ← p + 78	[p = 78]
<10, A> d ← d + 558	[d = 558]

Vector clock A = (11, 5)

Order deviation = 3

Numerical deviation = (2, 482)

Replica B

Conit

d = 412 // distance

g = 45 // gas

p = 70 // price

Operation		Result
< 5, B>	g ← g + 45	[g = 45]
< 6, B>	p ← p + 70	[p = 70]
< 7, B>	d ← d + 412	[d = 412]

Vector clock B = (0, 8)

Order deviation = 1

Numerical deviation = (3, 686)

- Shaded values are committed

Example (2)

- Order deviation is how many entries are not committed at a node
 - Can specify maximum allowable deviation
- Can also define a numerical limit on deviation
 - What update values have not been seen yet...
- How to know the deviation?
 - Basically need separate communication
 - » This assumes that this is cheaper than synchronizing replicas!