

**CompSci131**

# **Parallel and Distributed Systems**

**Prof. A. Veidenbaum**

# Today's topics

- Consistency and replication
- Consistency models
- Reading assignment:
  - Today: 7.2
  - Next time: The rest of 7.2

# Last Lecture Covered

- Uses of epidemic protocols
- The consistency problem
- Continuous consistency

# Consistency

- Existence of multiple replicas in a data store may lead to replicas having different “values” - *an inconsistency*
  - Any use of an inconsistent value may lead to program errors
  - *Strict consistency is impossible to achieve in DS*
- A **consistency unit** (conit) is a data unit on which consistency is defined
  - An “update unit”: a byte, a word, a cache line, a page, etc
- Use of conit can lead to “false sharing”
  - Conit contains multiple independently used data items
    - » Means consistency really did not have to be enforced at this level of granularity

# Notation

- Will define a Read set and a Write set at process  $i$
- $W_i(x)a$ : process  $i$  writes the value  $a$  to variable  $x$ 
  - $\text{Data\_Store}[x] = a$
- $R_i(x)b$ : process  $i$  reads  $x$ ,  $\text{Data\_Store}[x]$  returns a value  $b$ 
  - Initial value for data at an “address” is NIL
- This allows us to define Write/Read ordering and will help figuring out consistency
- Have been talking about read/write ordering for a **single item** –  $\text{Read}(x) / \text{Write}(x)$ 
  - **Consistency is more than that**, it also covers ordering of **Rd/Wr** for multiple items together

# (Cache) Coherence

- A consistency model defines ordering of accesses to all *data items* in the data store
  - *conit's*
- Coherence defines what happens on access to a single data item
  - A conit = cache line (cache block)
  - Can also have “false sharing”
- An example of false sharing
- In summary, want to achieve **a consistent ordering of operations** on shared, replicated data by all participating processes

# Relaxing Consistency Requirements

- Programmers can reason about concurrent programs even if the system is not strict
- For instance, the following is often assumed:
  - Concurrent programs should not assume anything about relative speed of processes
    - » Thus cannot know about event order
  - Synchronization must be used to achieve a certain event order
- *What other assumptions can be made that would still allow programmers to reason about concurrent or distributed programs?*

# Sequential Consistency

- Sequential Consistency was defined by Lamport for multiprocessors
- Definition
  1. The result of any execution is the same as if all Rd/Wr operations by all processes were executed in some sequential order

**AND**

  2. The operations of each individual process appear in this sequence in the order specified by its program
- In other words, the result of an execution may be any legal interleaving of accesses
  - Does not say that a read gets the result of most recent write
- *But all processes MUST see the same order*
- No reference to time!



# Sequential Consistency Examples

- a) Is it a sequentially consistent data store?
- b) What about this one?

Why?

- Time is the horizontal axis
  - Is it important?

P1:	W(x)a		
P2:	W(x)b		
P3:		R(x)b	R(x)a
P4:		R(x)b	R(x)a

(a)

P1:	W(x)a		
P2:	W(x)b		
P3:		R(x)b	R(x)a
P4:		R(x)a	R(x)b

(b)

## Let's look a little deeper

- 3 data items: x,y,z. Initialized to 0.
- Three concurrently executing processes
  - 6! Possible inter-leavings of the operations
  - Some are valid and some violate program order!
    - » Print before assignment
- A signature: values printed by processes P1,P2,P3
  - » In this order
    - 00 11 01

Process P1

Process P2

Process P3

---

x = 1;  
print ( y, z);

y = 1;  
print (x, z);

z = 1;  
print (x, y);

## 4 Valid Execution Sequences

- There are  $2^6$  possible *signatures* in this case
  - Not all of these are valid!
  - » 000000?
- Overall, many different ordering are perfectly legal!

```
x = 1;  
print (y, z);  
y = 1;  
print (x, z);  
z = 1;  
print (x, y);
```

Prints: 001011

Signature:  
001011

(a)

```
x = 1;  
y = 1;  
print (x,z);  
print(y, z);  
z = 1;  
print (x, y);
```

Prints: 101011

Signature:  
101011

(b)

```
y = 1;  
z = 1;  
print (x, y);  
print (x, z);  
x = 1;  
print (y, z);
```

Prints: 010111

Signature:  
110101

(c)

```
y = 1;  
x = 1;  
z = 1;  
print (x, z);  
print (y, z);  
print (x, y);
```

Prints: 111111

Signature:  
111111

(d)

# Reasoning About Consistency

- The example showed that multiple orderings are sequentially consistent
  - But not all
- A program should work with any valid order
  - A program that works for some but not all is incorrect
    - » It violates the contract with the data store
- Program order deals with how accesses to different data items are interleaved
- Data coherence looks at access to the same item