

CompSci131

Parallel and Distributed Systems

Prof. A. Veidenbaum

Today's topics

- Causal Consistency
- Grouping operations
- Reading assignment:
 - Today: the rest of 7.2
 - Next time: Cache coherence (lecture notes only)
 - After that: 7.3

Last Lecture Covered

- Sequential consistency

Causal Consistency

- Relaxes some of the sequential consistency rules
 - To improve performance
- Already saw causality in logical clock events
 - Data in a message used in computation, result sent out
- Can something like this be used to relax rules?
 - P1: $W(\underline{a})\underline{x}$
 - P2: $W(\underline{b})\underline{y}$
- The two writes are completely independent
 - Are not causally related, may be seen in different order...

Causal Consistency

- On the other hand, if
 - P1 performs $W(x)$,
 - P2 performs $Rd(x)$, $Wr(y)$
 - The two writes are causally related, may be dependent
- A necessary condition for causal consistency:
 - Writes that are potentially causally related must be seen by all processes in the same order.
 - Concurrent writes may be seen in a different order by different processes.

Causal Consistency Example

P1:	W(<u>x</u>) <u>a</u>			W(<u>x</u>) <u>c</u>
P2:		R(<u>x</u>) <u>a</u>	W(<u>x</u>) <u>b</u>	
P3:		R(<u>x</u>) <u>a</u>		R(<u>x</u>) <u>c</u> R(<u>x</u>) <u>b</u>
P4:		R(<u>x</u>) <u>a</u>		R(<u>x</u>) <u>b</u> R(<u>x</u>) <u>c</u>

- Is this valid for a sequentially consistent store?
- Is this a valid causally consistent store?

More Causal Consistency

P1:	W(x)a		
P2:		R(x)a	W(x)b
P3:			R(x)b <u>R(x)a</u>
P4:		<u>R(x)a</u>	<u>R(x)b</u>

(a)

P1:	W(x)a		
P2:		W(x)b	
P3:			R(x)b R(x)a
P4:		R(x)a	R(x)b

(b)

- a) Is this a casually-consistent store?
- b) What about this one?

More Causal Consistency

- Previous examples looked at one variable – x
- True causality is better illustrated by the next example
 - Assume prior values of x and y are NIL

P1:	W(x)a		
P2:		R(x)a	W(y)b
P3:			R(y)b R(x)?
P4:			R(x)a R(y)?

- What values can R₃(x) and R₄(y) return?
- How can this be implemented?
 - Maintain a dependency graph between accesses

Grouping Operations

- Even less strict models are possible and useful
- So far dealt with a conit R/W order
 - May not match the granularity of applications
 - Not all writes are important
- An example: a Critical Section (CS)
 - Only one process in the CS
 - operates on N variables in the CS
 - » only these N variables are important, not the rest
 - But the system may not know a process is in a CS!
- Solution: a system aware of synchronization

- Define “Enter_CS” and “Leave_CS” operations
 - Enter_CS guarantees the local store is up-to-date
 - » Stalls until this is true
 - Initiate update upon/prior to Leave_CS
- All reads and writes are grouped into an indivisible unit accessed and updated by one process
 - *The rest are not guaranteed to be consistent*
- Need to define the semantics of Enter and Leave
 - Done using shared synchronization variables, aka locks
 - » Operations on locks – Acquire and Release
 - Shared data items are associated with each lock
- - » Each shared data item is associated with at most one lock

- **Define a Synchronization Variable (SV)**
 - Each has a current “owner”
 - » the process that last acquired it
 - This process may enter and exit the CS many times
- **A process not currently owning the SV has to**
 - Send message to the current owner
 - » Requesting SV ownership **AND** all associated data
- **Several processes may simultaneously own an SV, *but in a non-exclusive mode***
 - They can only Read the SV, *not write it.*
- ***Note that we are basically requiring sequential consistency for locks***

- **The following criteria need to be met:**
 - **A lock acquire is not allowed to perform until all the updates to its guarded shared data have been performed**
 - » **With respect to the acquiring process**
 - **An exclusive mode access to an SV is allowed to perform only when no other process holds the SV**
 - » **Exclusively or non-exclusively**
 - **No read or write operation on data items are allowed to be performed until lock acquire has been performed.**
 - **A non-exclusive access to the SV may be performed only after any previous exclusive access has completed**
 - » **E.g. the SV owner updated the data and released the lock**

~~m~~

Entry Consistency

- A way to implement a form of grouping
- Associate a lock with each data item
 - $L(x)$ is a lock op for data item x
 - $U(x)$ is an unlock op for data item x

P1: $L(x)$ $W(x)_a$ $L(y)$ $W(y)_b$ $U(x)$ $U(y)$

P2: $L(x)$ $R(x)_a$ $R(y)$ NIL

P3: $L(y)$ $R(y)_b$

- Each process has a copy of a variable, BUT
 - This copy is not instantly or atomically updated
 - » Updated on locking
- *Not easy to program due to multiple locks!*