

CompSci131

Parallel and Distributed Systems

Prof. A. Veidenbaum

Today's topics

- Cache Coherence
- Coherence protocols
- Reading assignment:
 - Today: Cache coherence (lecture notes only)
 - Next time: more coherence, the rest of 7.2
 - After that: 7.3

Last Lecture Covered

- Causal Consistency
- Grouping operations

(c) A. Veidenbaum

Cache coherence vs consistency

- Caches are replicas of shared memory
 - But not full replicas, as $\text{Cache_size} \ll \text{Memory_size}$
 - A process (core) accesses its local cache only
 - » A unit here is a 64Byte block (aligned), aka cache line
 - Identified by its memory address
- Thus there is a consistency problem
 - But caches are hardware managed, so need simplicity
- Coherence is a special case of consistency
 - Deals with consistency for a single address/block
- Finally, a cache may not contain a requested block
 - Has to be brought in from memory or other replicas

Cache Terminology

- Processor P_i issues Rd/Wr requests to its local cache
- A cache hit:
 - P_i reads/writes an address A from C_i (i.e. finds it there)
- A cache miss:
 - P_i tries accessing A , but it's not in C_i
 - » Is read from memory into C_i at this point
- A write-back:
 - The space occupied by a modified block is needed, so the block is written back to memory
 - » Not when the write happens (a write-back cache)

Cache coherence protocols


- Coherence guarantees update of all existing replicas
 - *Of a given block*
- Many shared memory coherence protocols
 - *With hardware support*
- General idea –a write causes coherence enforcement
- Also possible in software
 - *Middleware-based systems with caches use software protocols for coherence*

SMP Protocols

- Some slides courtesy of Prof. J.-L. Baer
- Caches and memory are connected by a bus
 - This means every bus access acts as a broadcast
 - » E.g. like an ethernet cable
 - » Bus access is arbitrated, one unit at a time
- A number of possible design choices
 - They are variants of MOESI or MESI protocols
 - » 1st letters of states – Modified, Exclusive, Shared, Invalid

Cache Coherence: Review

- Initial state: P2 reads A; P3 reads A
 - From memory

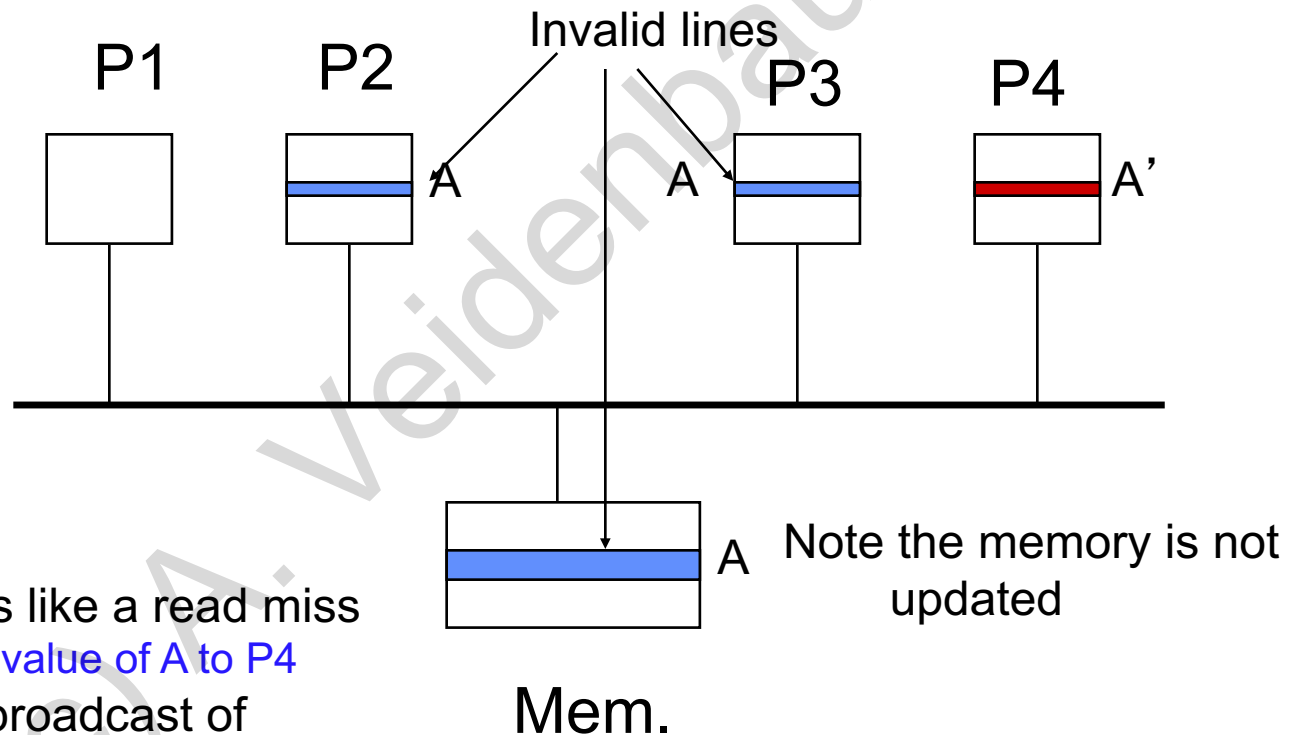


Cache coherence (shared-bus)

- Now P4 wants to write to A – *and writes are the issue!*
 - *It is a write miss for P4, it fetches the block from memory*
- Broadcast an invalidation message for address A
 - the address is snooped by the P2,P3 caches (and all others)
 - Seeing a write, P2 and P3 invalidate their copies
 - » This is called a write-invalidate protocol
 - Note that the copy in memory is not up-to-date any longer
 - Because it is write-back cache
- Now P4 writes to block A in its cache
 - *An alternative implementation is to request ownership on a read miss*

A Write-invalidate example

- P2 and P3 have read line A
- P4 has a write miss on an element of line A



- A write miss looks like a read miss
 - Brings the old value of A to P4
- Followed by the broadcast of an invalidation and a write

Memory Update

- What happens next time P2 wants to read A?
 - *Cannot fetch the block from memory!*
- But the read request is broadcast on the bus!
 - the address is *snooped* by P4's caches
 - » Now is a good time to update memory!
 - C4 can make C2 abort its read, then write back to memory
 - Now C2 will retry and get it from memory!

Snoopy Cache Coherence Protocols

- Need protocols to perform all these actions
 - And in the right order
- Use FSM and associate a state with each cache line
- Let us start with 3 states:
 - Invalid (I) - not present in this cache
 - Shared (S) - one or more unmodified, up to date copies
 - Modified (M) – The only up-to-date copy
- 4th or 5th state can be added for performance
 - MESI: E stands for Exclusive
 - MOESI protocols: adds O for Ownership

State Transitions for a Given Cache Line

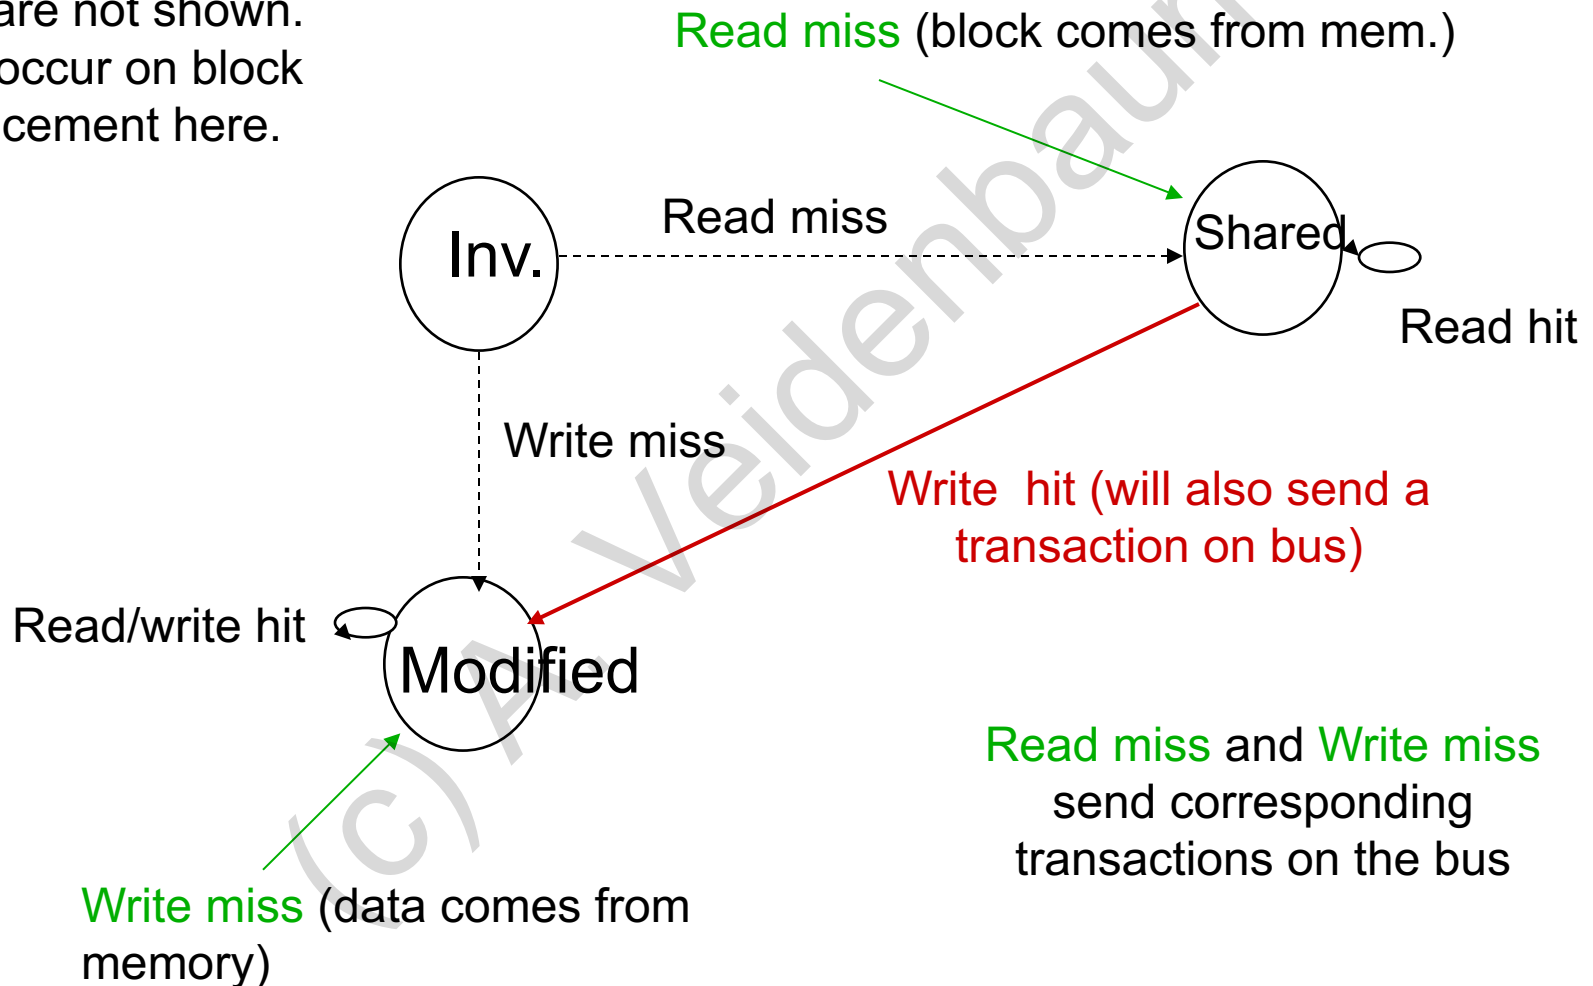
- Occur in a given cache in response to
 1. **Processor requests to its cache**
 - Read miss, write miss, write on shared line
 2. **Bus snooping of other processor actions on the bus**
 - Read miss by Q would make P's line transit from M to S
 - Write miss by Q might make P's line transit from S to I
 - » write invalidate protocol

Basic Write-invalidate Protocol (write-back write-allocate caches)

- Uses 3 states for each cache line
 - Invalid
 - Shared (read only – can be shared)
 - Modified (only valid copy in the system)
- Different state transitions are
 - Induced by the processor attached to the cache
 - Induced by snooping on the bus

3-State Protocol: Processor Actions

Transitions to Invalid state are not shown. They occur on block replacement here.



Basic 3 State Protocol: Transitions from Bus Snooping

The bus reads and writes
are by other caches

